

Parallel Programming in C with MPI and OpenMP

Michael J. Quinn



Chapter 2

Parallel Architectures

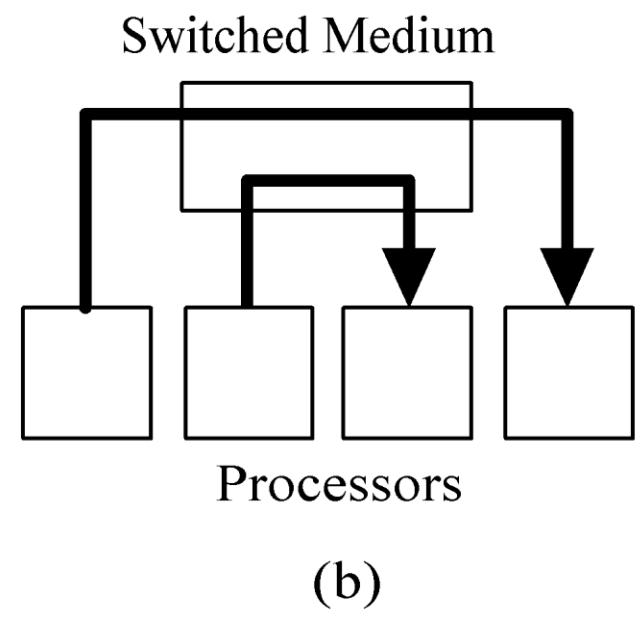
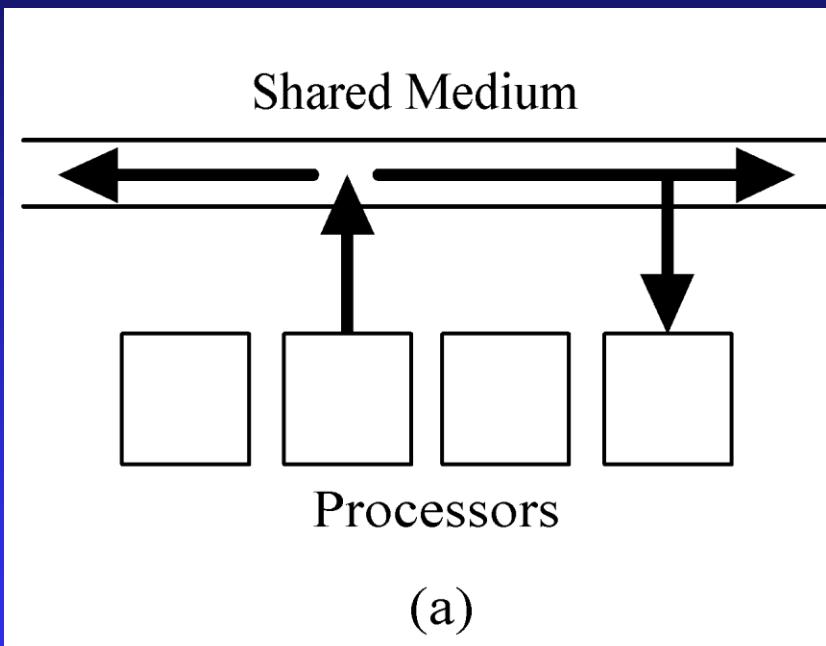
Outline

- Interconnection networks
- Processor arrays
- Multiprocessors
- Multicomputers
- Flynn's taxonomy

Interconnection Networks

- Uses of interconnection networks
 - ◆ Connect processors to shared memory
 - ◆ Connect processors to each other
- Interconnection media types
 - ◆ Shared medium
 - ◆ Switched medium

Shared versus Switched Media



Shared Medium

- Allows only message at a time
- Messages are broadcast
- Each processor “listens” to every message
- Arbitration is decentralized
- Collisions require resending of messages
- Ethernet is an example

Switched Medium

- Supports point-to-point messages between pairs of processors
- Each processor has its own path to switch
- Advantages over shared media
 - ◆ Allows multiple messages to be sent simultaneously
 - ◆ Allows scaling of network to accommodate increase in processors

Switch Network Topologies

- View switched network as a graph
 - ◆ Vertices = processors or switches
 - ◆ Edges = communication paths
- Two kinds of topologies
 - ◆ Direct
 - ◆ Indirect

Direct Topology

- Ratio of switch nodes to processor nodes is 1:1
- Every switch node is connected to
 - ◆ 1 processor node
 - ◆ At least 1 other switch node

Indirect Topology

- Ratio of switch nodes to processor nodes is greater than 1:1
- Some switches simply connect other switches

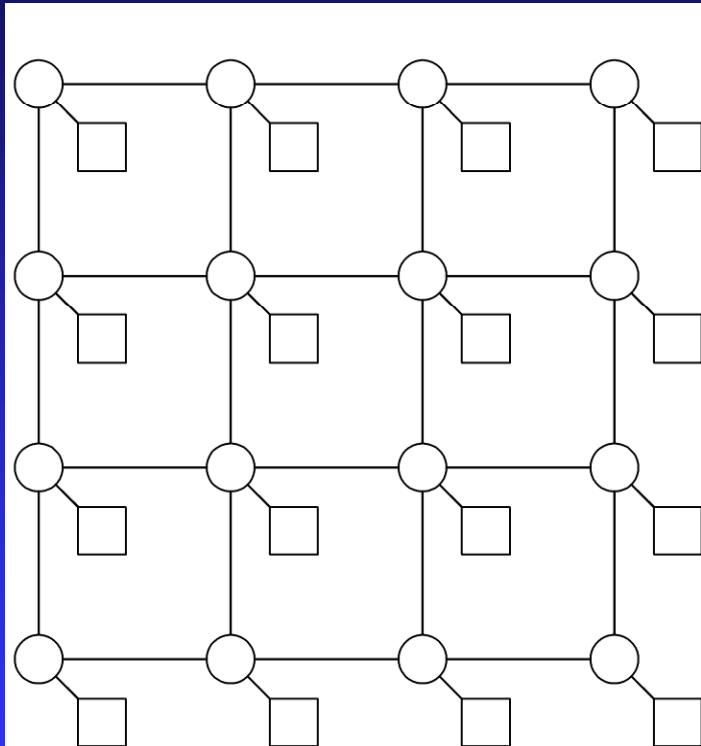
Evaluating Switch Topologies

- Diameter
- Bisection width
- Number of edges / node
- Constant edge length? (yes/no)

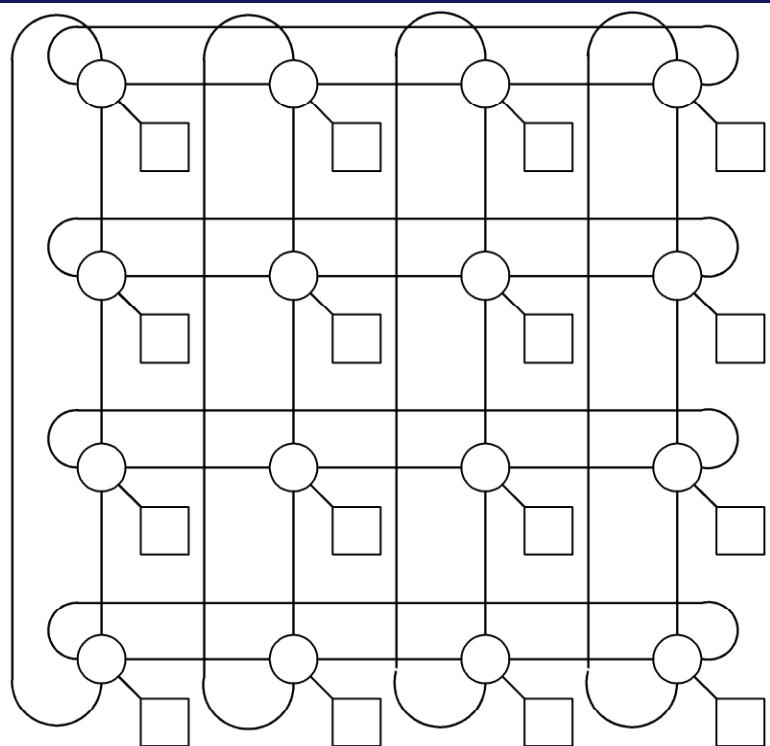
2-D Mesh Network

- Direct topology
- Switches arranged into a 2-D lattice
- Communication allowed only between neighboring switches
- Variants allow wraparound connections between switches on edge of mesh

2-D Meshes



(a)



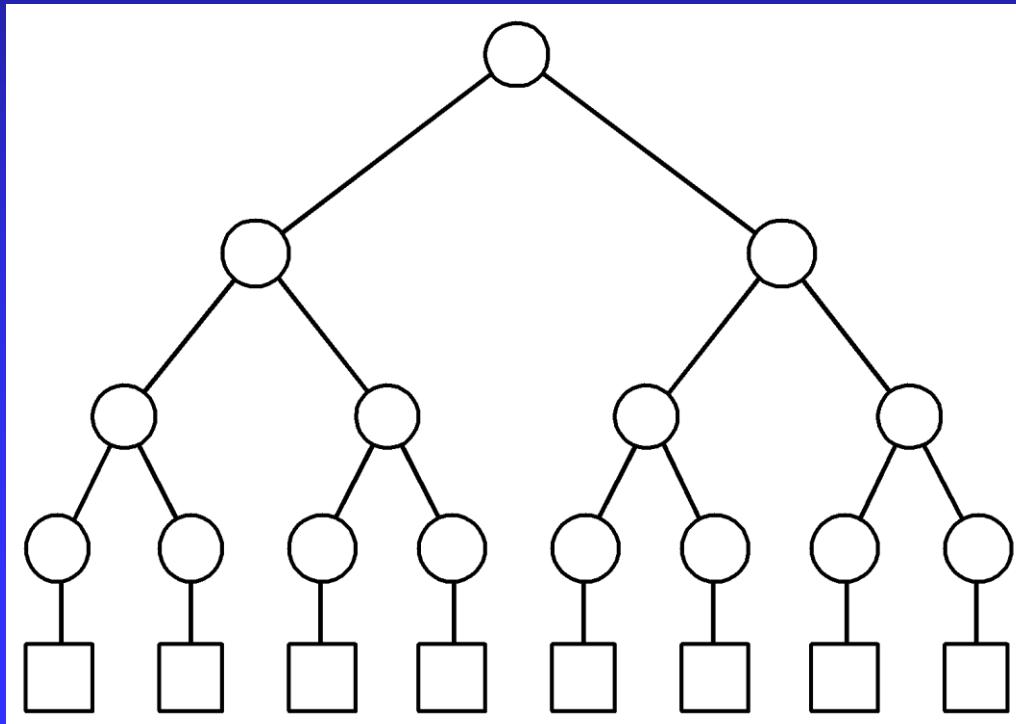
(b)

Evaluating 2-D Meshes

- Diameter: $\Theta(n^{1/2})$
- Bisection width: $\Theta(n^{1/2})$
- Number of edges per switch: 4
- Constant edge length? Yes

Binary Tree Network

- Indirect topology
- $n = 2^d$ processor nodes, $n-1$ switches



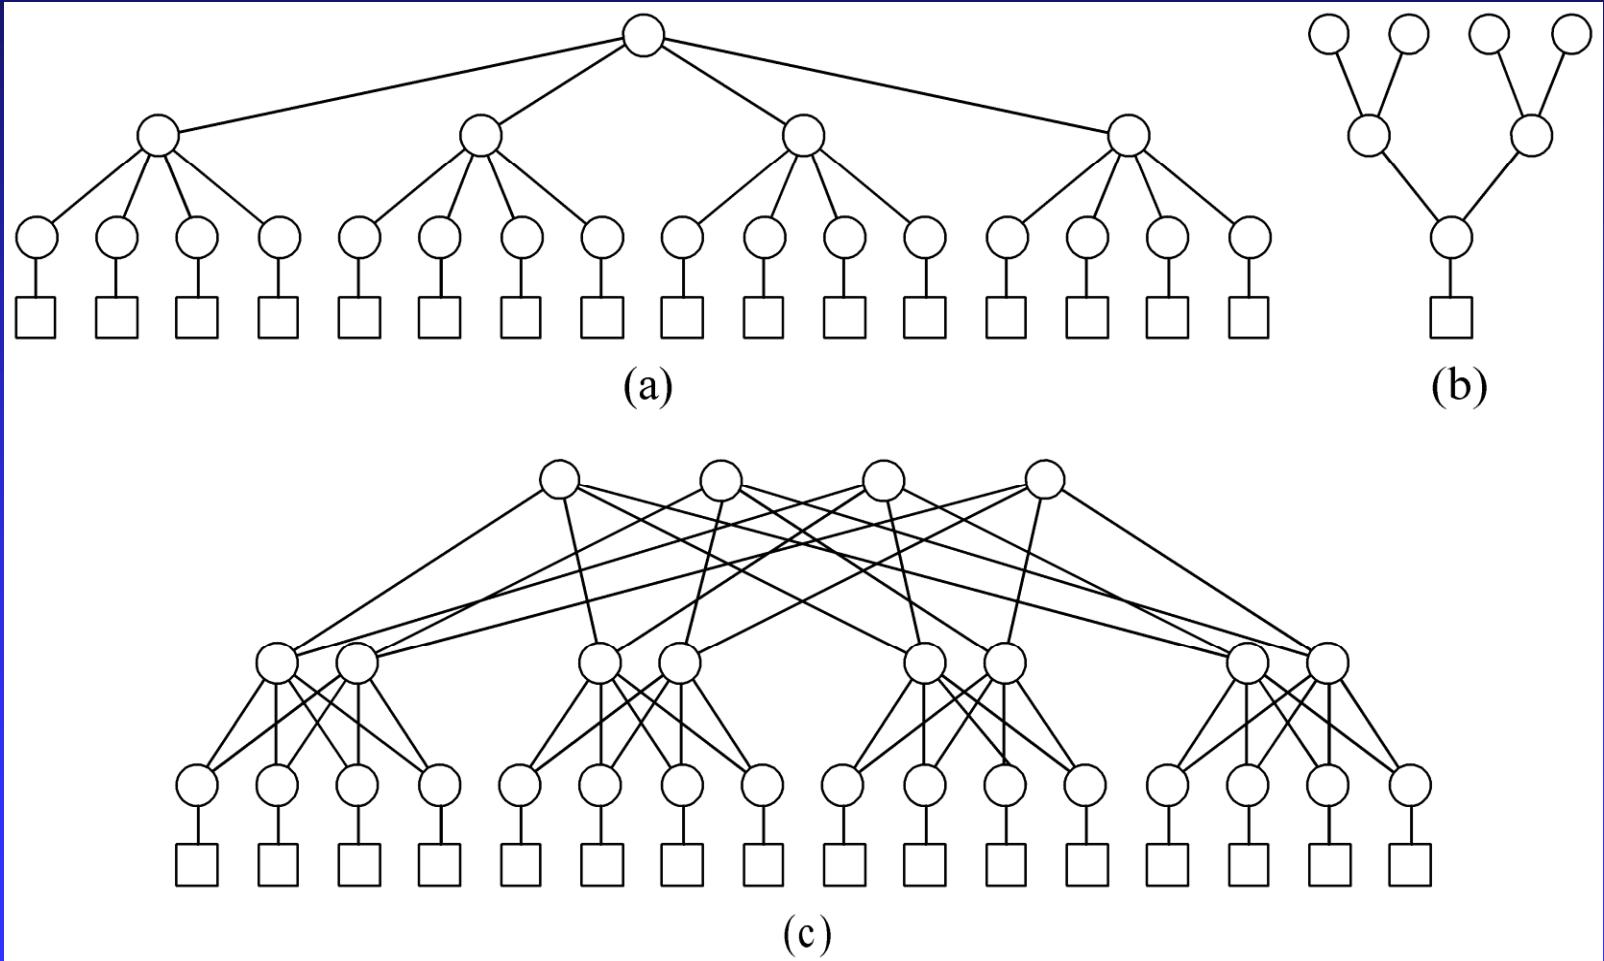
Evaluating Binary Tree Network

- Diameter: $2 \log n$
- Bisection width: 1
- Edges / node: 3
- Constant edge length? No

Hypertree Network

- Indirect topology
- Shares low diameter of binary tree
- Greatly improves bisection width
- From “front” looks like k -ary tree of height d
- From “side” looks like upside down binary tree of height d

Hypertree Network

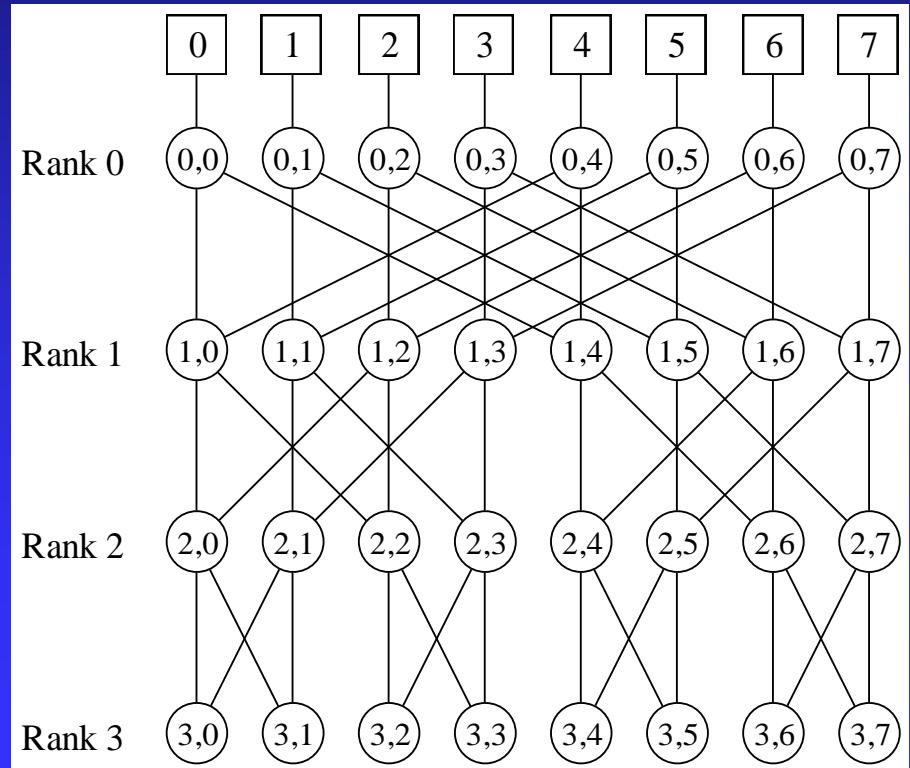


Evaluating 4-ary Hypertree

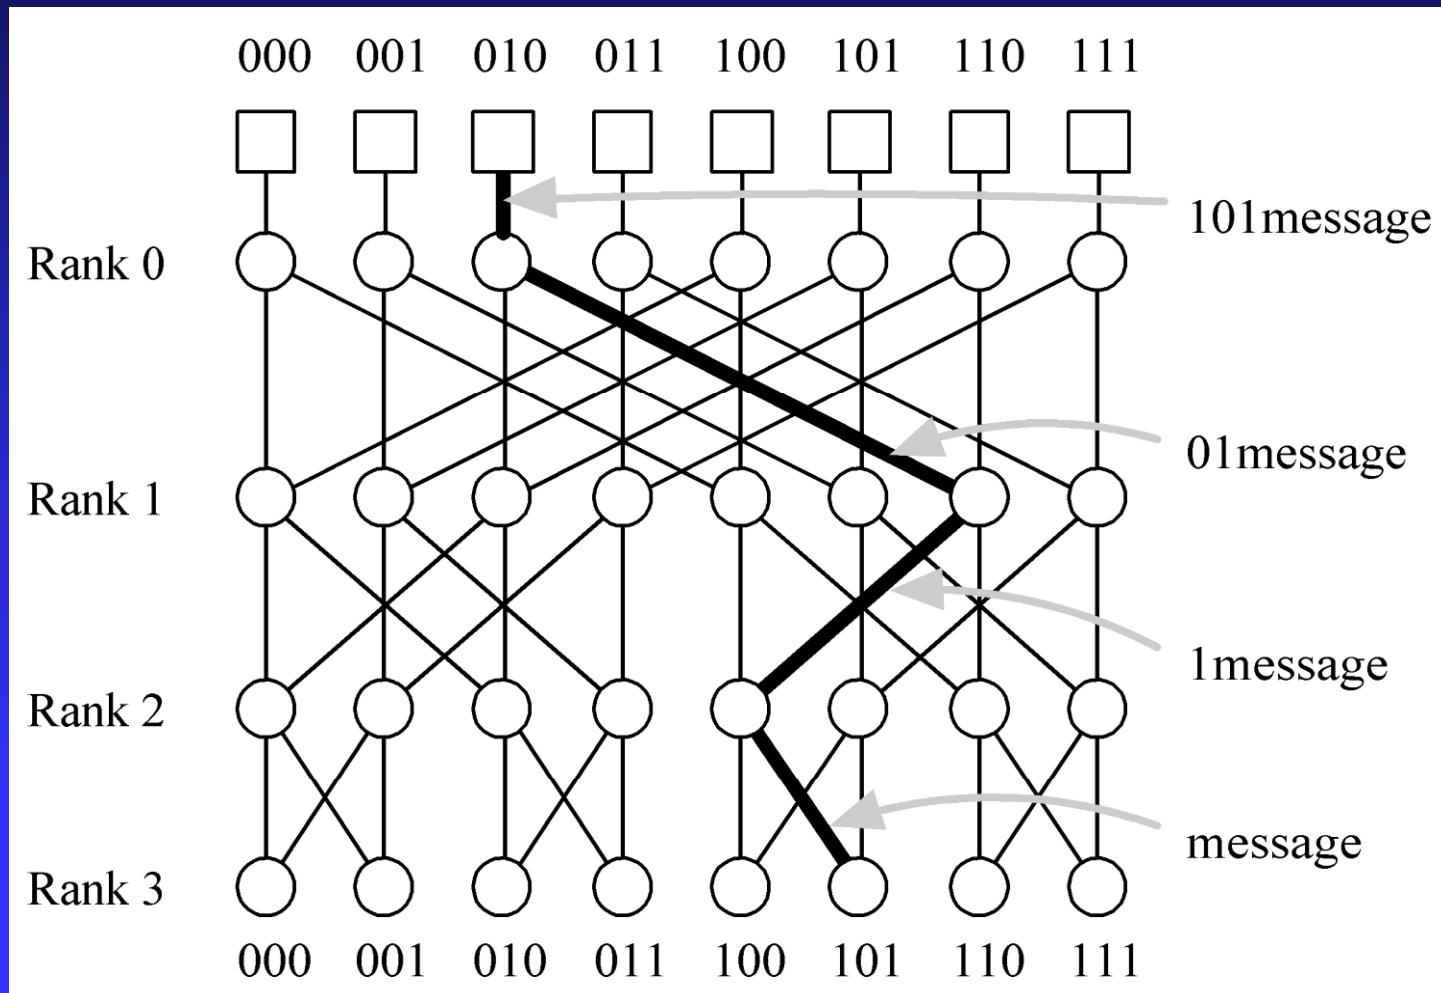
- Diameter: $\log n$
- Bisection width: $n / 2$
- Edges / node: 6
- Constant edge length? No

Butterfly Network

- Indirect topology
- $n = 2^d$ processor nodes connected by $n(\log n + 1)$ switching nodes



Butterfly Network Routing



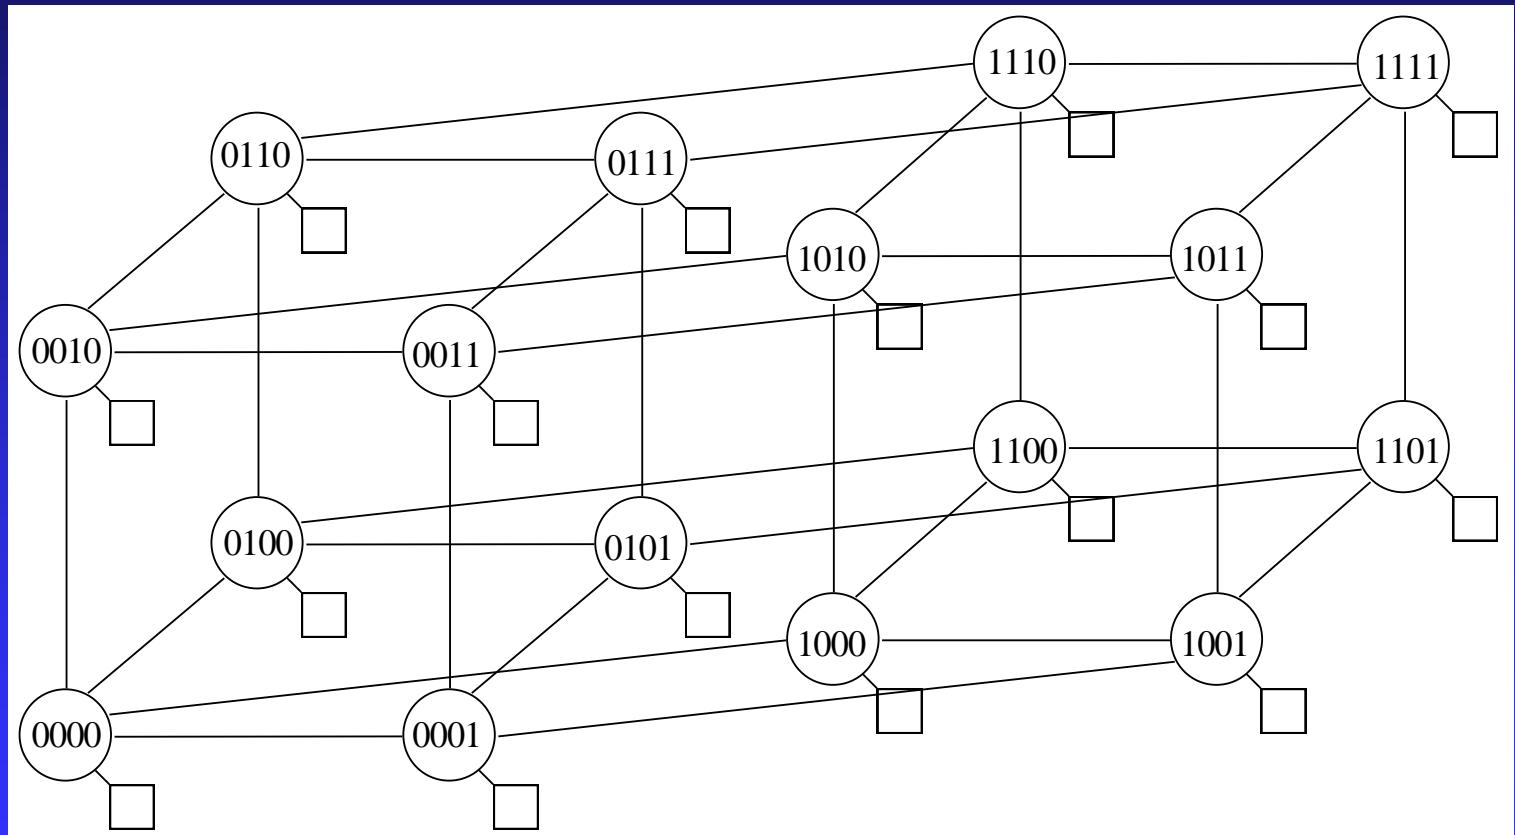
Evaluating Butterfly Network

- Diameter: $\log n$
- Bisection width: $n / 2$
- Edges per node: 4
- Constant edge length? No

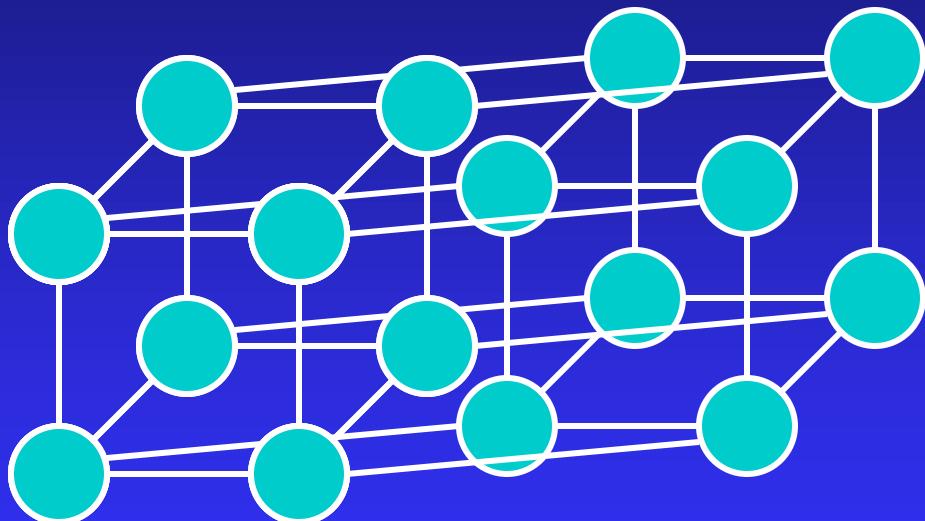
Hypercube

- Directory topology
- $2 \times 2 \times \dots \times 2$ mesh
- Number of nodes a power of 2
- Node addresses 0, 1, ..., $2^k - 1$
- Node i connected to k nodes whose addresses differ from i in exactly one bit position

Hypercube Addressing



Hypercubes Illustrated



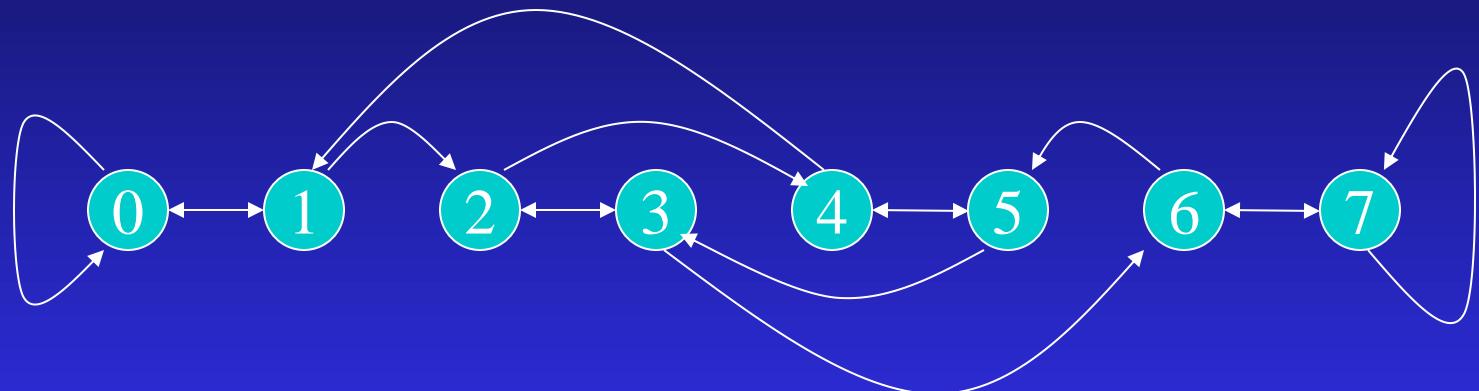
Evaluating Hypercube Network

- Diameter: $\log n$
- Bisection width: $n / 2$
- Edges per node: $\log n$
- Constant edge length? No

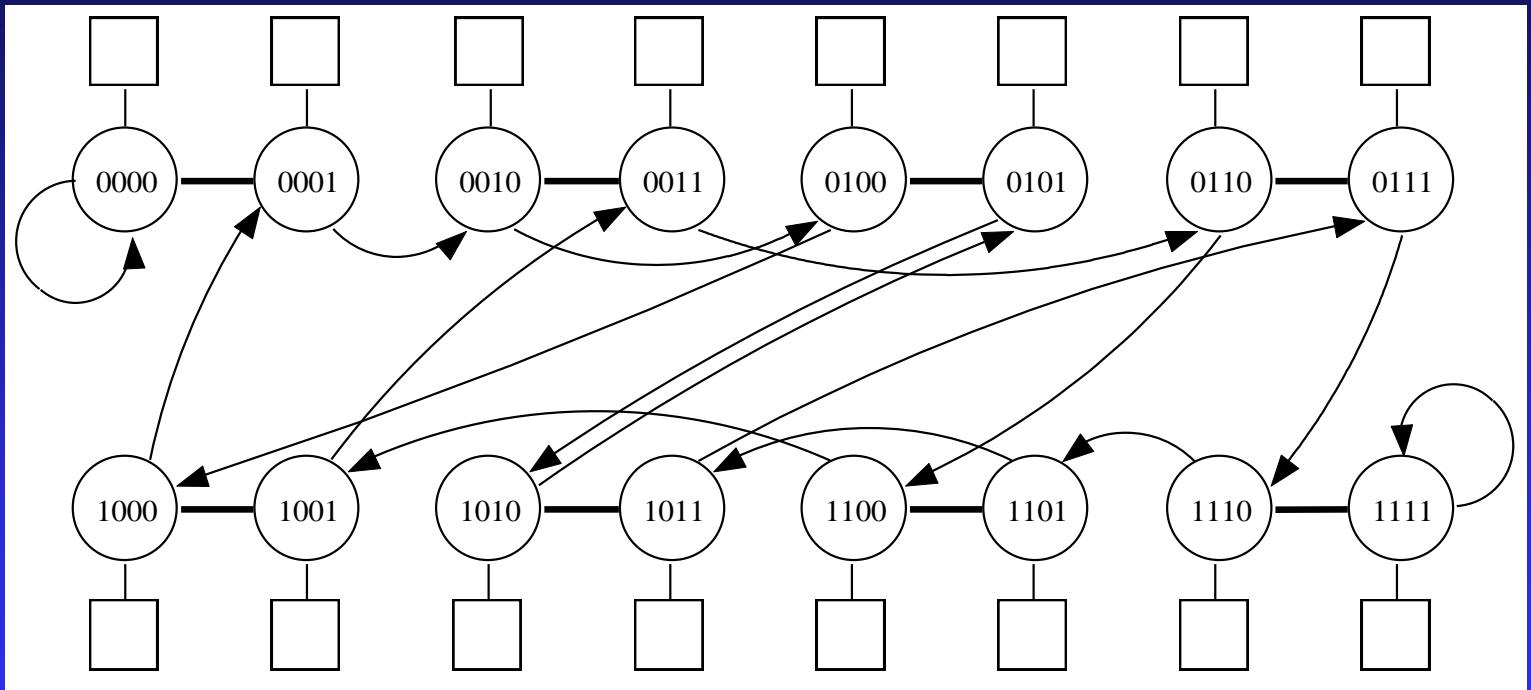
Shuffle-exchange

- Direct topology
- Number of nodes a power of 2
- Nodes have addresses $0, 1, \dots, 2^k - 1$
- Two outgoing links from node i
 - ◆ Shuffle link to node $\text{LeftCycle}(i)$
 - ◆ Exchange link to node $[\text{xor}(i, 1)]$

Shuffle-exchange Illustrated



Shuffle-exchange Addressing



Evaluating Shuffle-exchange

- Diameter: $2\log n - 1$
- Bisection width: $\approx n / \log n$
- Edges per node: 2
- Constant edge length? No

Comparing Networks

- All have logarithmic diameter except 2-D mesh
- Hypertree, butterfly, and hypercube have bisection width $n / 2$
- All have constant edges per node except hypercube
- Only 2-D mesh keeps edge lengths constant as network size increases

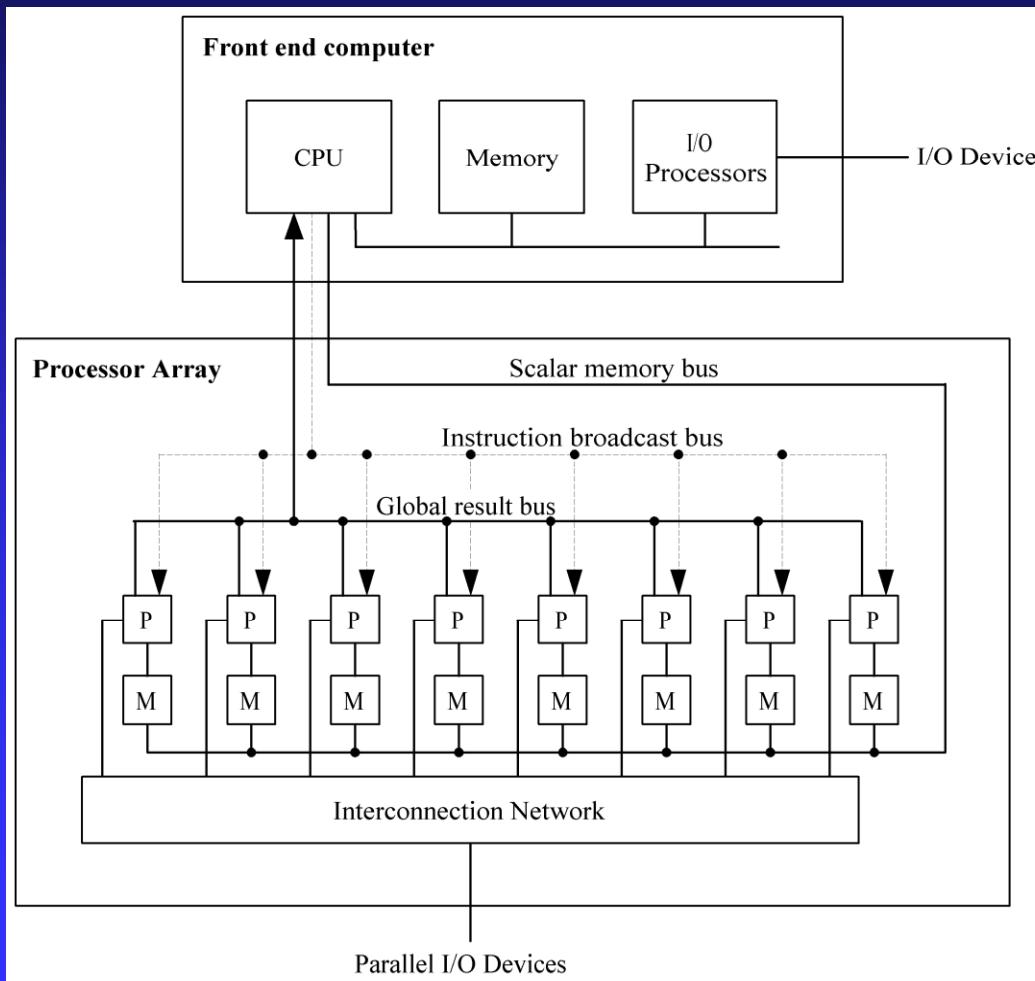
Vector Computers

- Vector computer: instruction set includes operations on vectors as well as scalars
- Two ways to implement vector computers
 - ◆ Pipelined vector processor: streams data through pipelined arithmetic units
 - ◆ Processor array: many identical, synchronized arithmetic processing elements

Why Processor Arrays?

- Historically, high cost of a control unit
- Scientific applications have data parallelism

Processor Array



Data/instruction Storage

- Front end computer
 - ◆ Program
 - ◆ Data manipulated sequentially
- Processor array
 - ◆ Data manipulated in parallel

Processor Array Performance

- Performance: work done per time unit
- Performance of processor array
 - ◆ Speed of processing elements
 - ◆ Utilization of processing elements

Performance Example 1

- 1024 processors
- Each adds a pair of integers in 1 μ sec
- What is performance when adding two 1024-element vectors (one per processor)?

$$\text{Performance} = \frac{1024 \text{operations}}{1\mu\text{sec}} = 1.024 \times 10^9 \text{ops/sec}$$

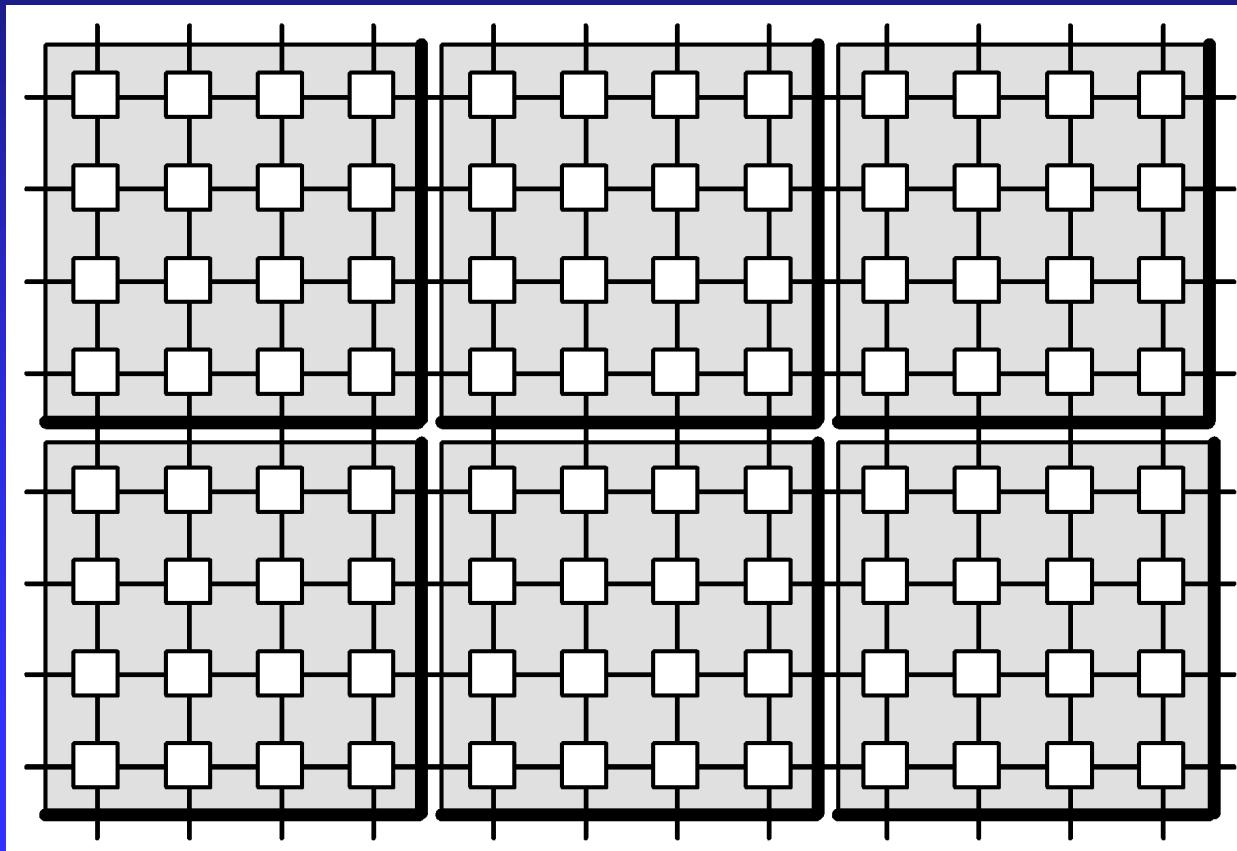
Performance Example 2

- 512 processors
- Each adds two integers in 1 μ sec
- Performance adding two vectors of length 600?

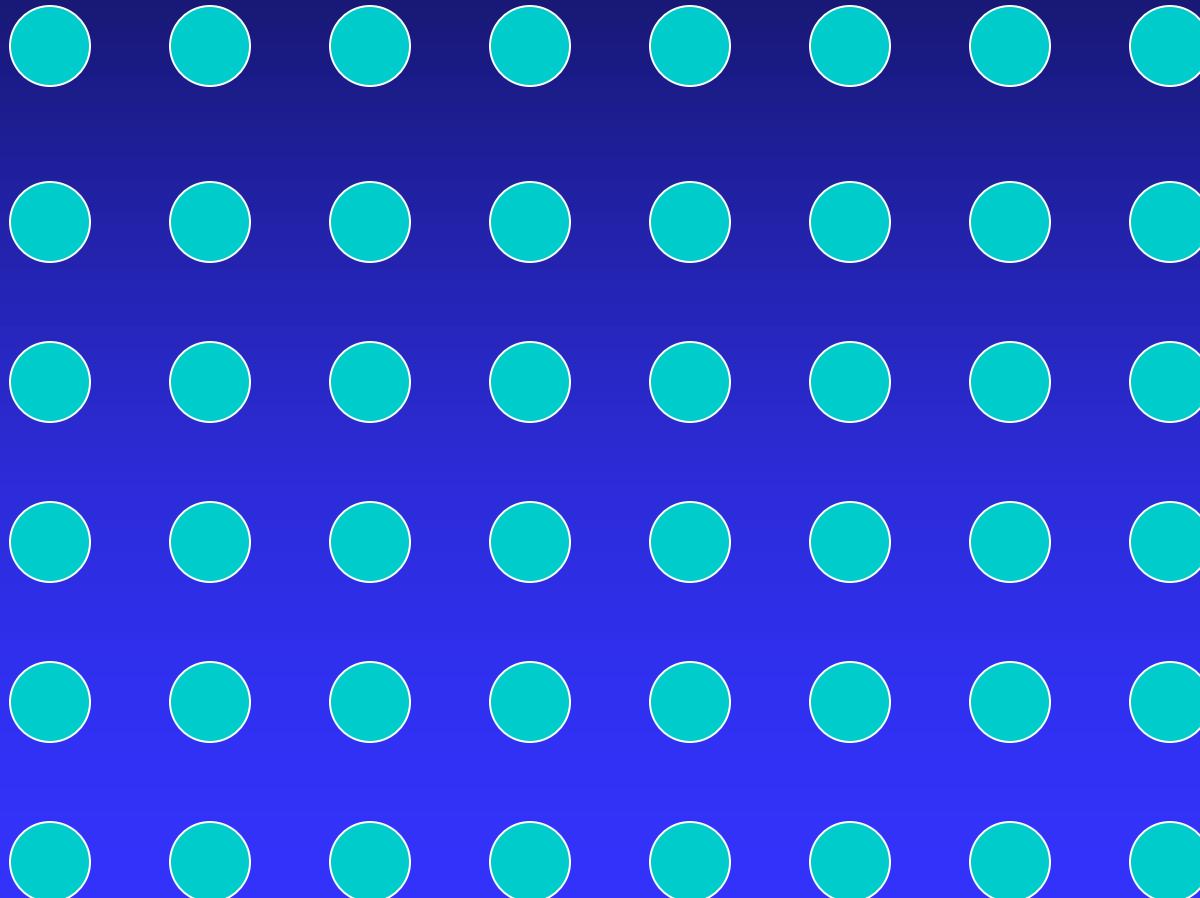
$$\text{Performance} = \frac{600 \text{operations}}{2 \mu\text{sec}} = 3 \times 10^6 \text{ops/sec}$$

2-D Processor Interconnection Network

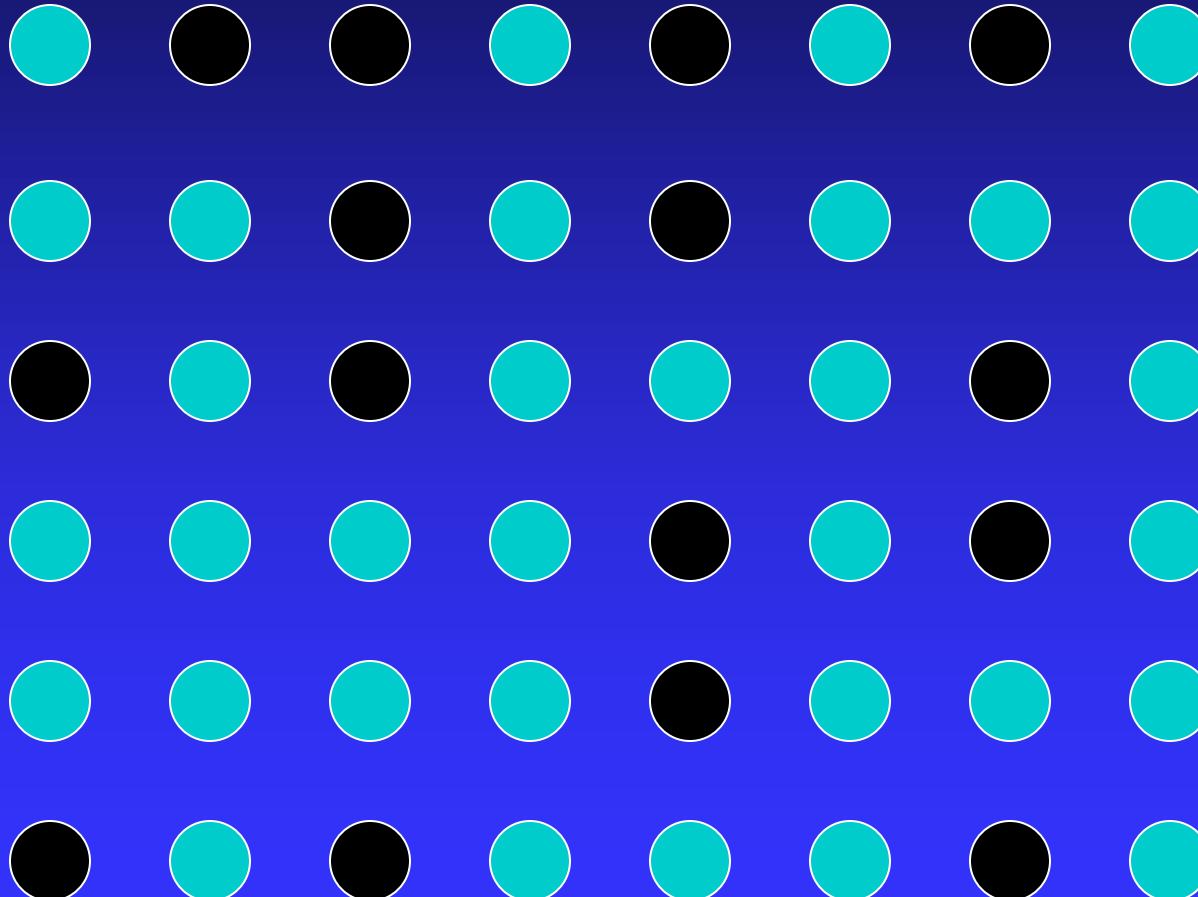
Each VLSI chip has 16 processing elements



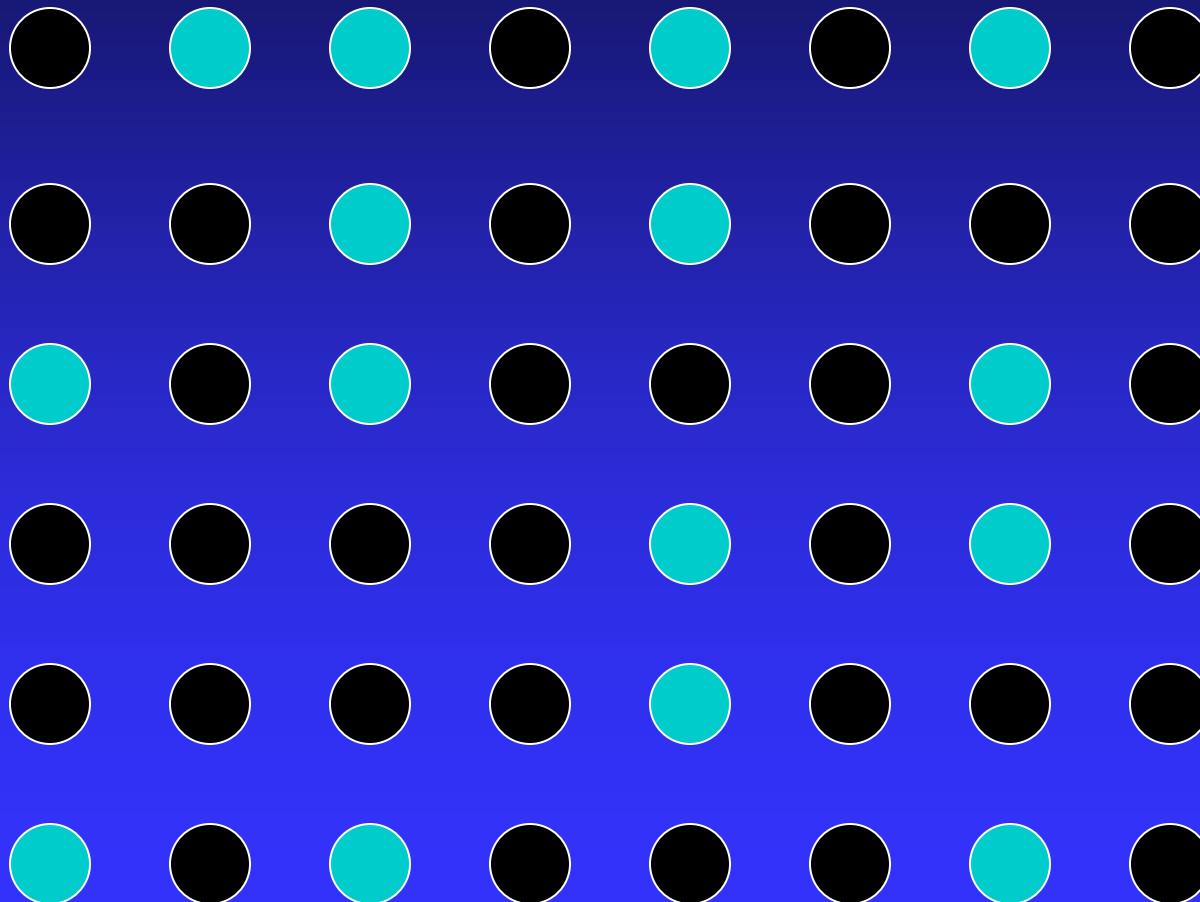
if (COND) then A else B



if (COND) then A else B



if (COND) then A else B



Processor Array Shortcomings

- Not all problems are data-parallel
- Speed drops for conditionally executed code
- Don't adapt to multiple users well
- Do not scale down well to “starter” systems
- Rely on custom VLSI for processors
- Expense of control units has dropped

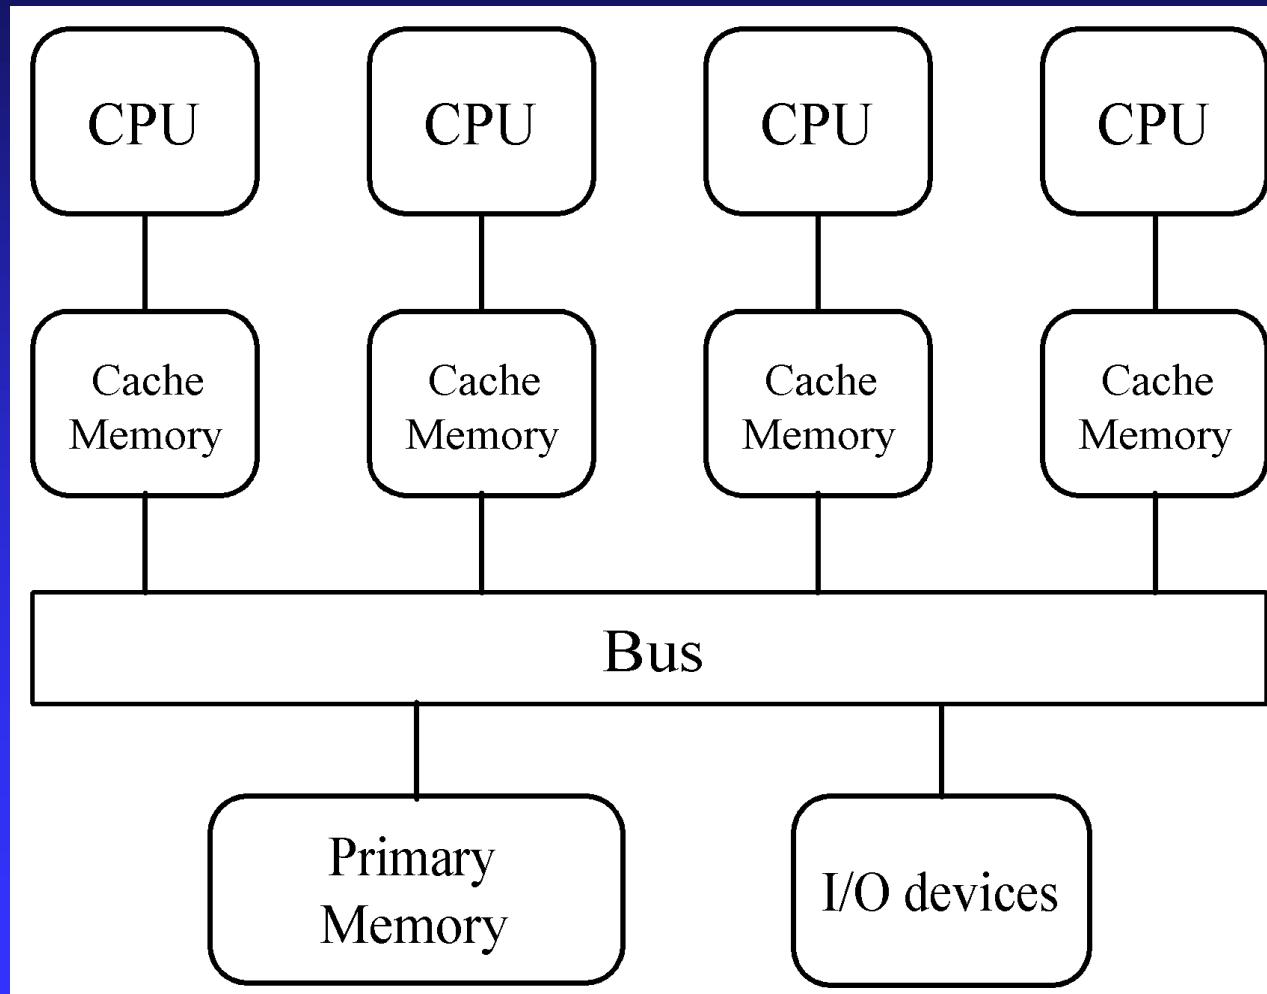
Multiprocessors

- Multiprocessor: multiple-CPU computer with a shared memory
- Same address on two different CPUs refers to the same memory location
- Avoid three problems of processor arrays
 - ◆ Can be built from commodity CPUs
 - ◆ Naturally support multiple users
 - ◆ Maintain efficiency in conditional code

Centralized Multiprocessor

- Straightforward extension of uniprocessor
- Add CPUs to bus
- All processors share same primary memory
- Memory access time same for all CPUs
 - ◆ Uniform memory access (UMA) multiprocessor
 - ◆ Symmetrical multiprocessor (SMP)

Centralized Multiprocessor



Private and Shared Data

- Private data: items used only by a single processor
- Shared data: values used by multiple processors
- In a multiprocessor, processors communicate via shared data values

Problems Associated with Shared Data

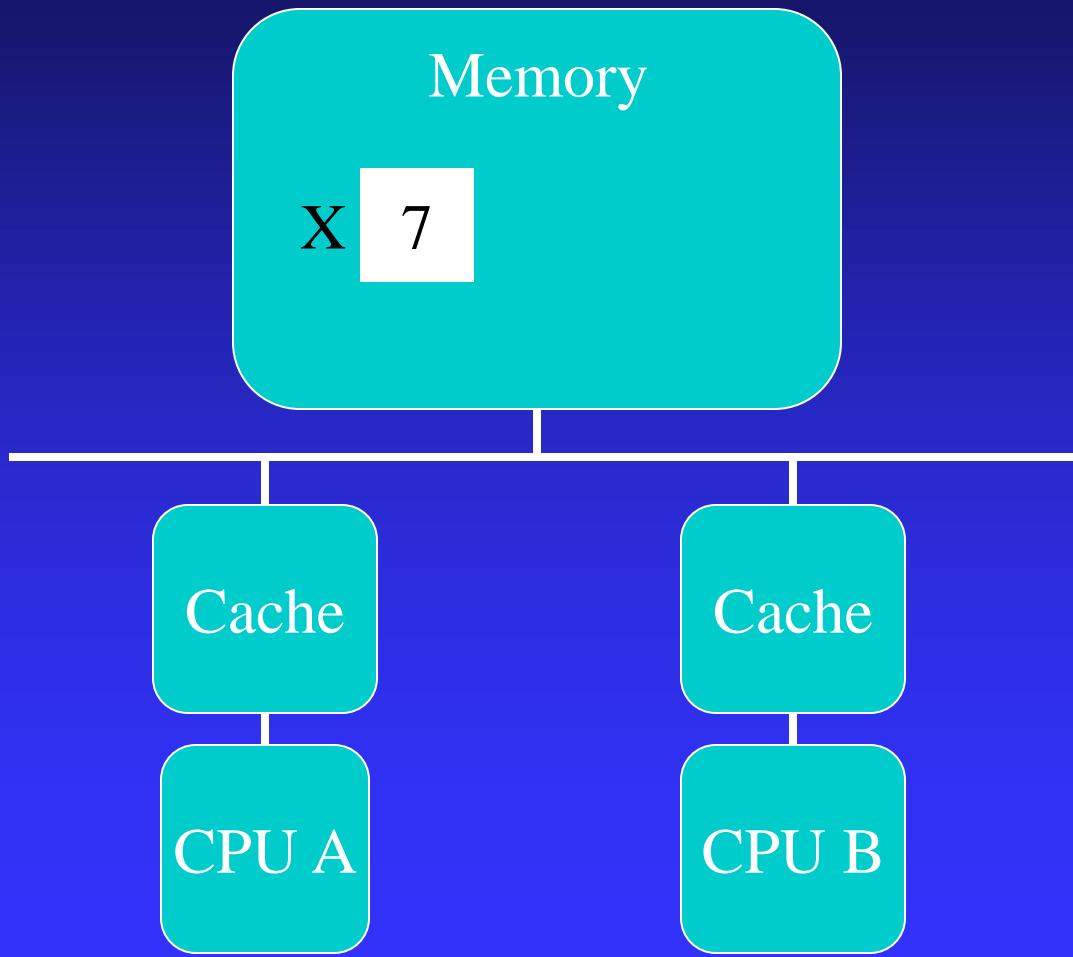
■ Cache coherence

- ◆ Replicating data across multiple caches reduces contention
- ◆ How to ensure different processors have same value for same address?

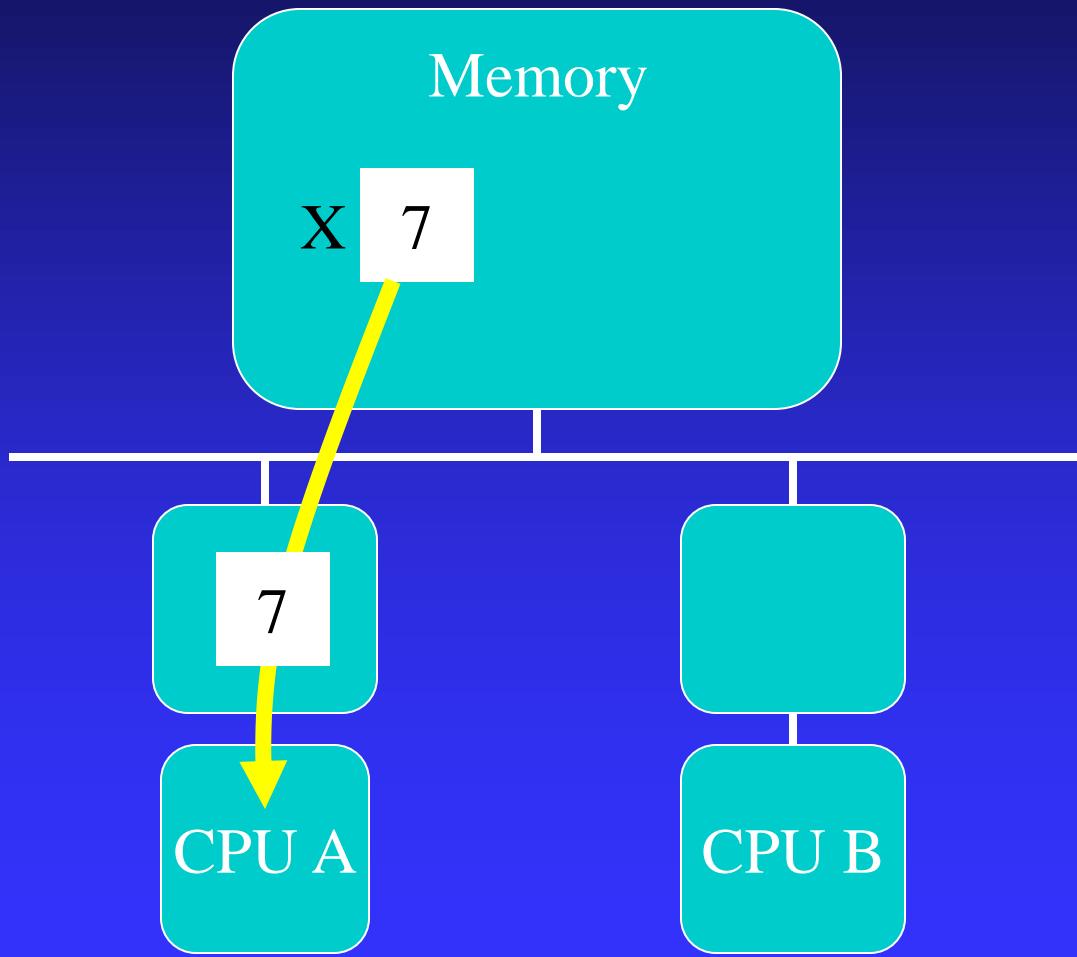
■ Synchronization

- ◆ Mutual exclusion
- ◆ Barrier

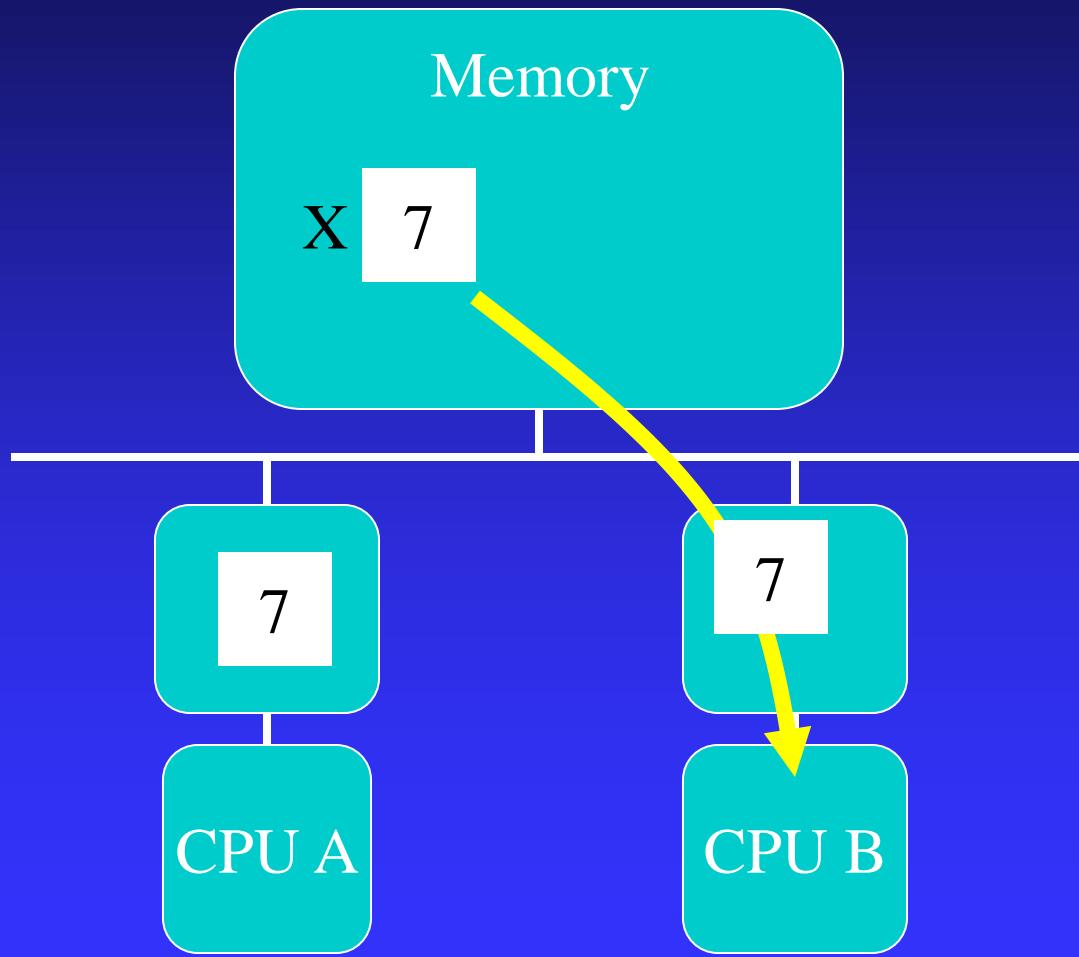
Cache-coherence Problem



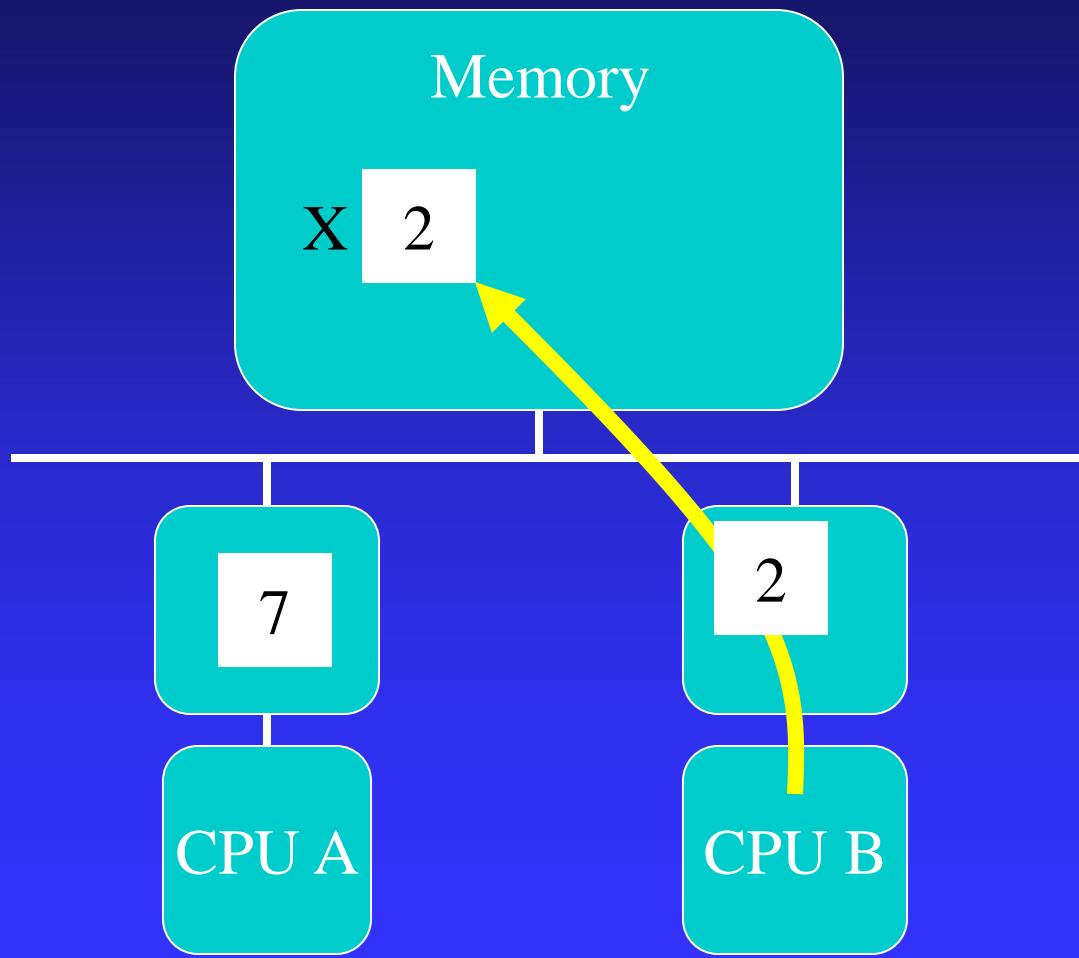
Cache-coherence Problem



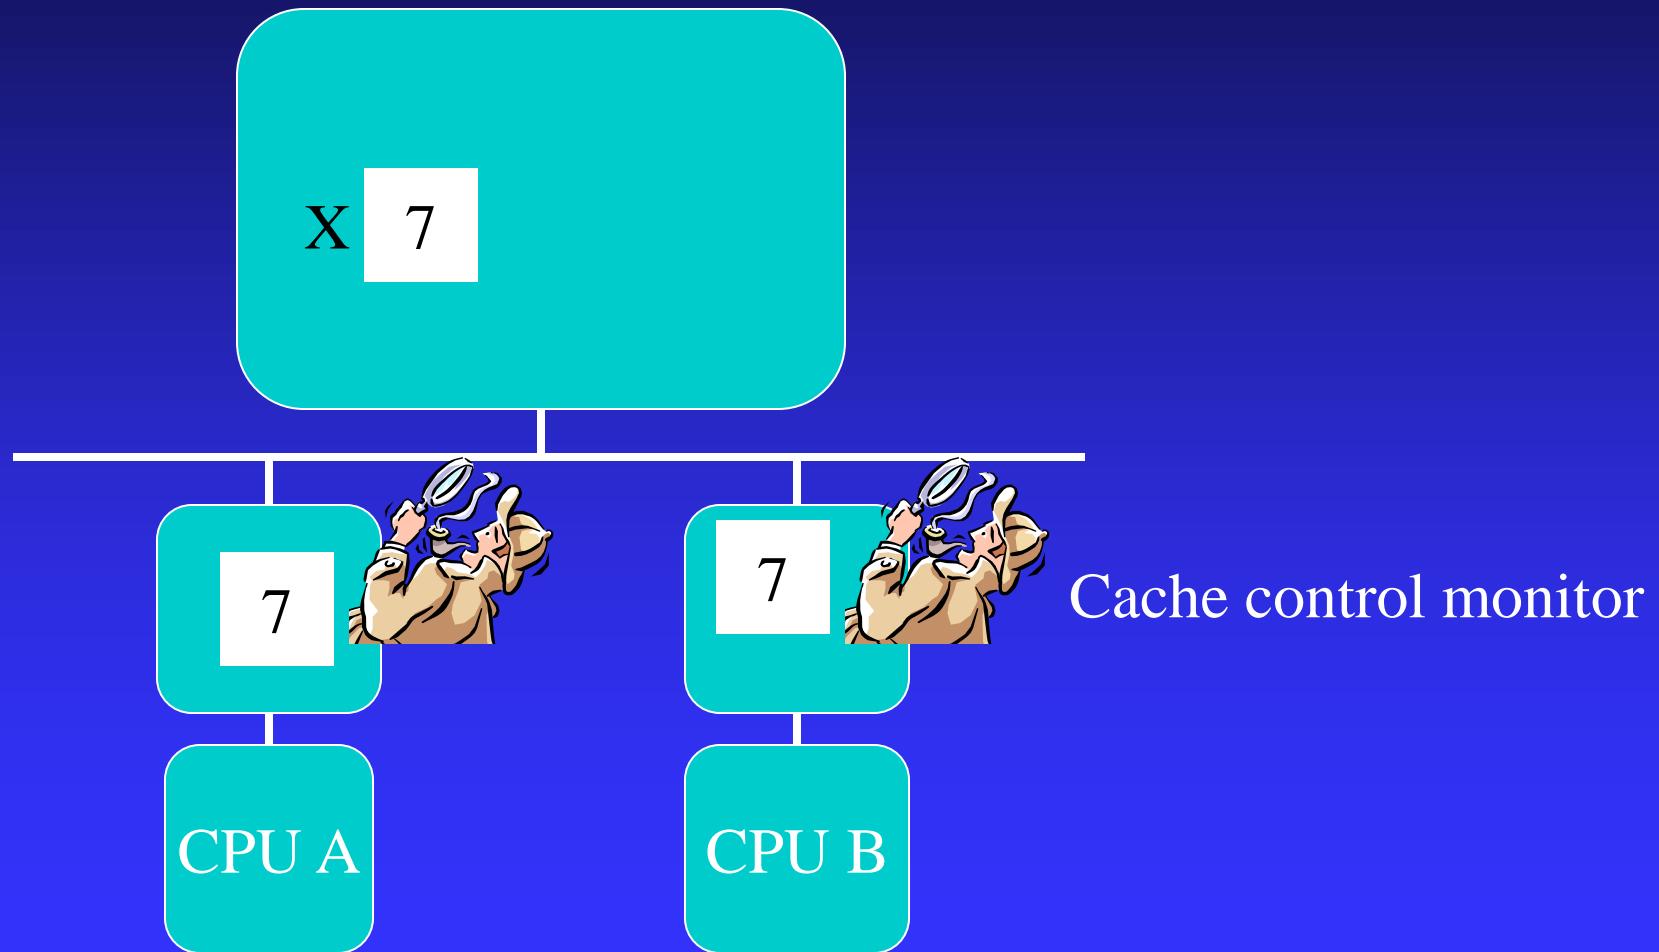
Cache-coherence Problem



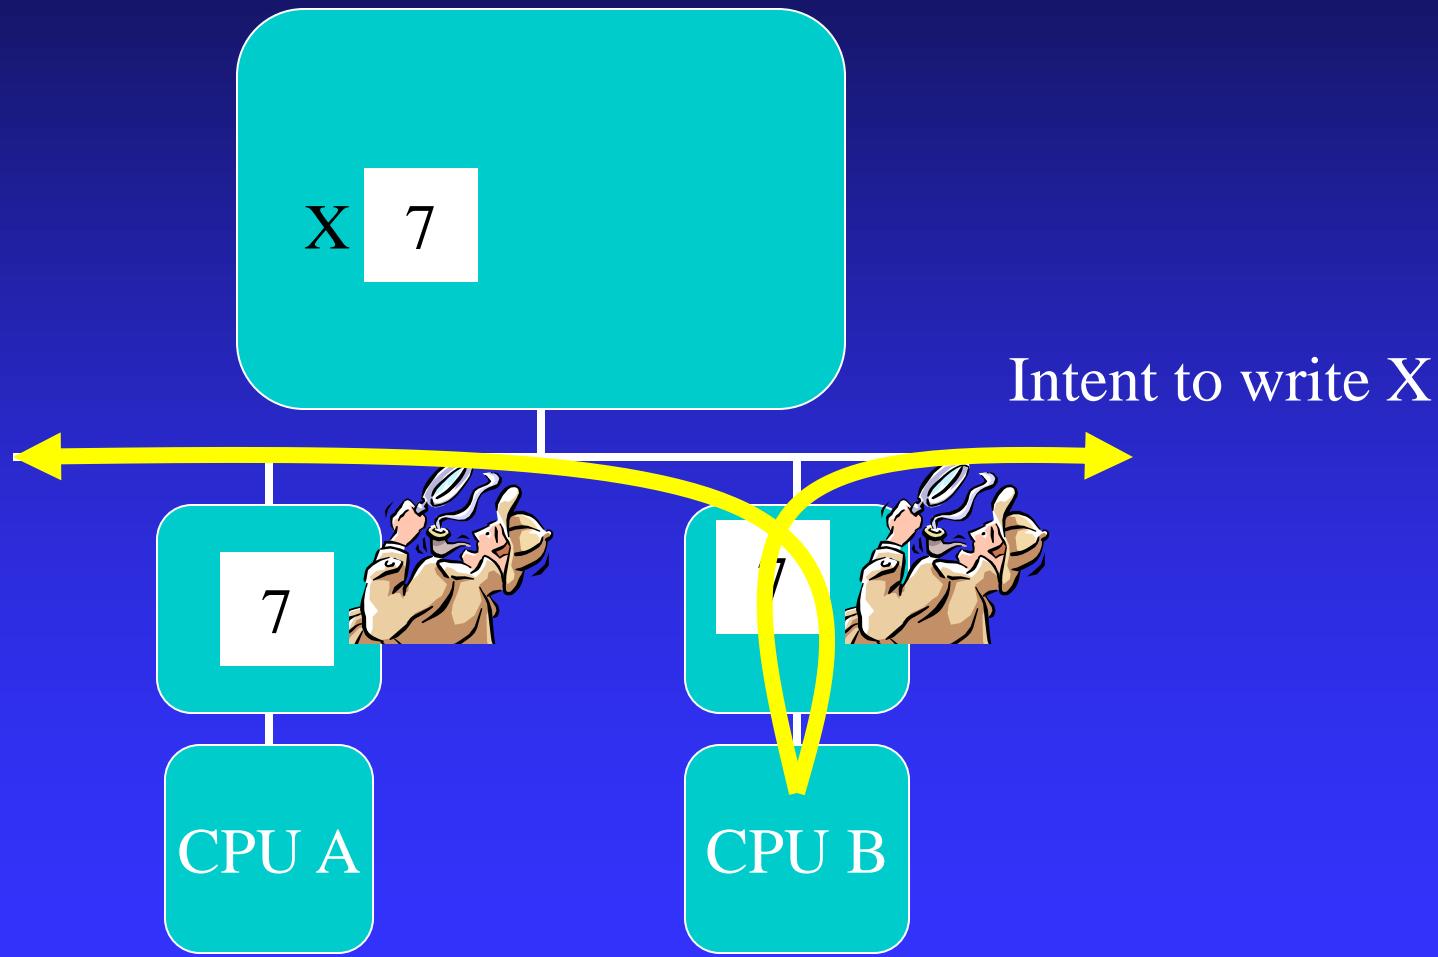
Cache-coherence Problem



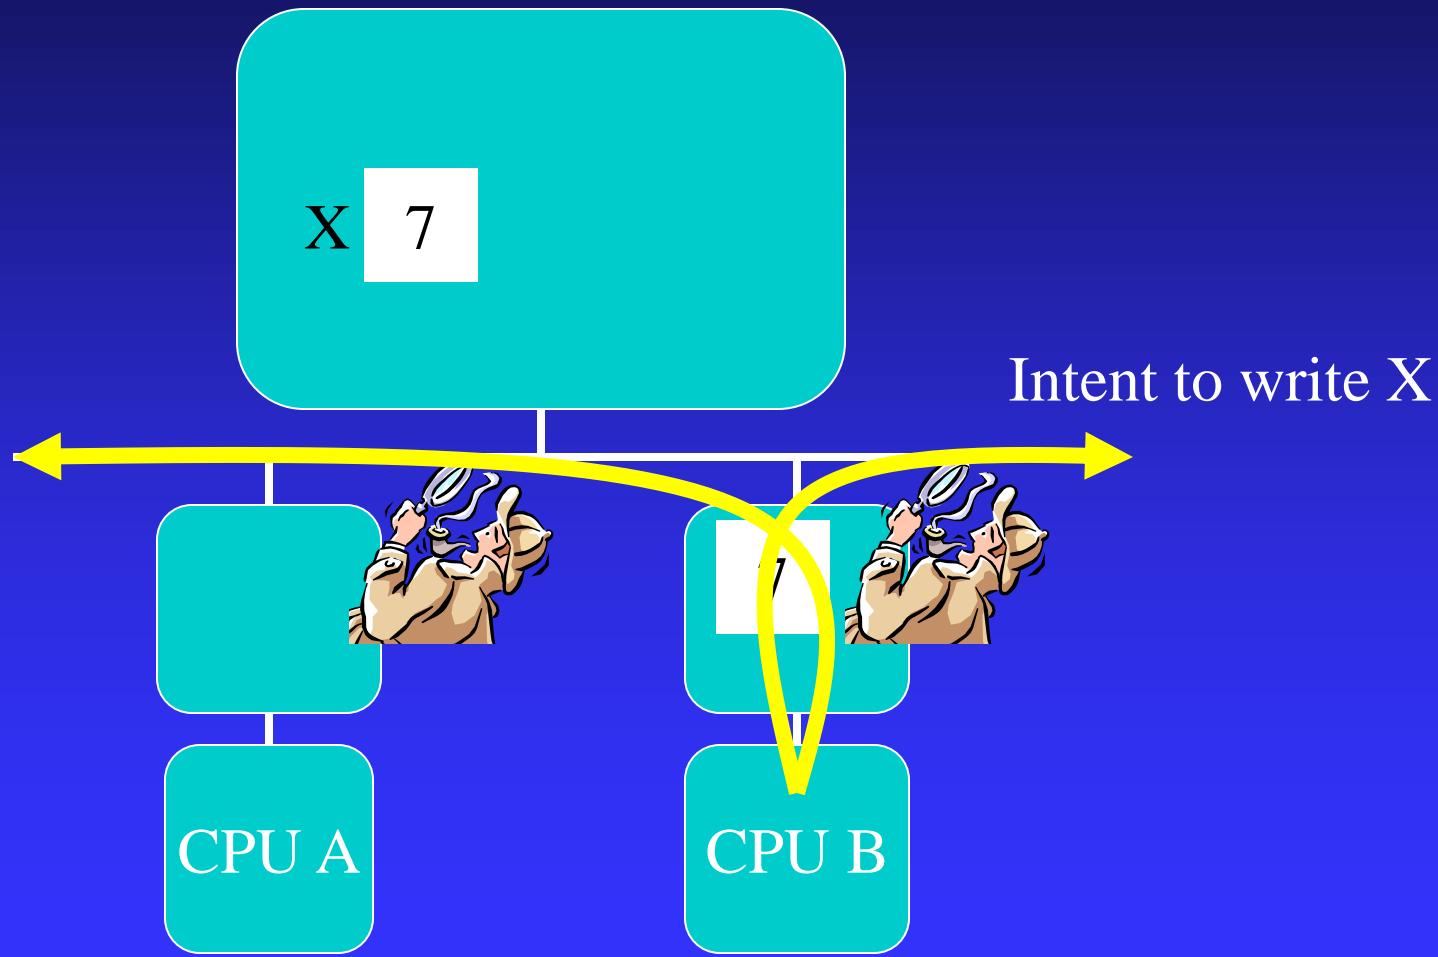
Write Invalidate Protocol



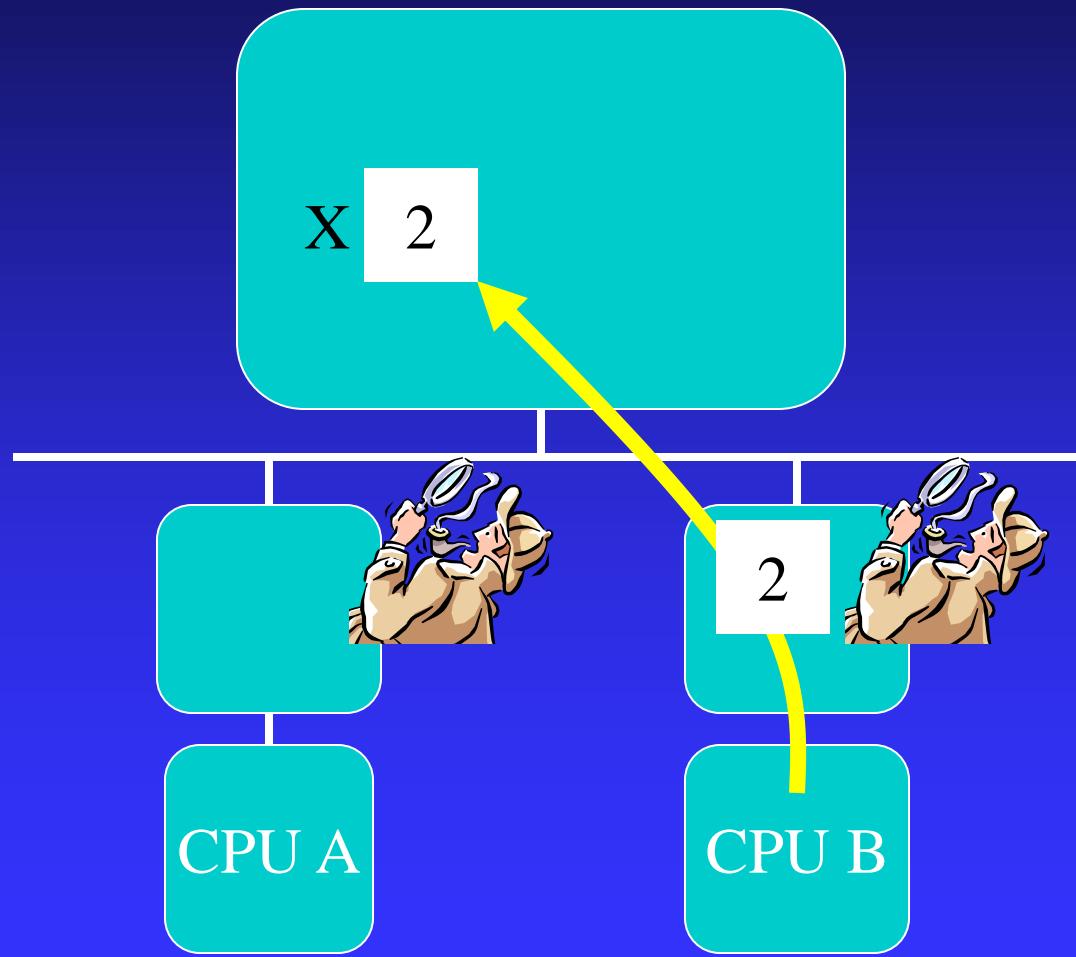
Write Invalidate Protocol



Write Invalidate Protocol



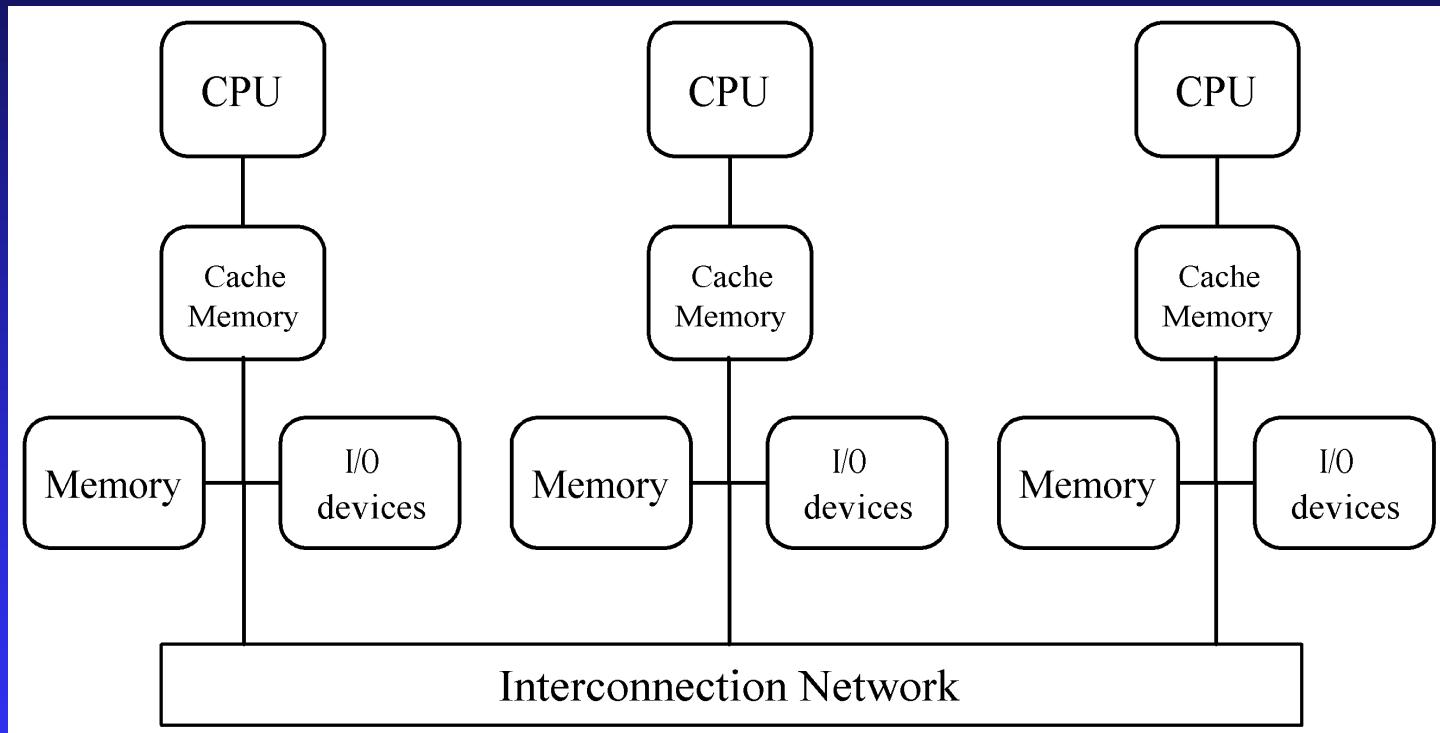
Write Invalidate Protocol



Distributed Multiprocessor

- Distribute primary memory among processors
- Increase aggregate memory bandwidth and lower average memory access time
- Allow greater number of processors
- Also called non-uniform memory access (NUMA) multiprocessor

Distributed Multiprocessor



Cache Coherence

- Some NUMA multiprocessors do not support it in hardware
 - ◆ Only instructions, private data in cache
 - ◆ Large memory access time variance
- Implementation more difficult
 - ◆ No shared memory bus to “snoop”
 - ◆ Directory-based protocol needed

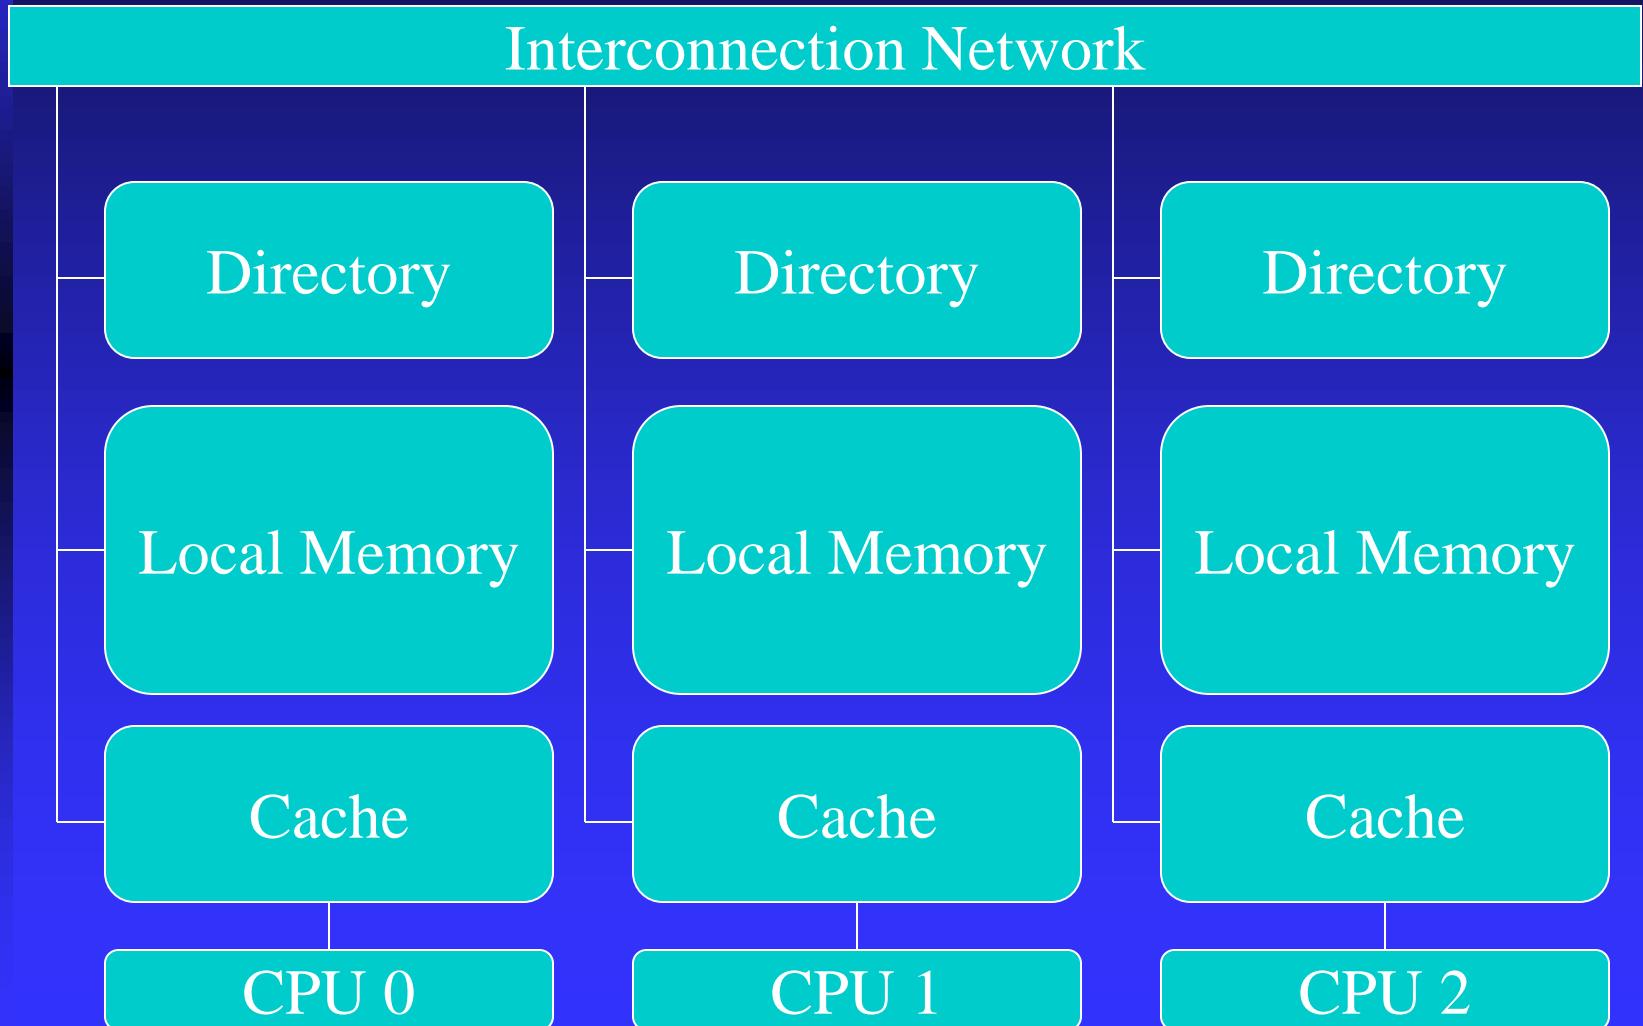
Directory-based Protocol

- Distributed directory contains information about cacheable memory blocks
- One directory entry for each cache block
- Each entry has
 - ◆ Sharing status
 - ◆ Which processors have copies

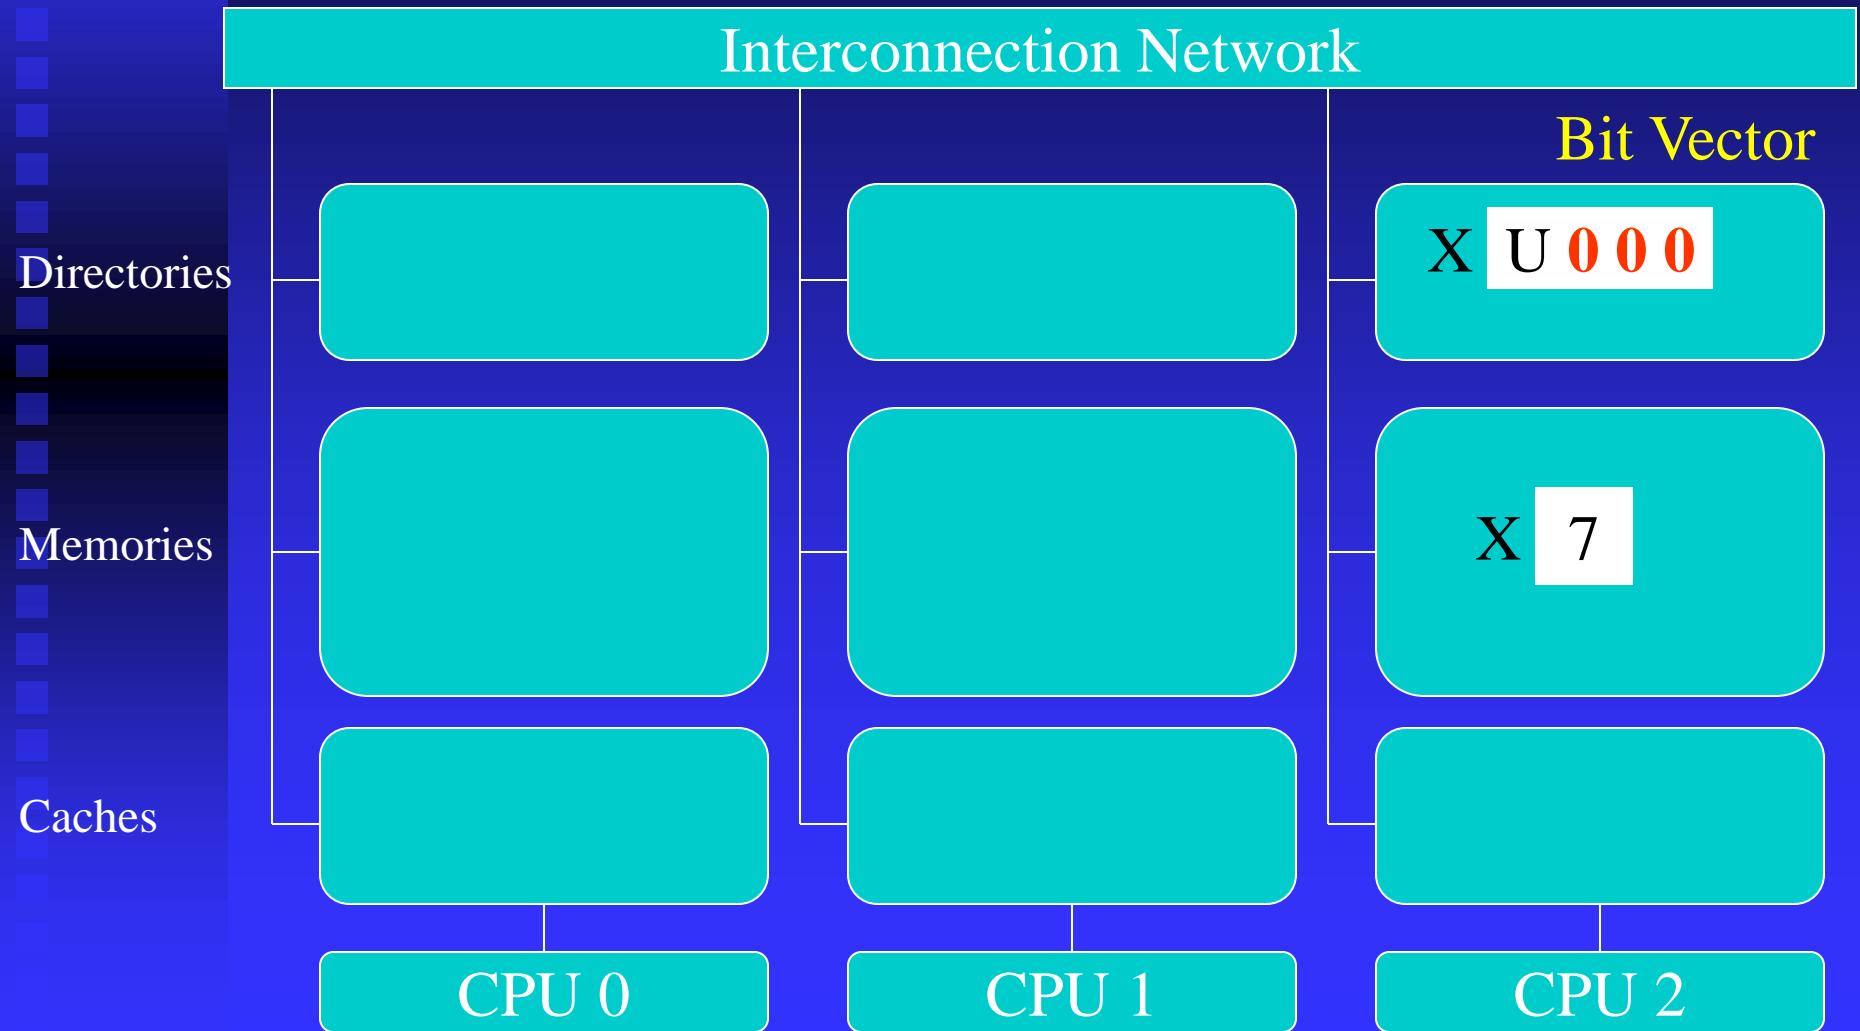
Sharing Status

- Uncached
 - ◆ Block not in any processor's cache
- Shared
 - ◆ Cached by one or more processors
 - ◆ Read only
- Exclusive
 - ◆ Cached by exactly one processor
 - ◆ Processor has written block
 - ◆ Copy in memory is obsolete

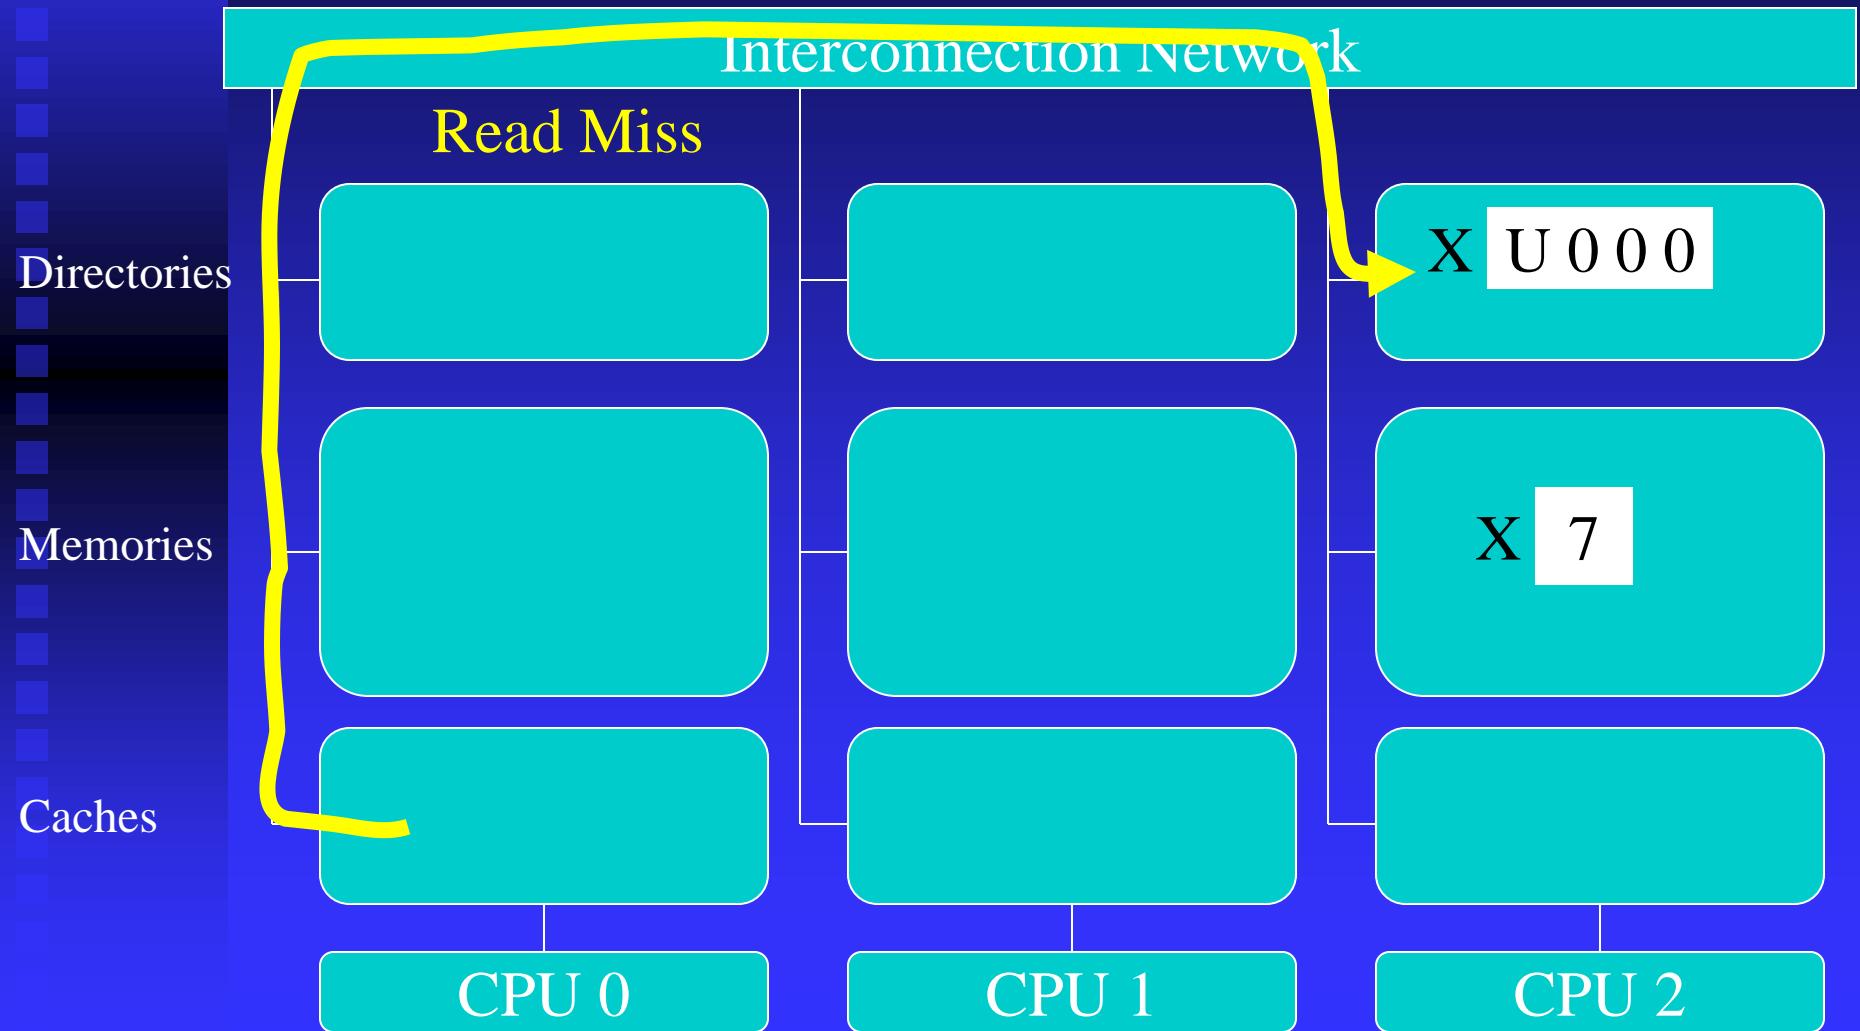
Directory-based Protocol



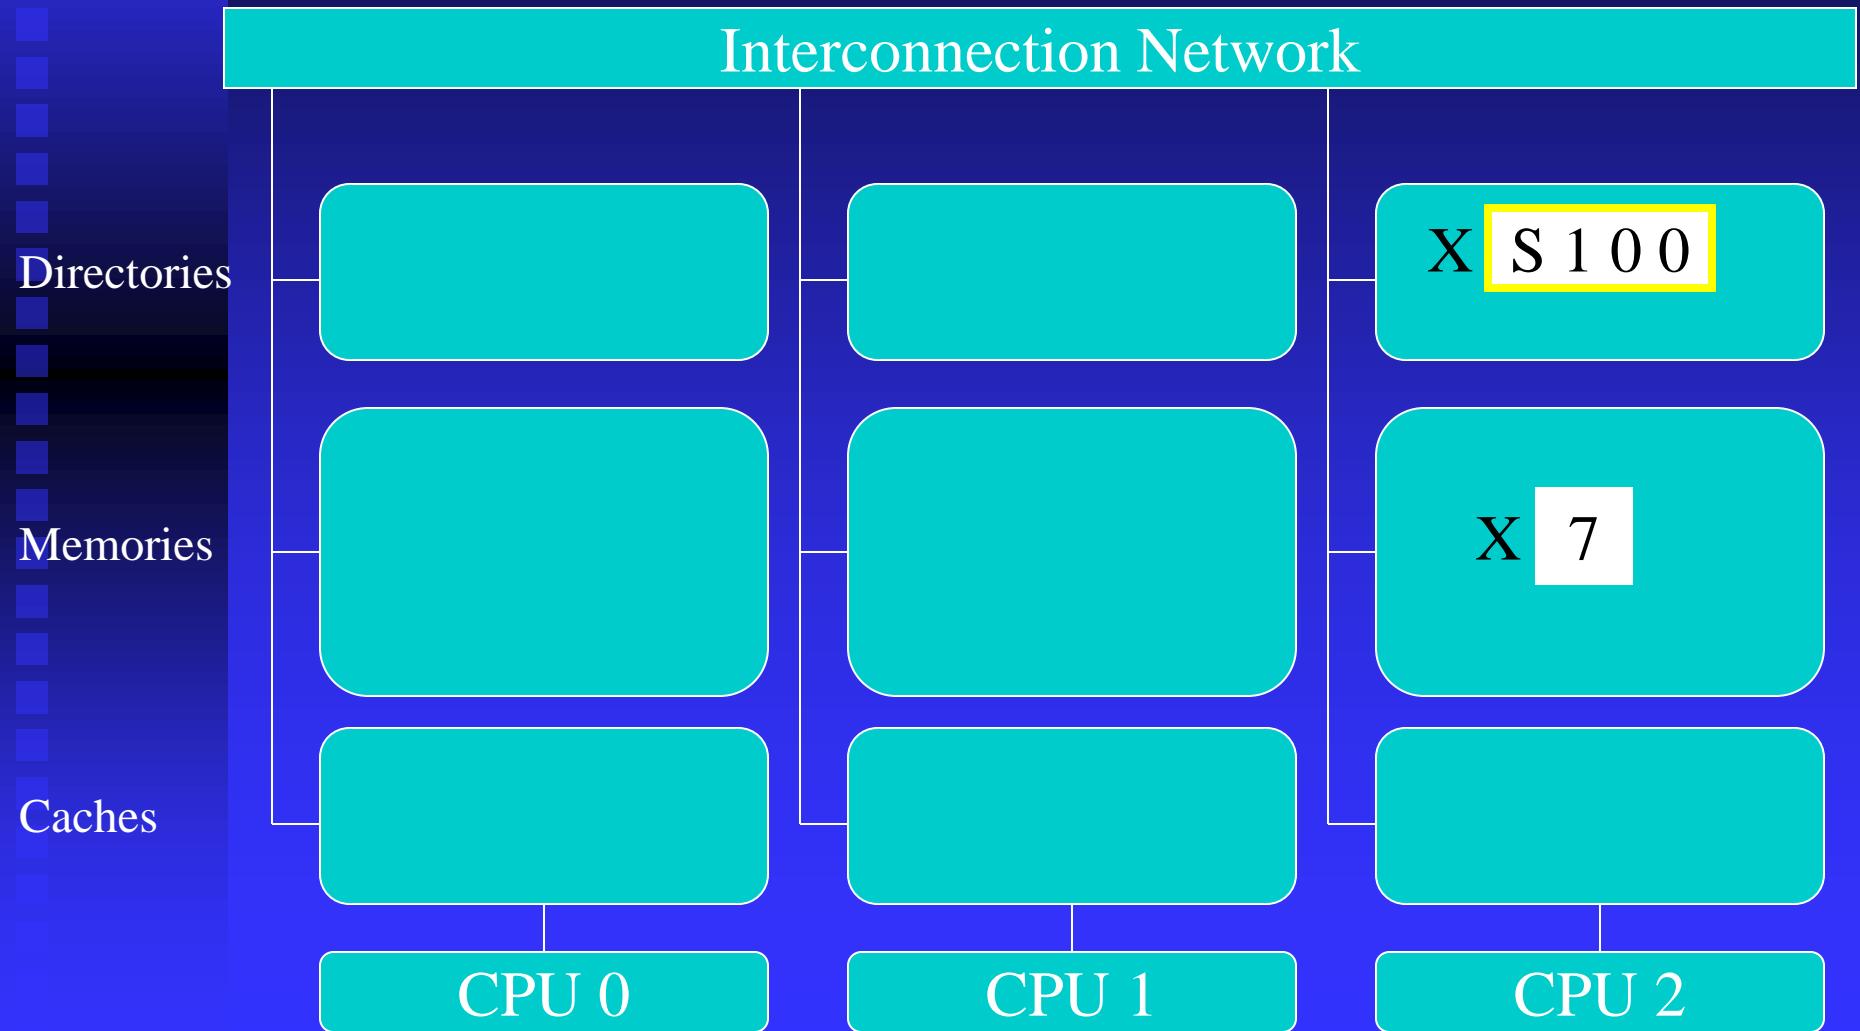
Directory-based Protocol



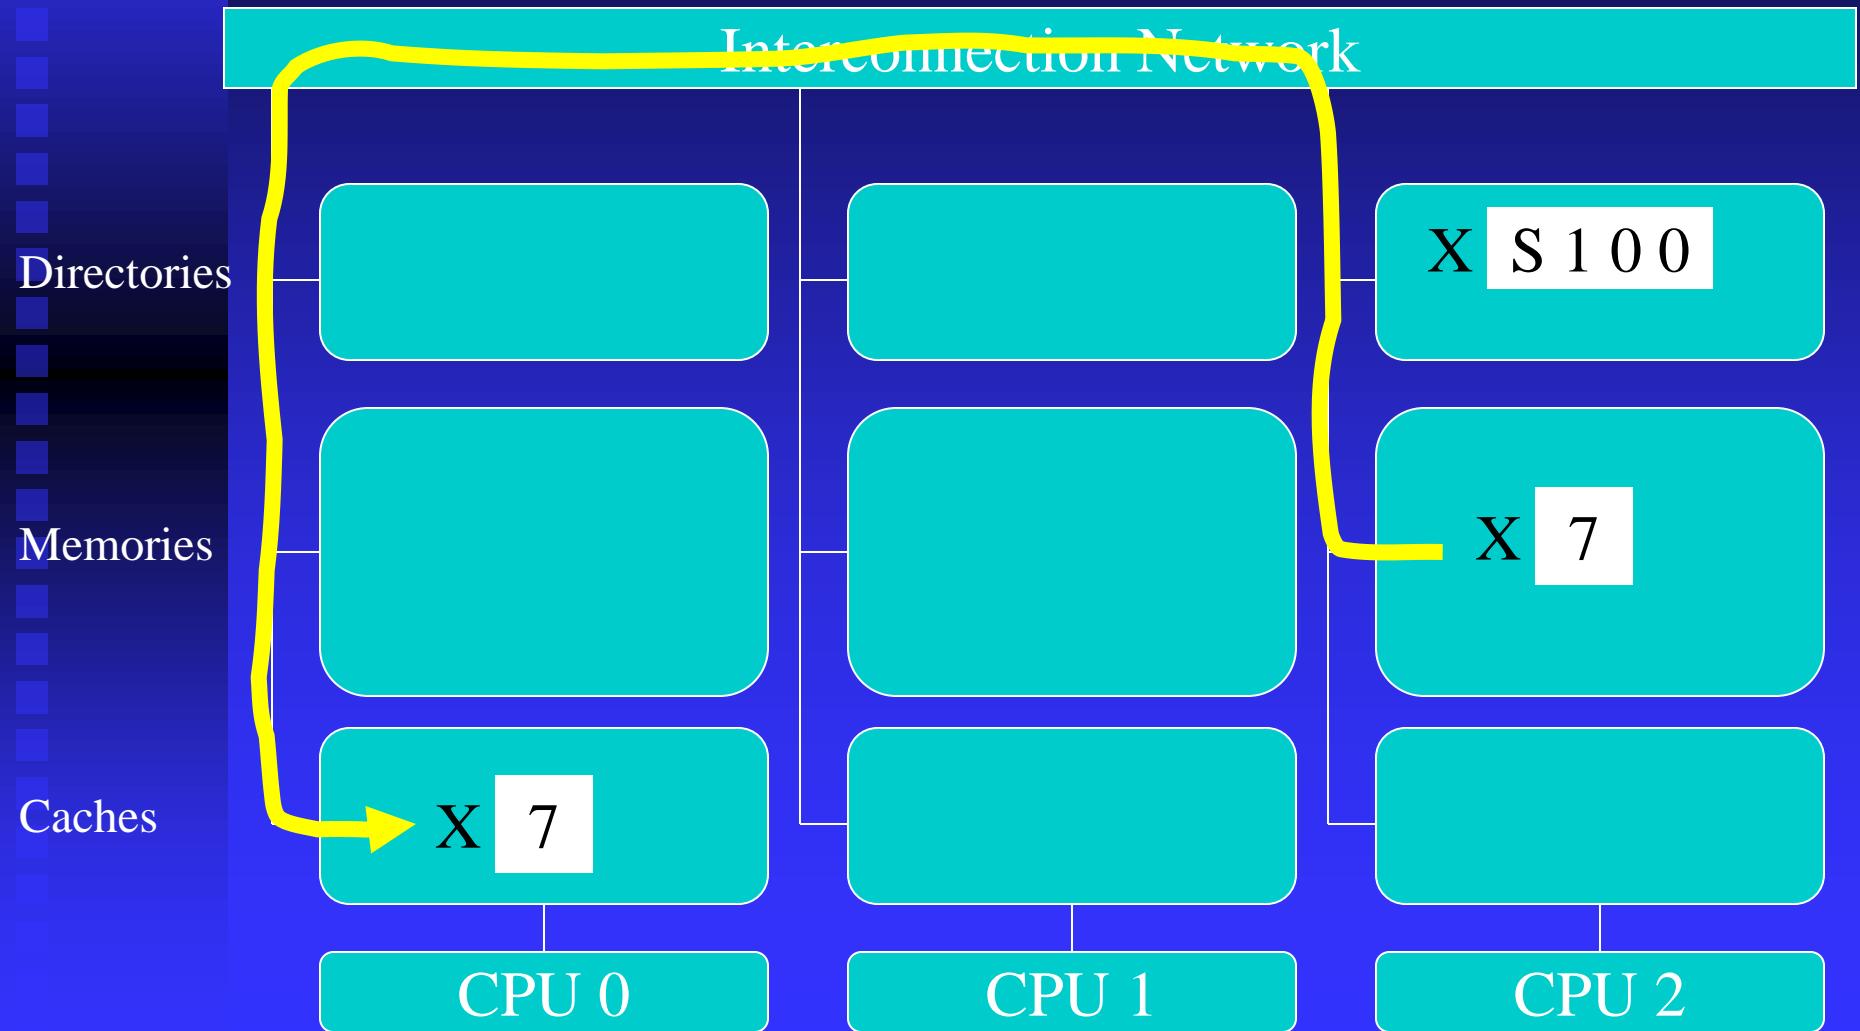
CPU 0 Reads X



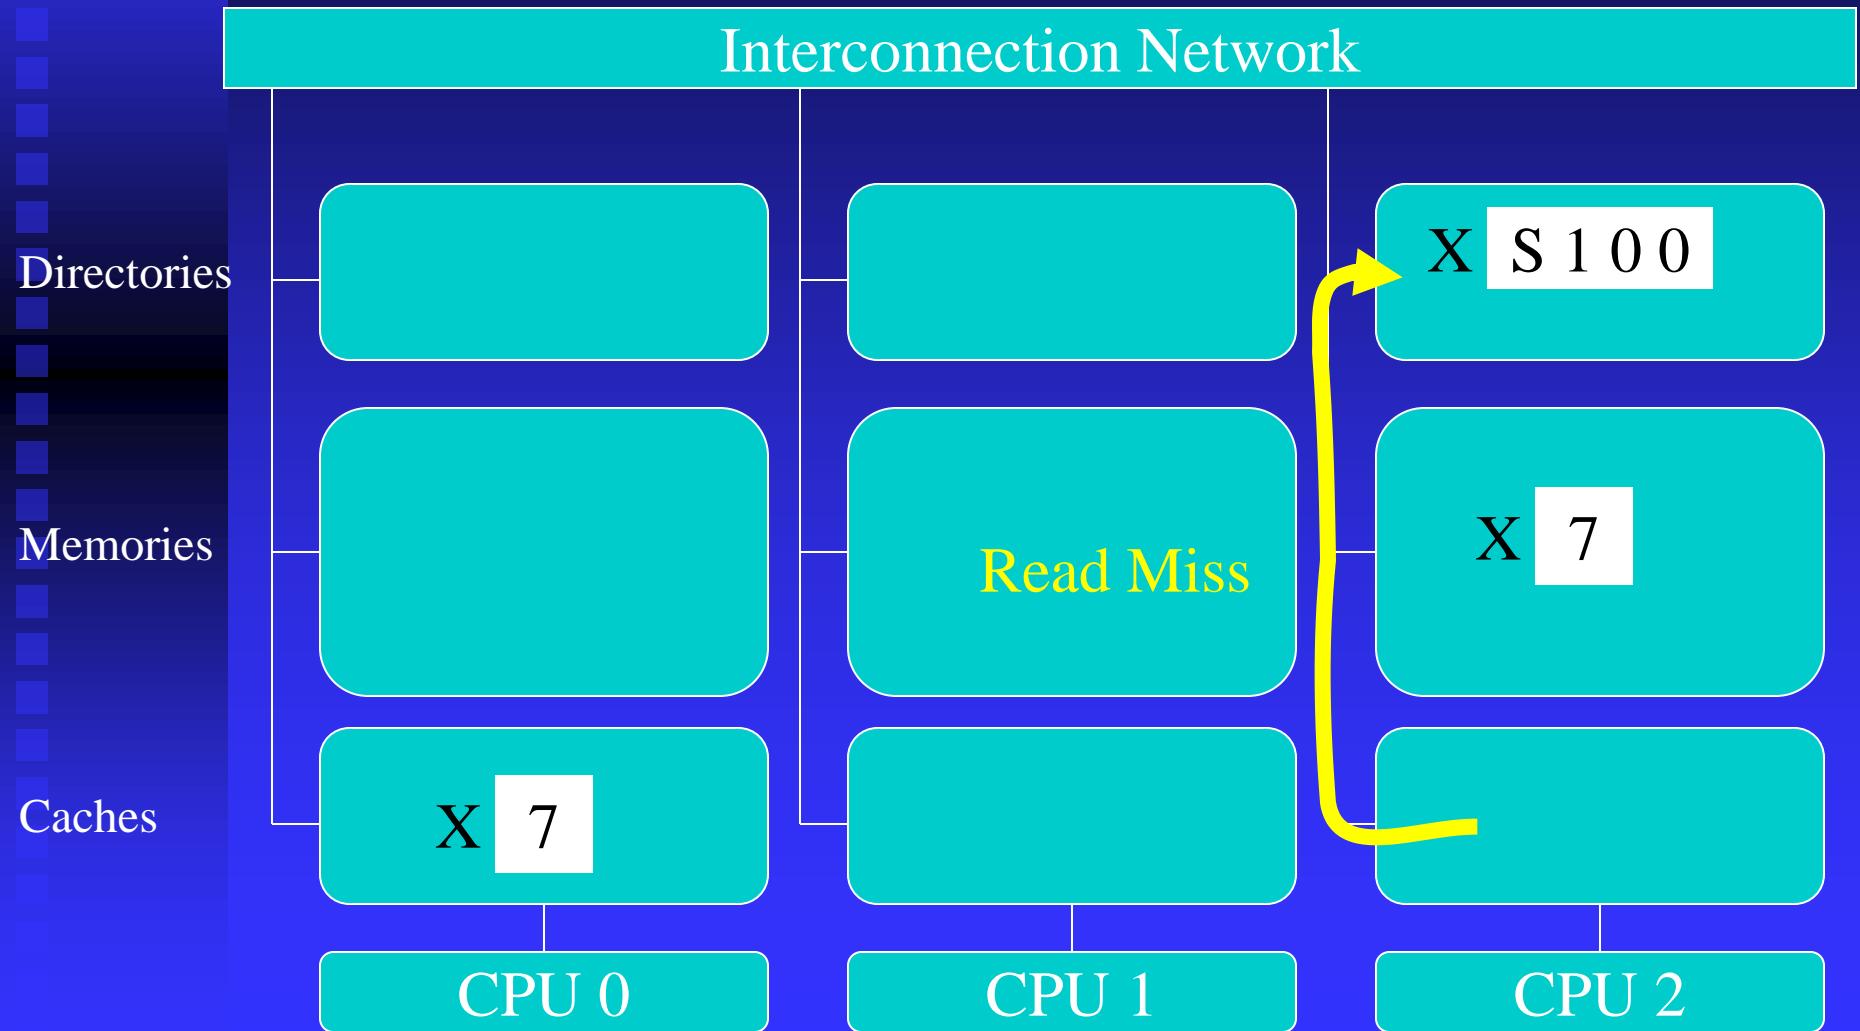
CPU 0 Reads X



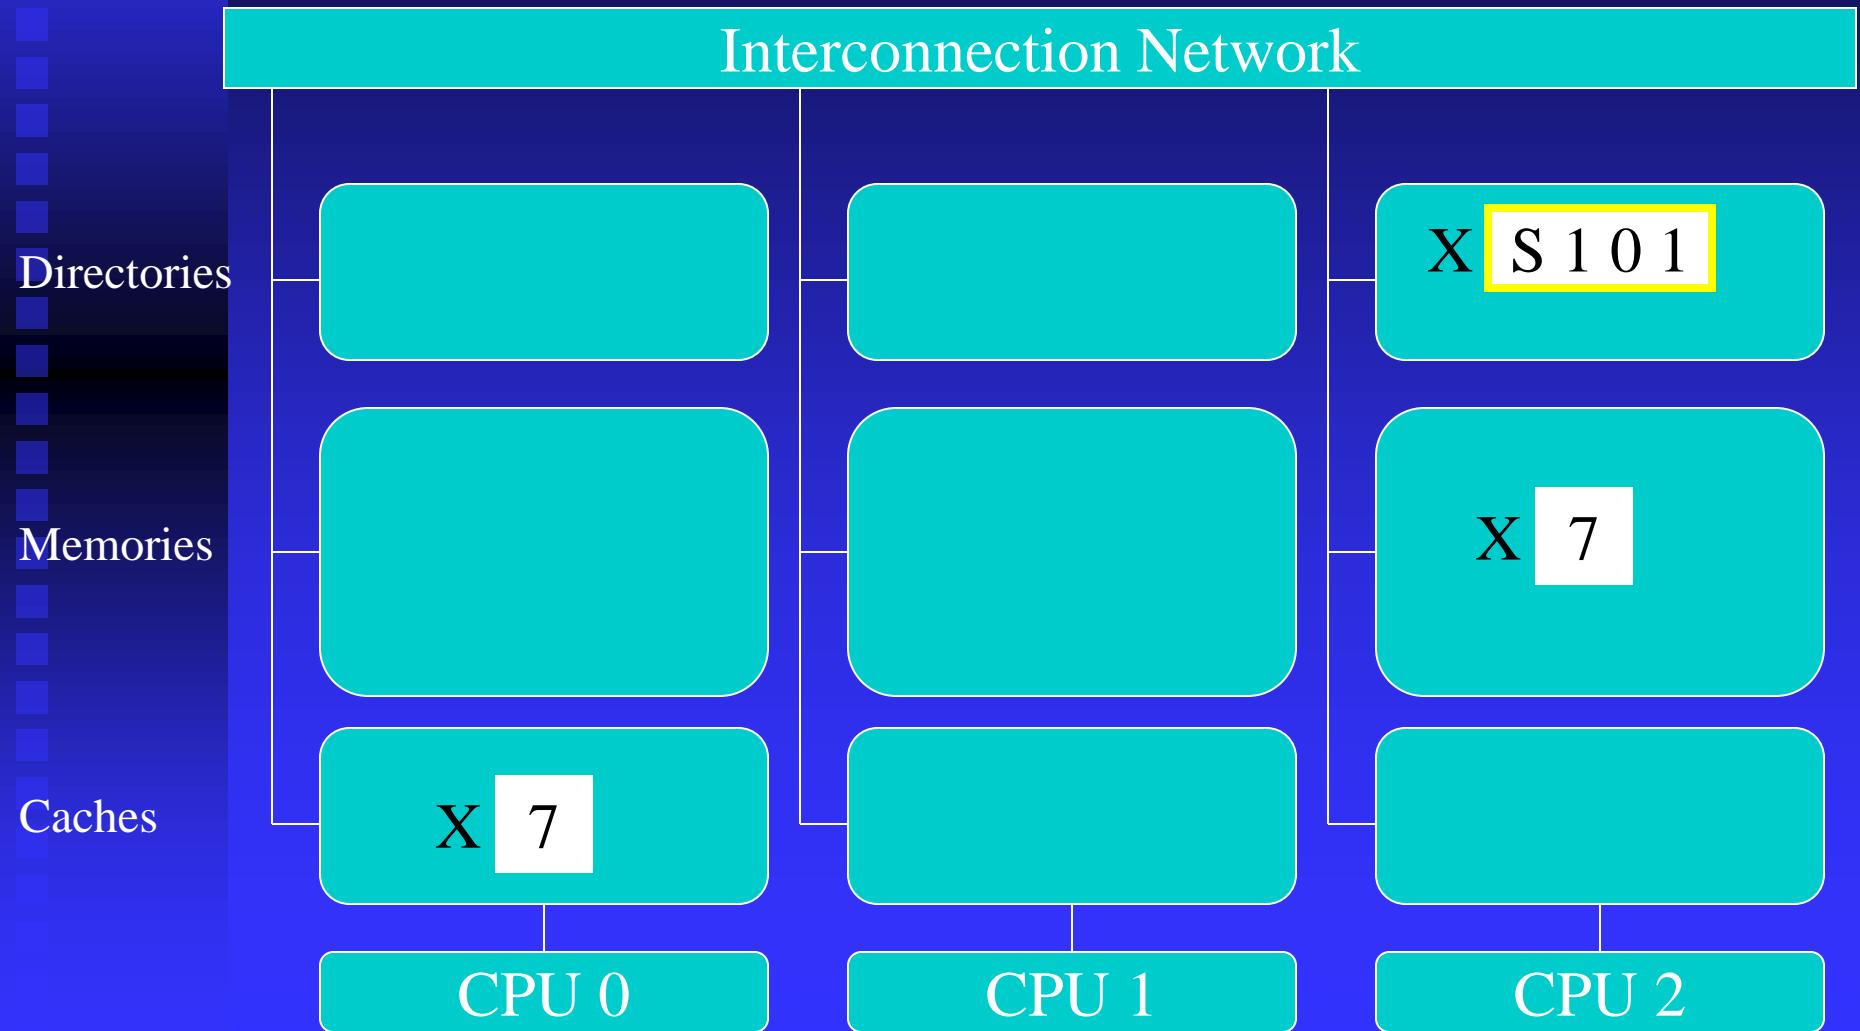
CPU 0 Reads X



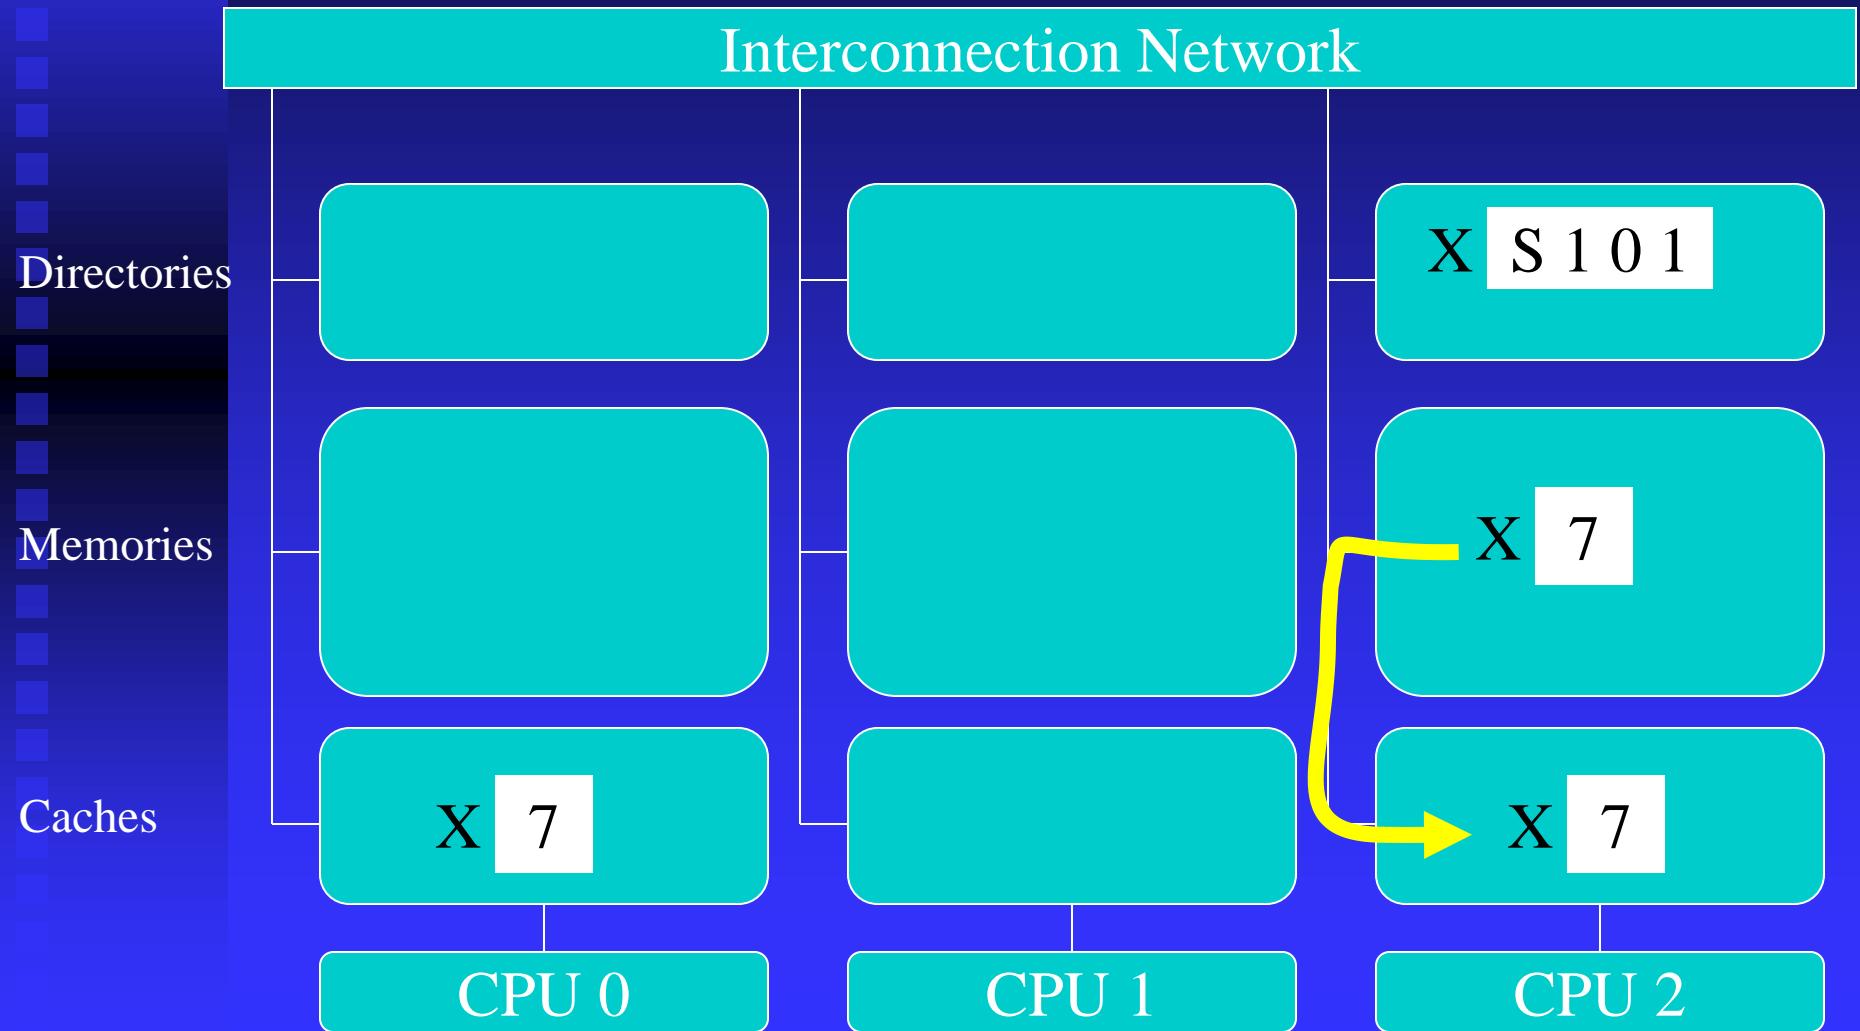
CPU 2 Reads X



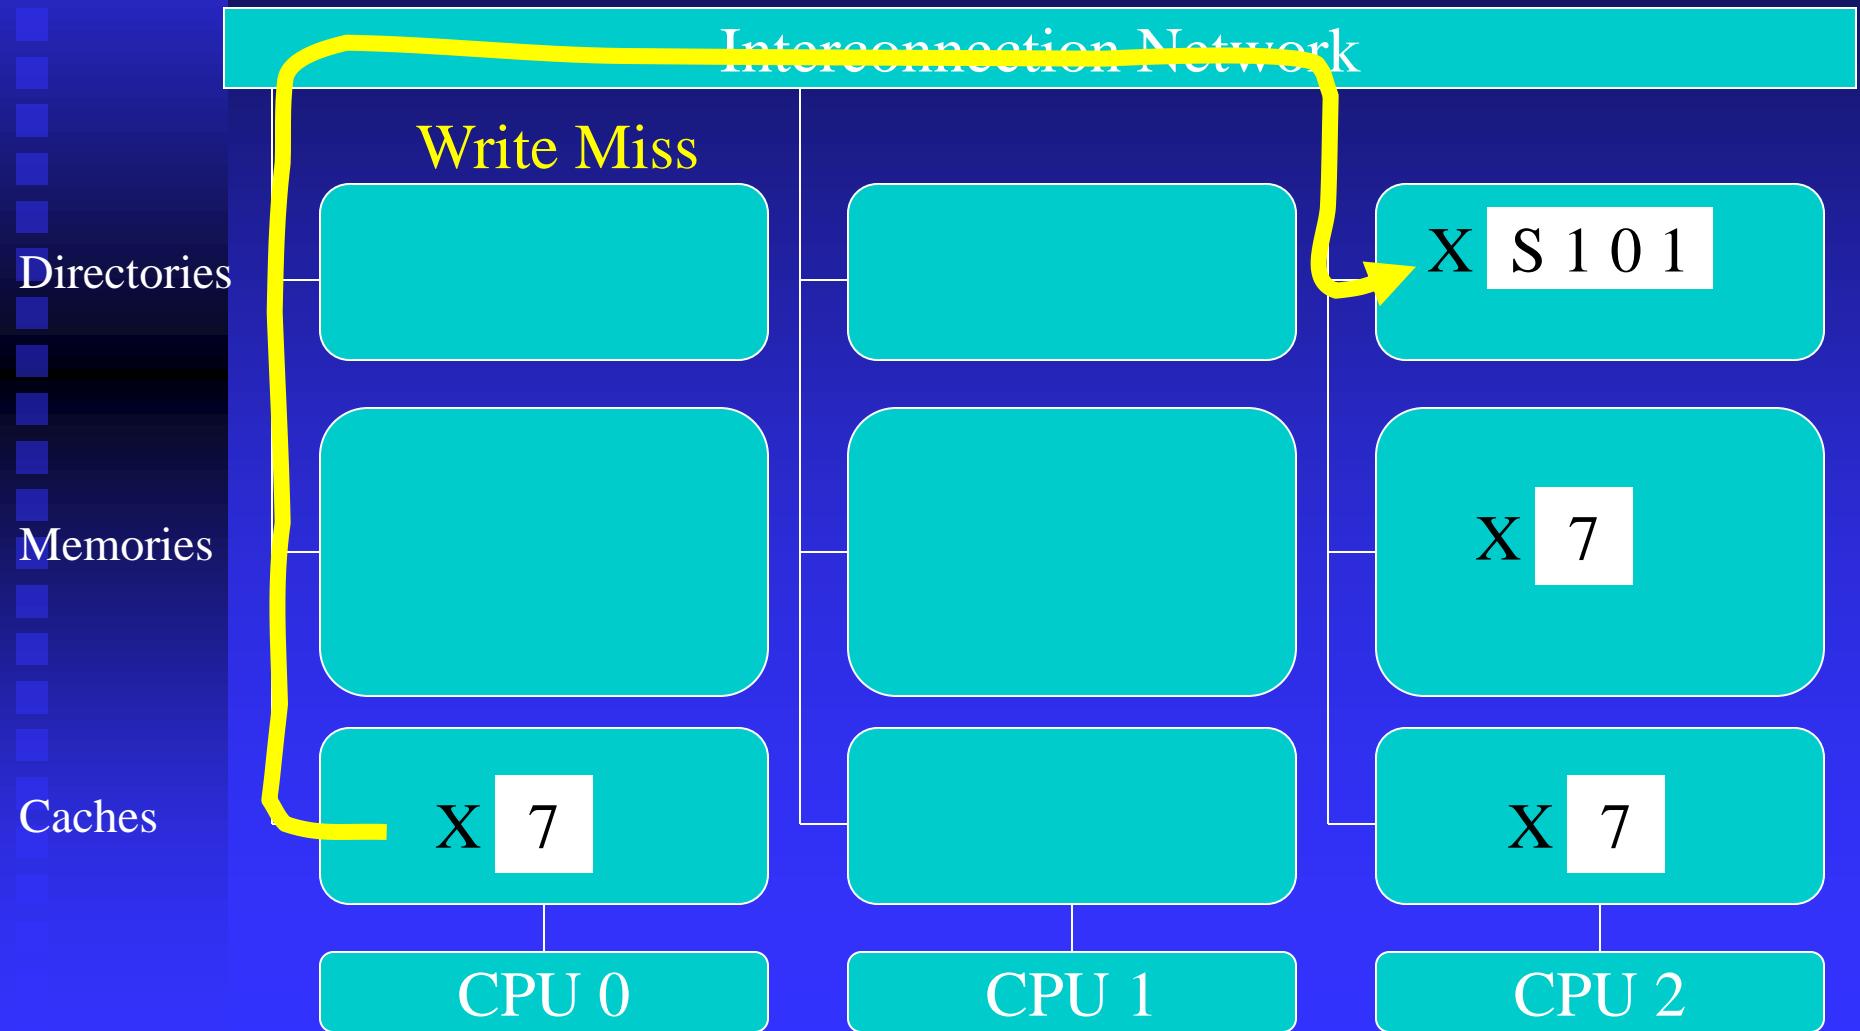
CPU 2 Reads X



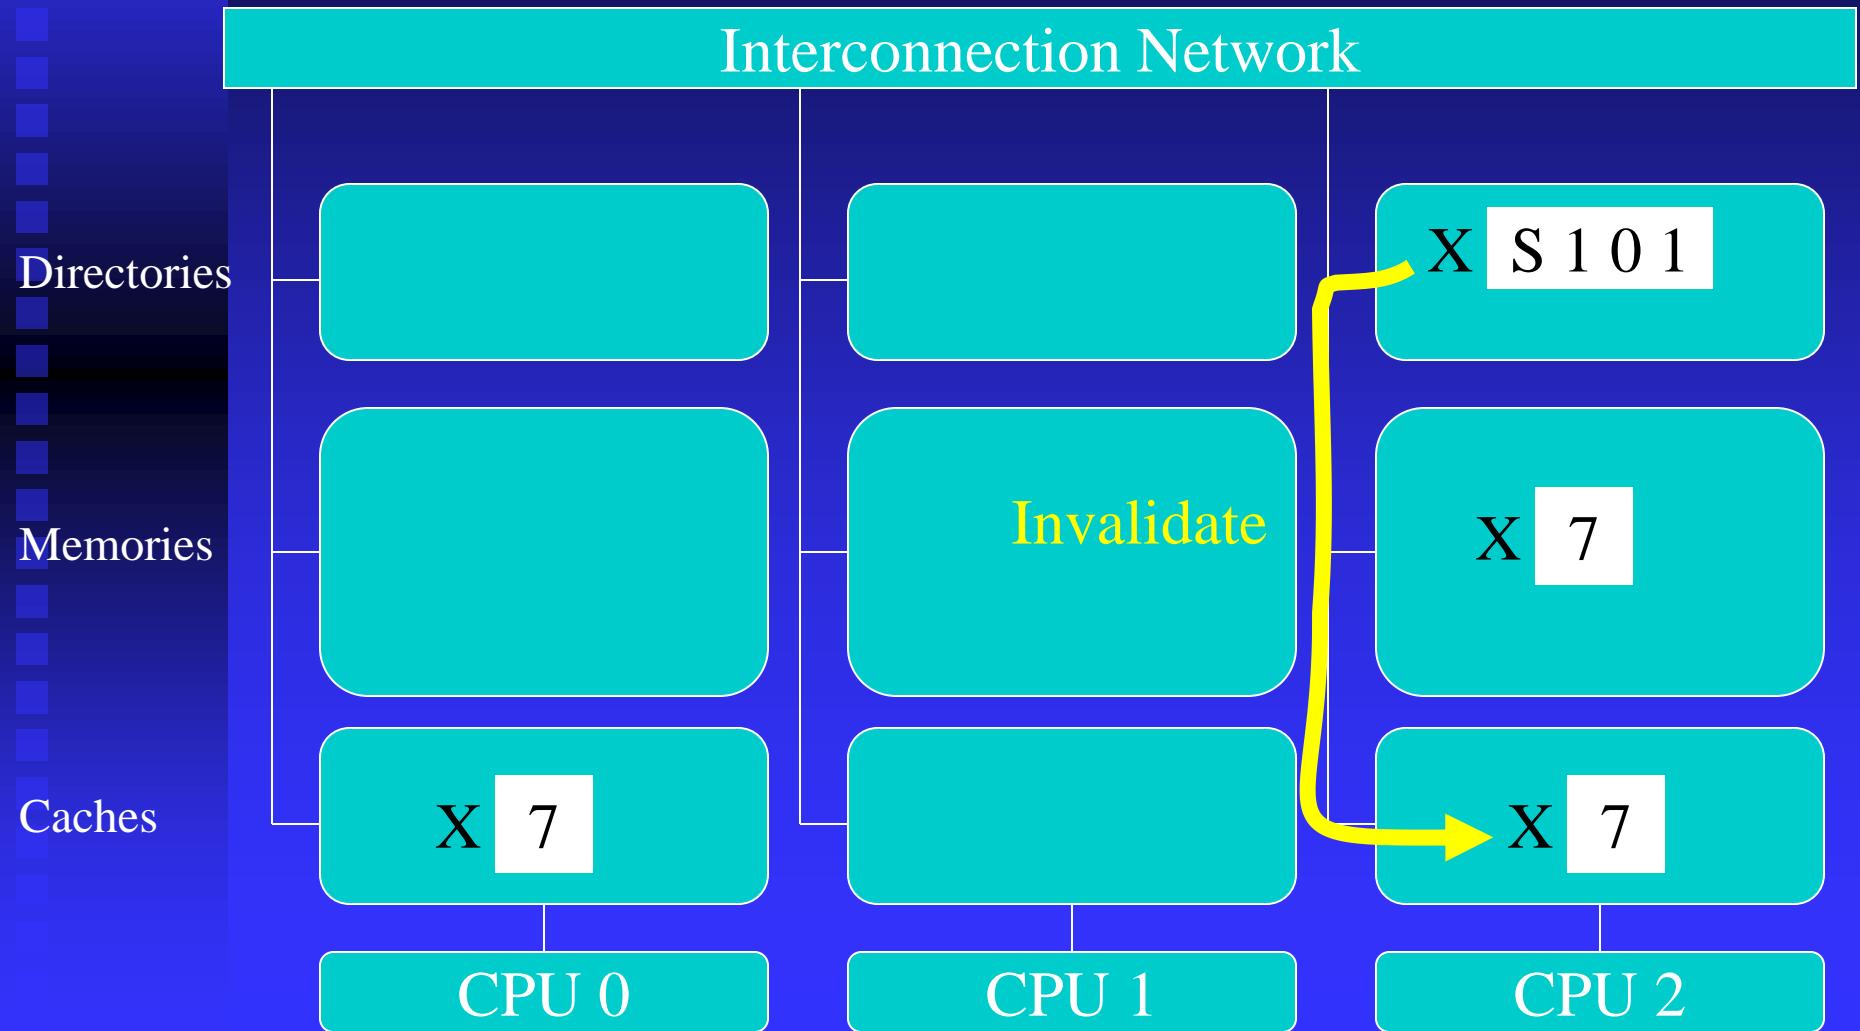
CPU 2 Reads X



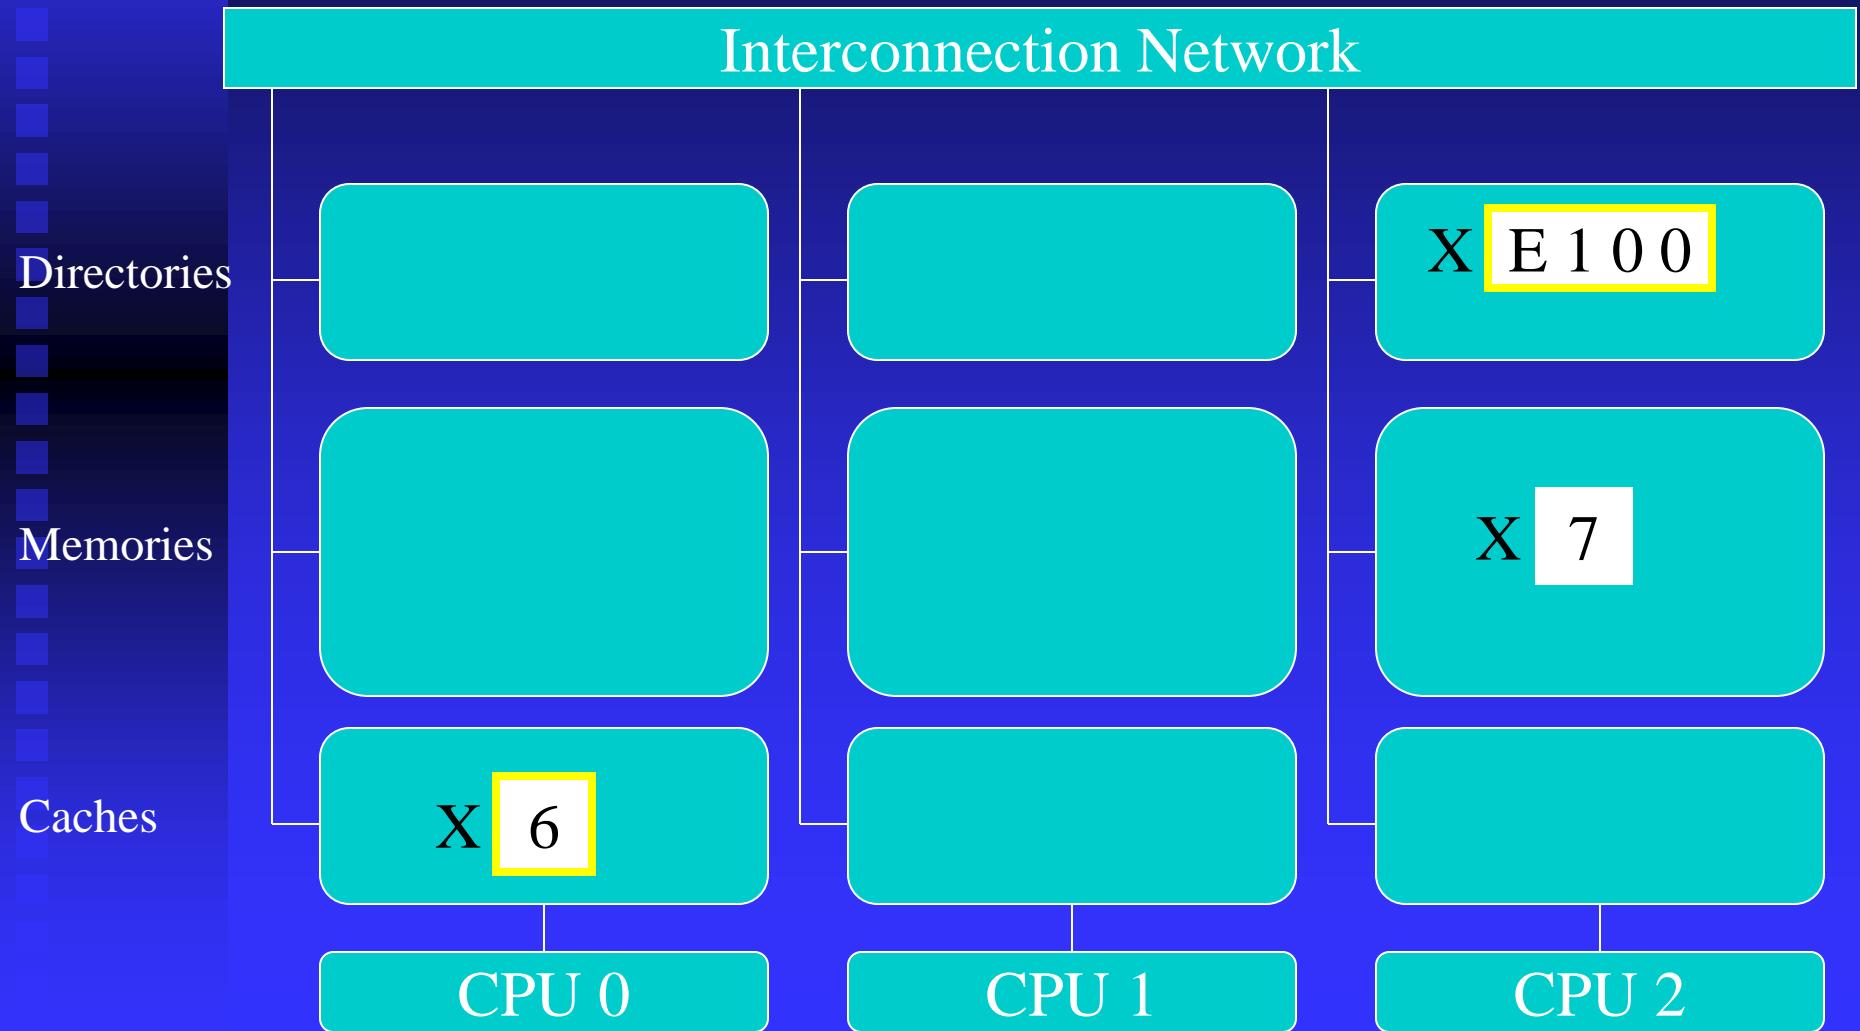
CPU 0 Writes 6 to X



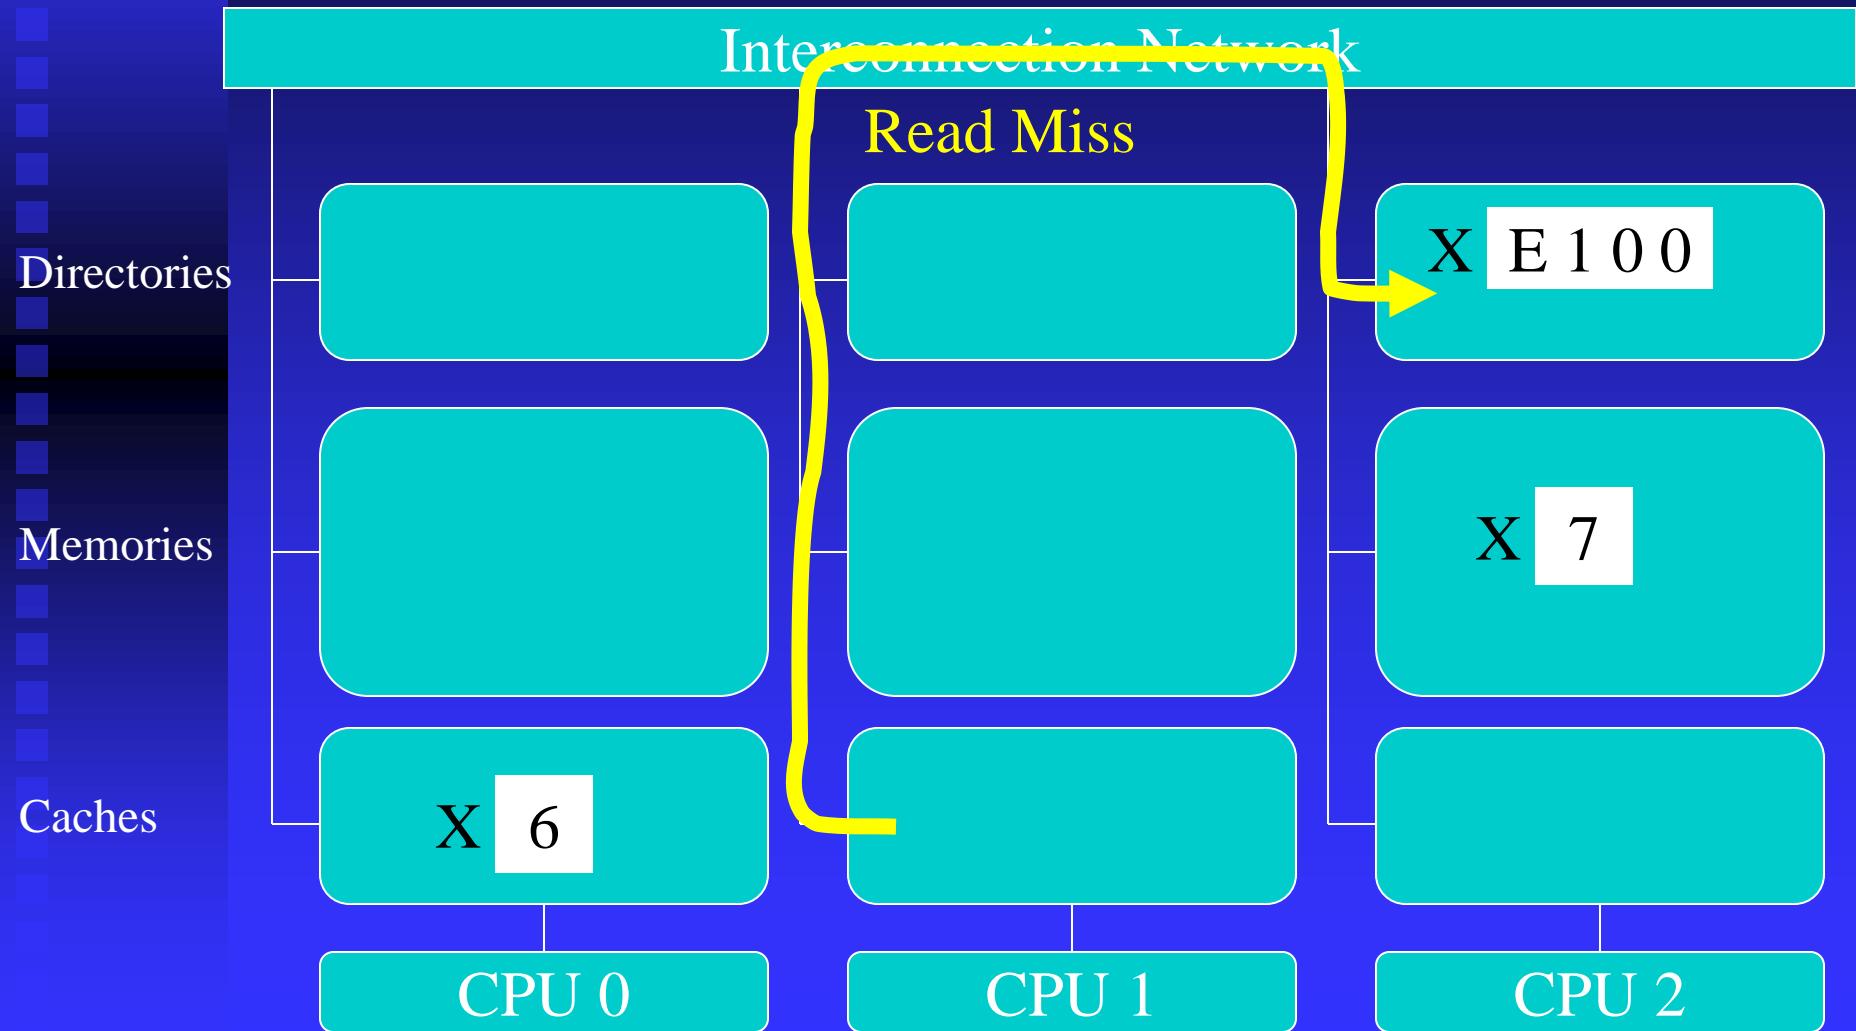
CPU 0 Writes 6 to X



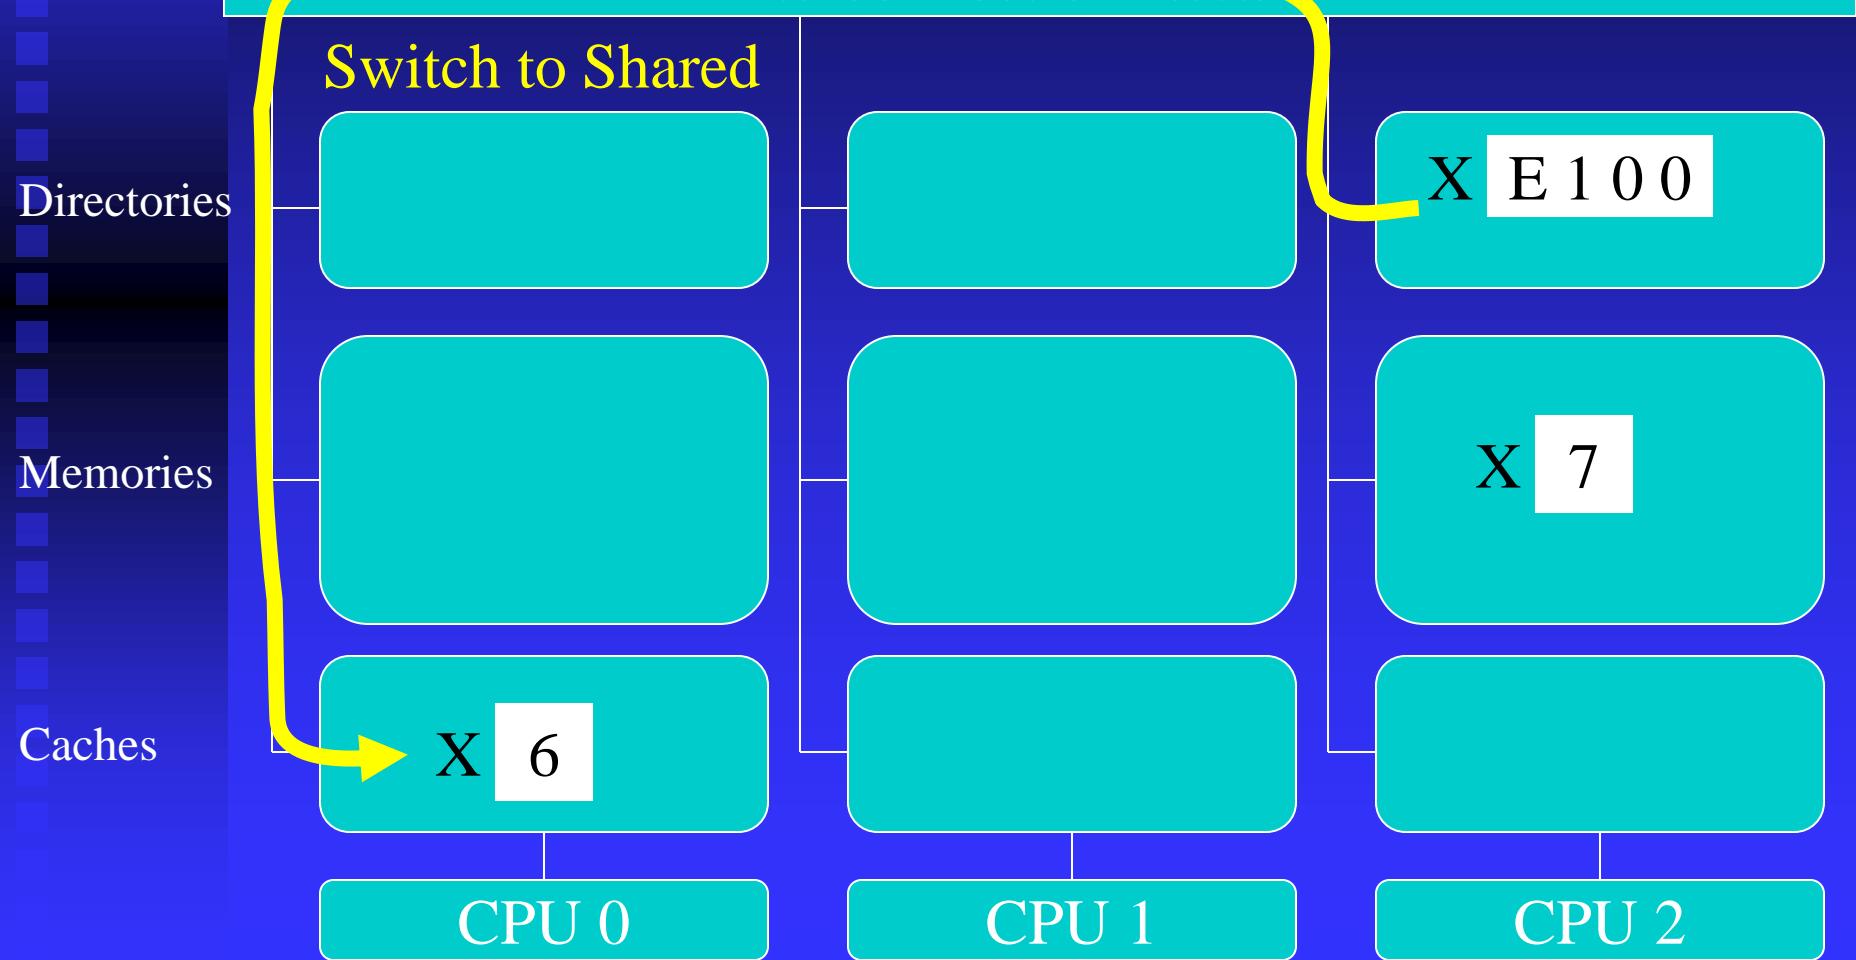
CPU 0 Writes 6 to X



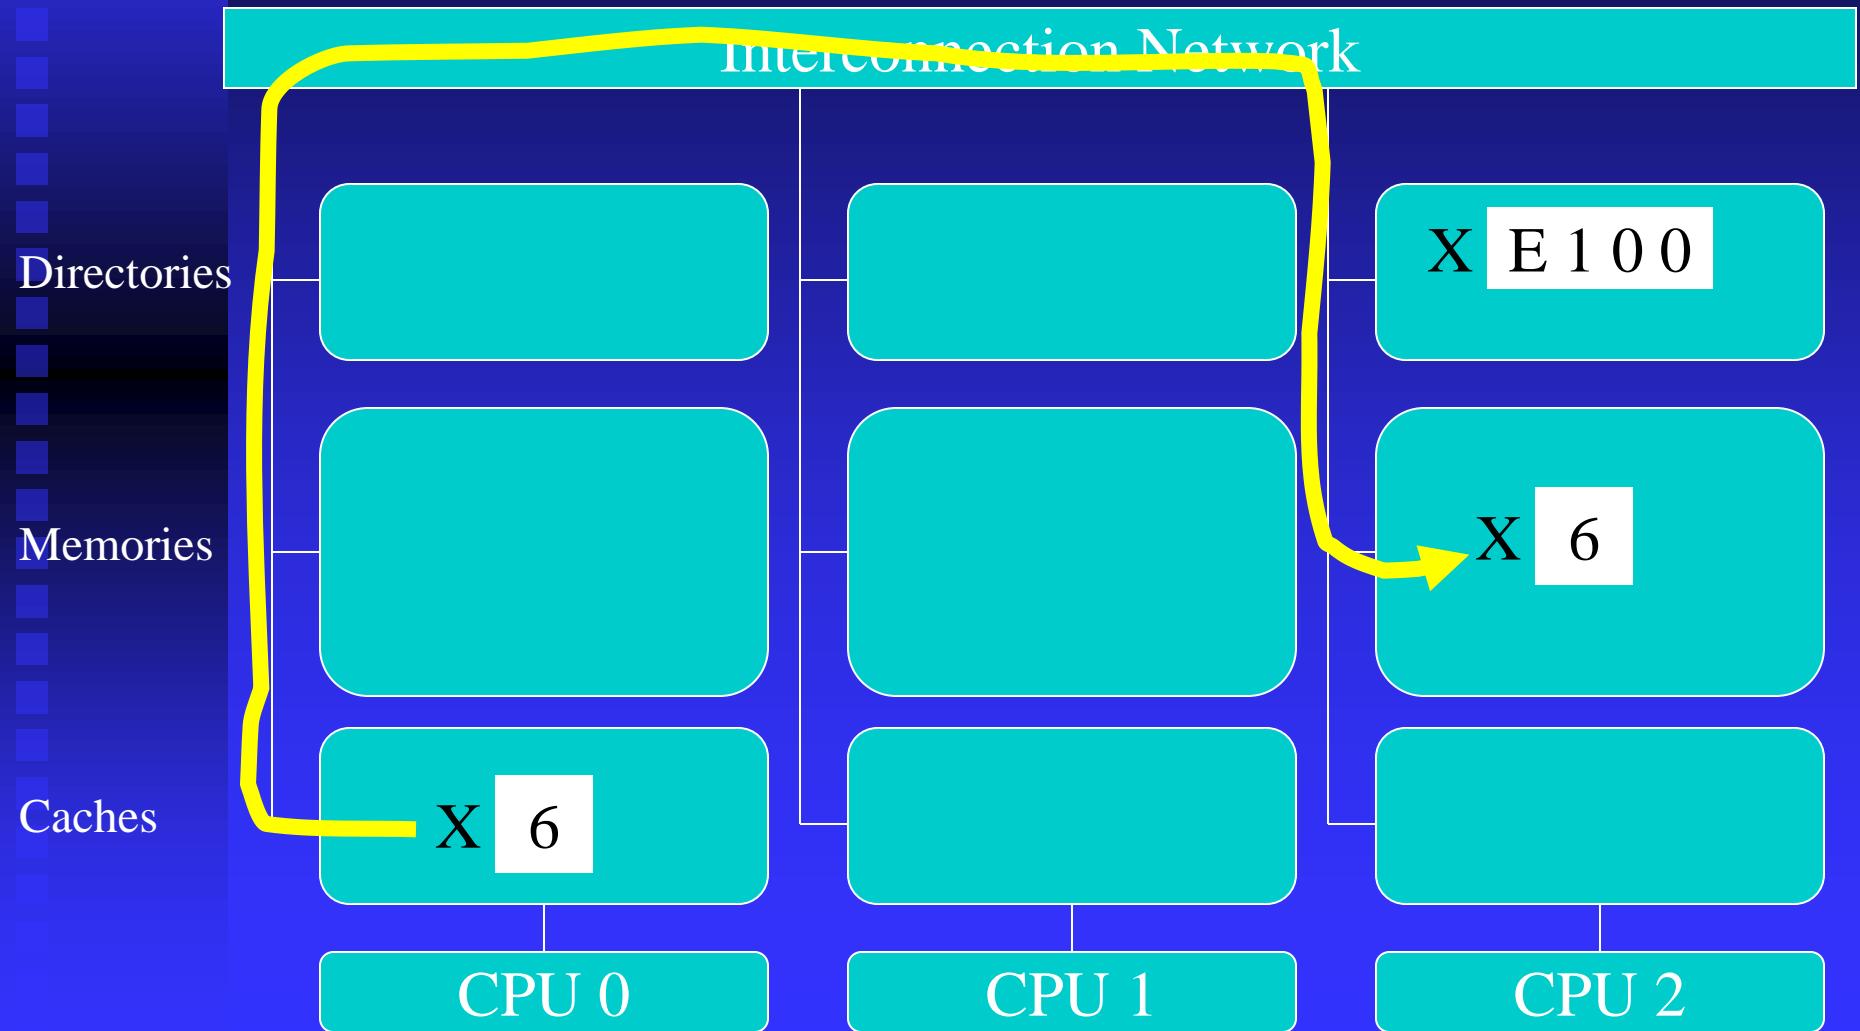
CPU 1 Reads X



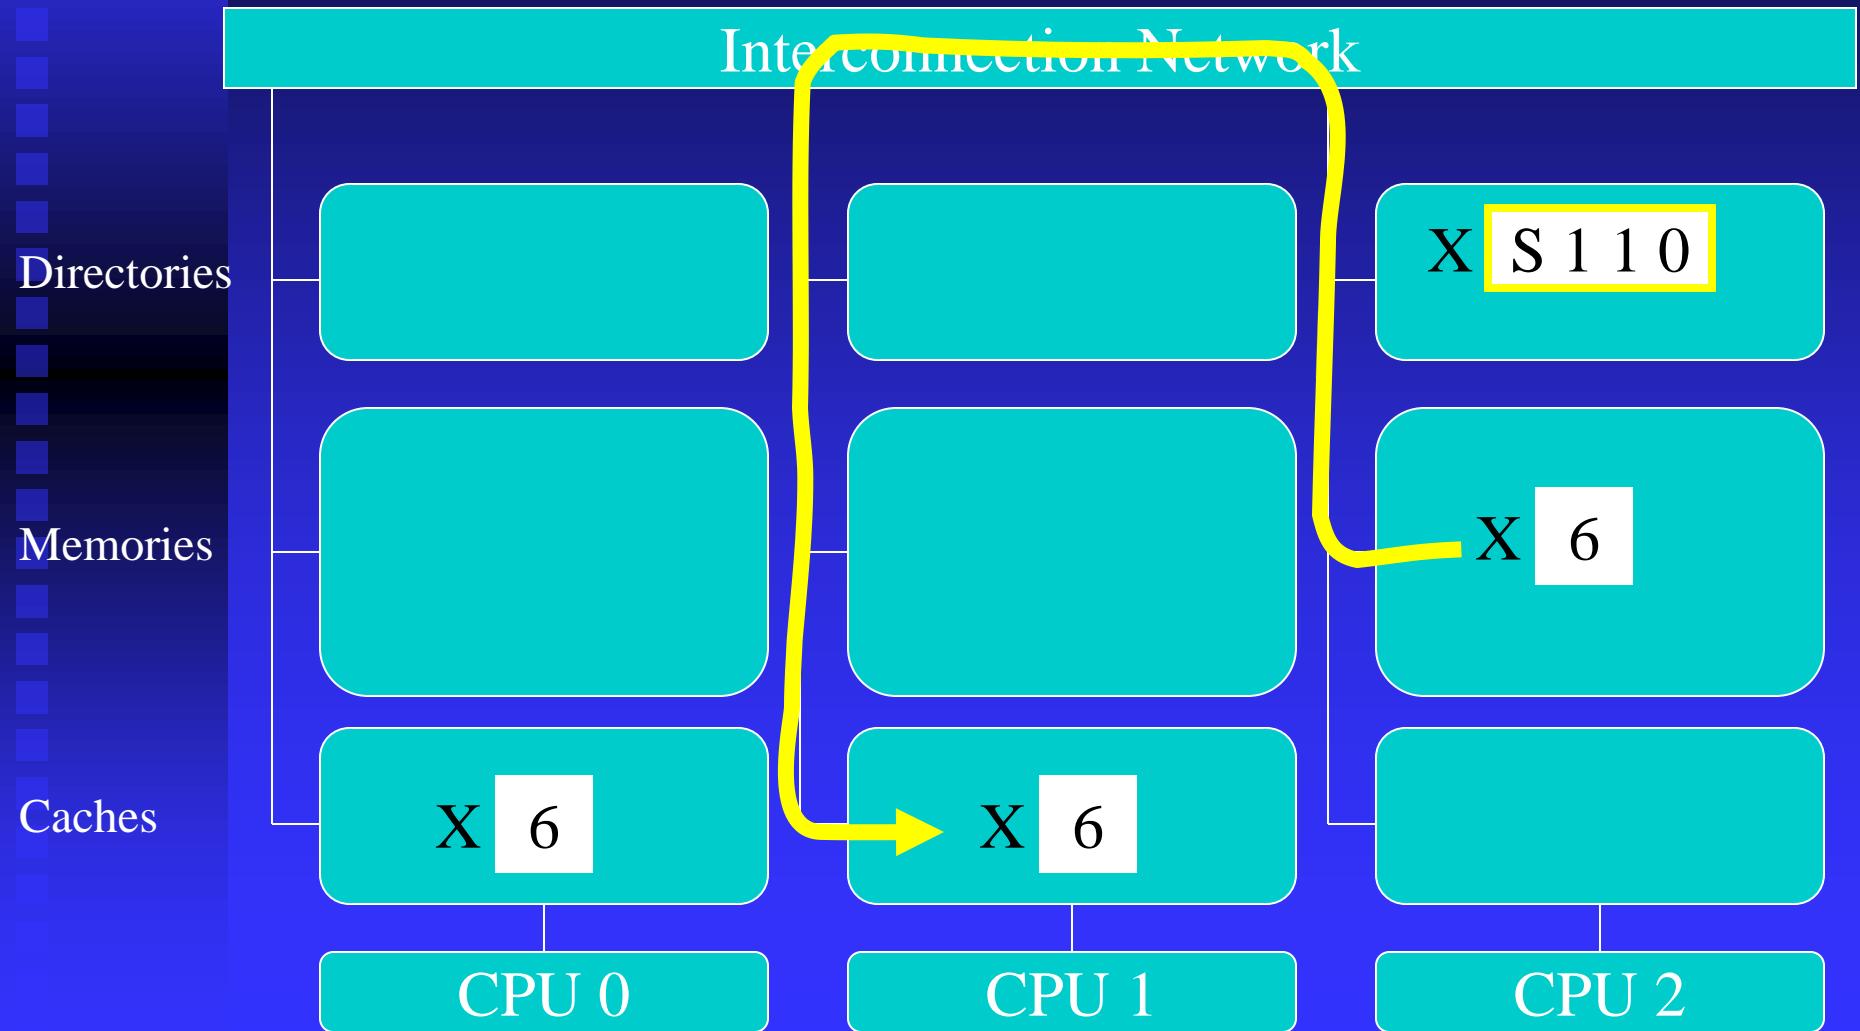
CPU 1 Reads X



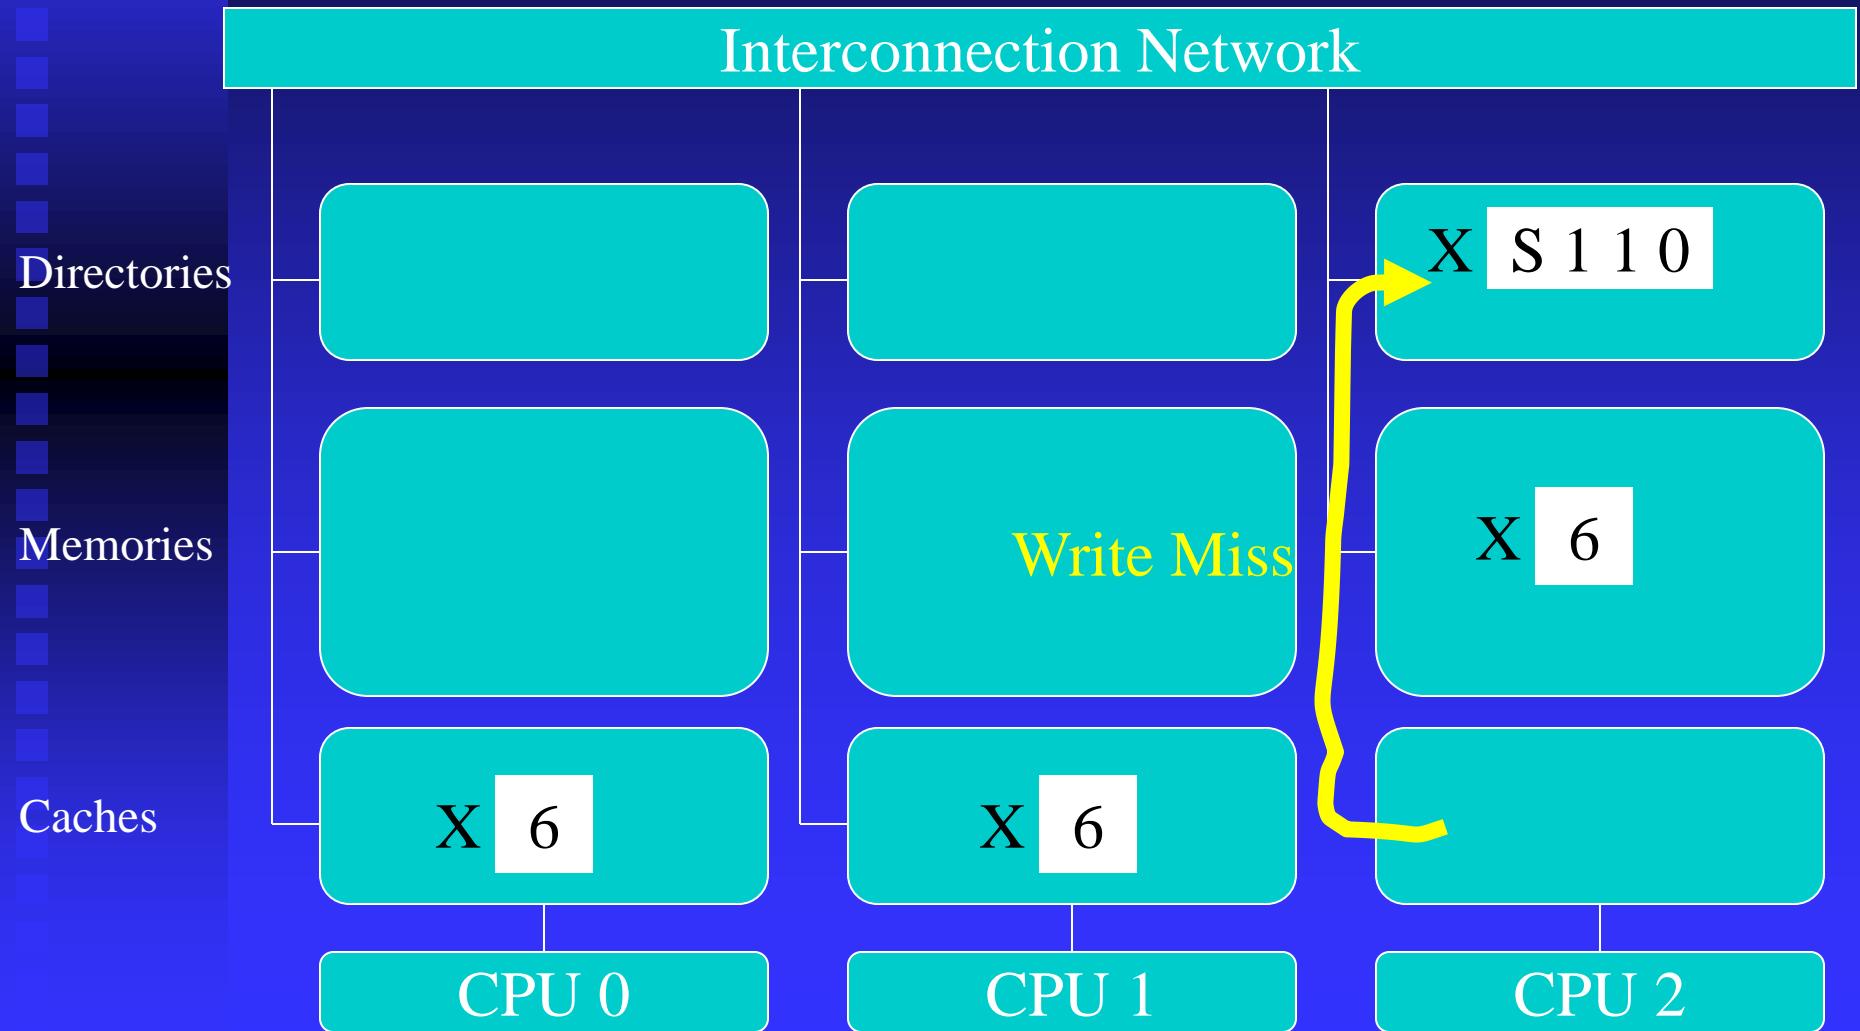
CPU 1 Reads X



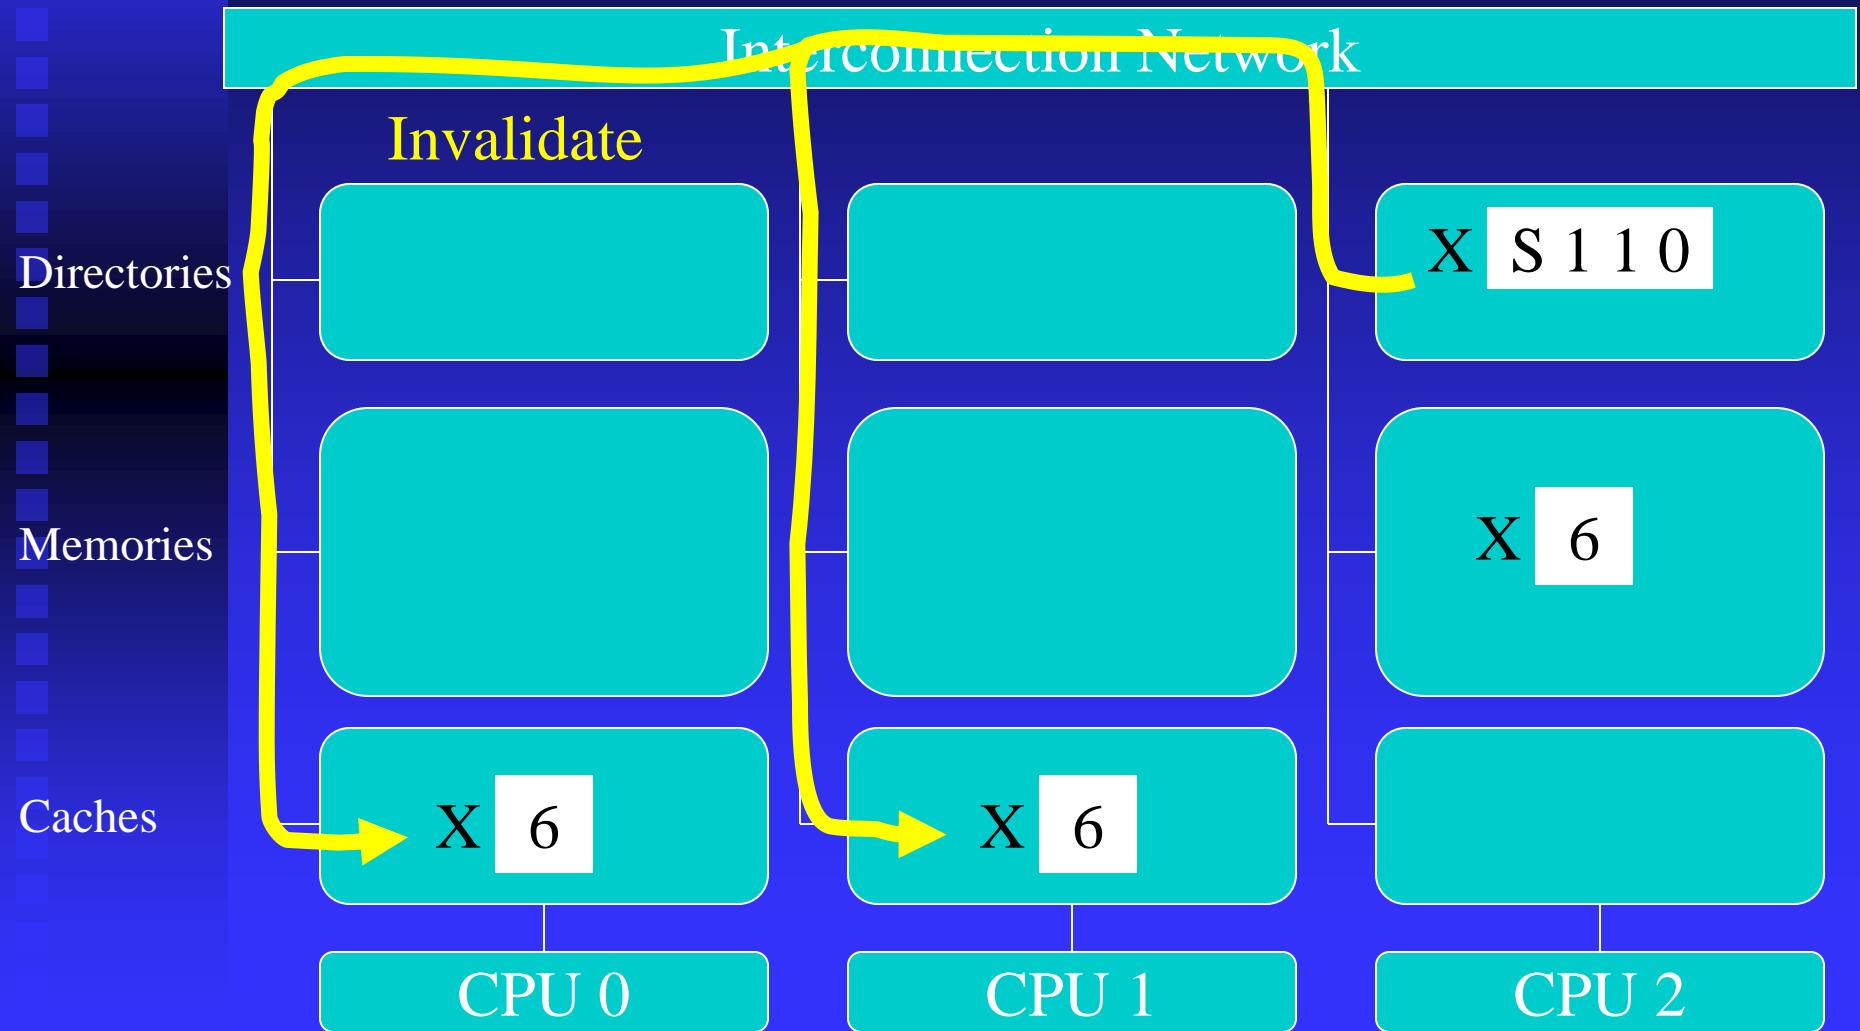
CPU 1 Reads X



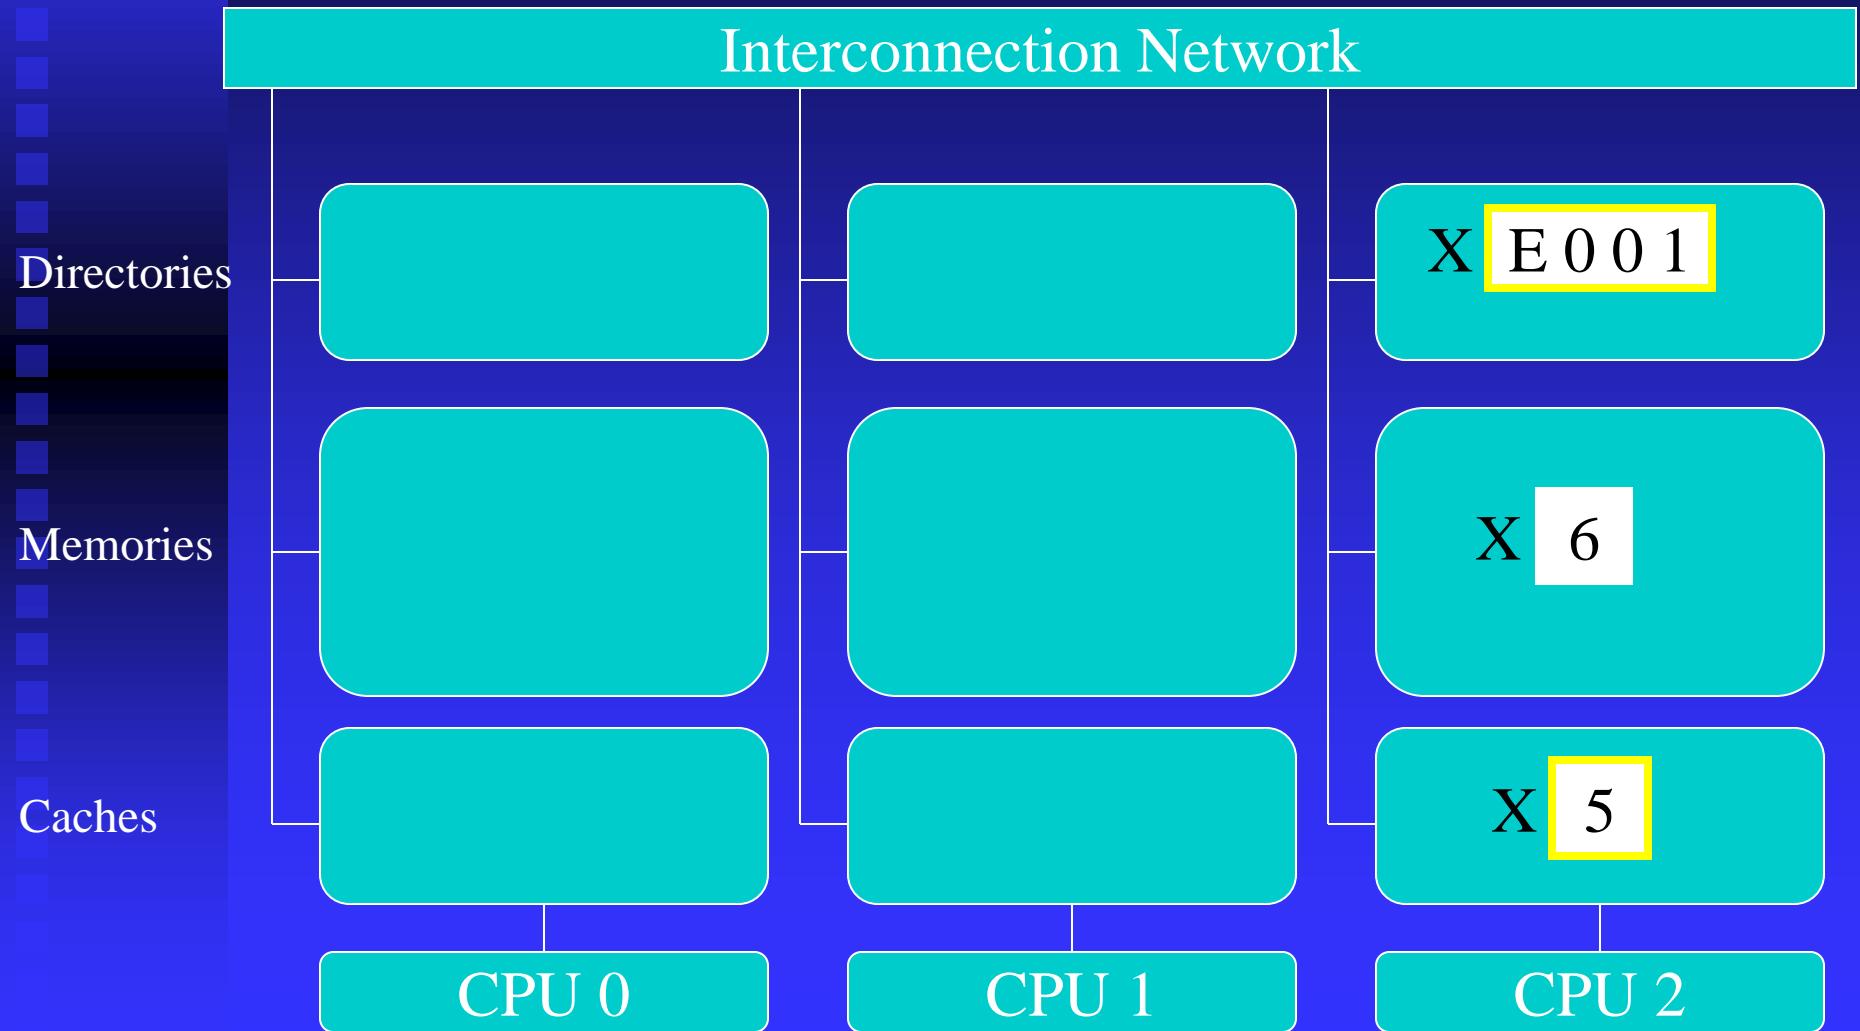
CPU 2 Writes 5 to X



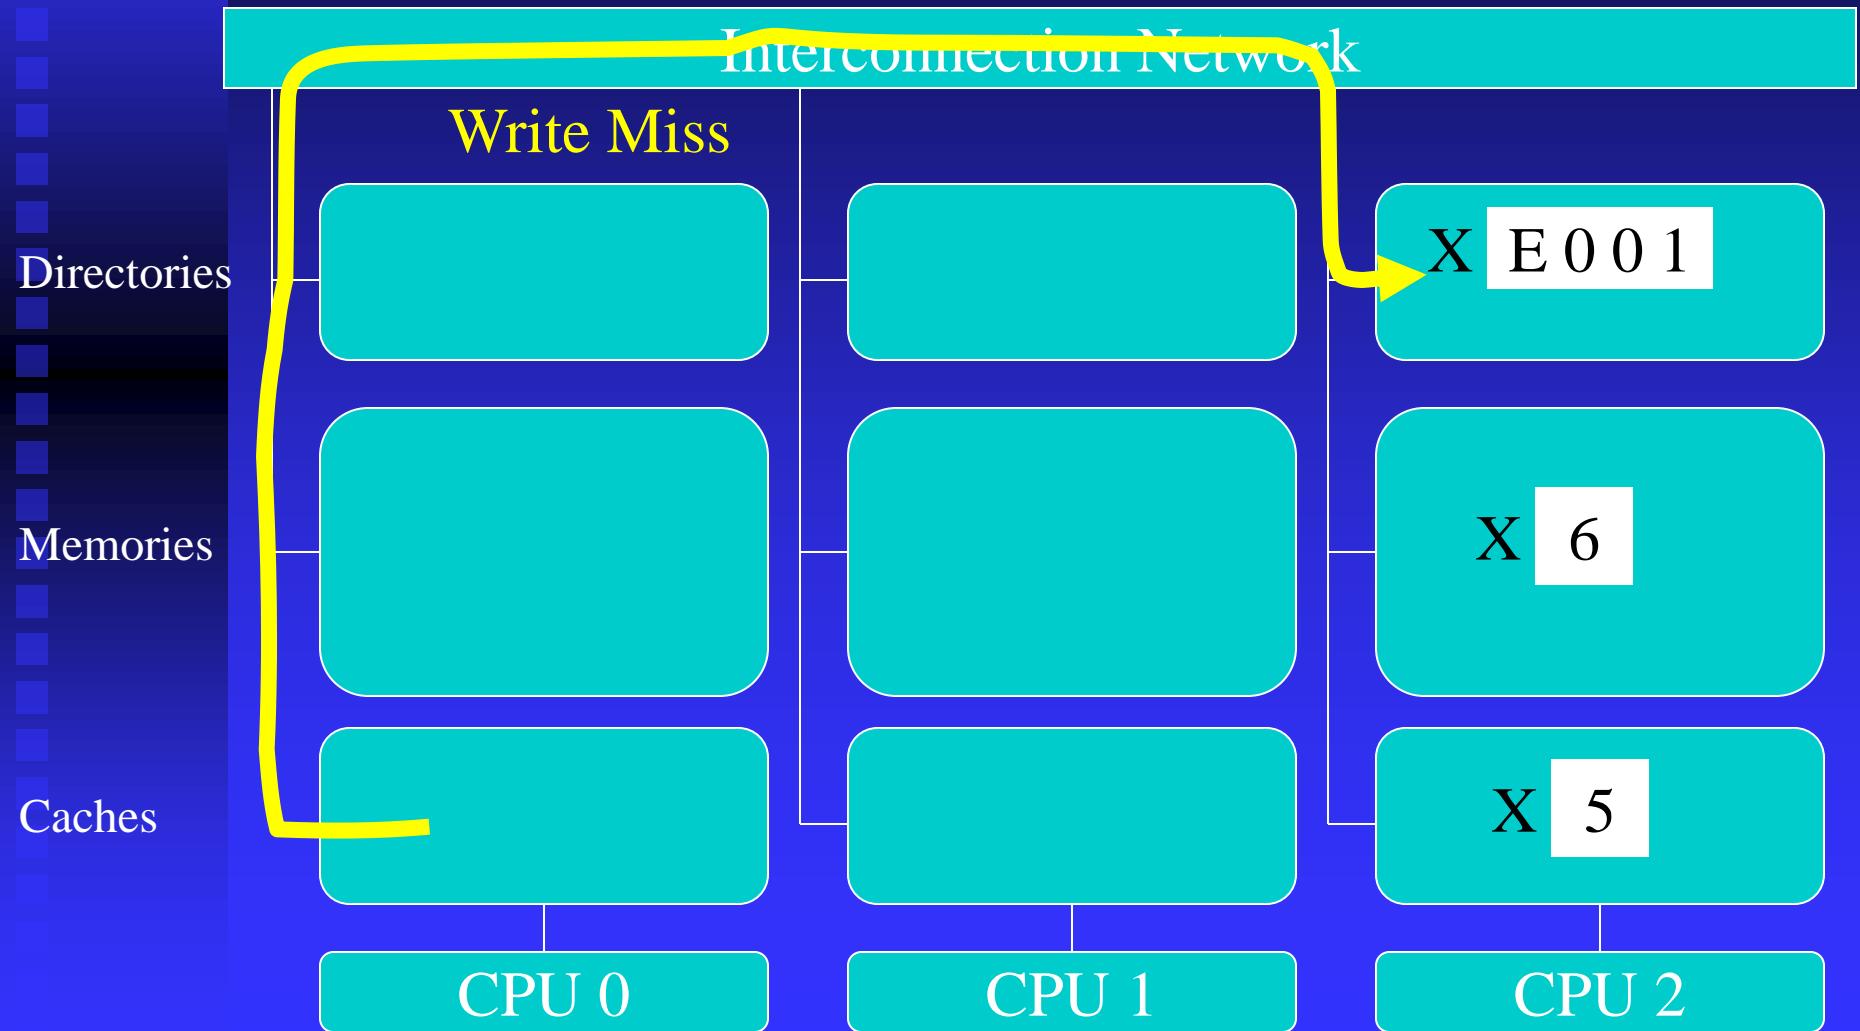
CPU 2 Writes 5 to X



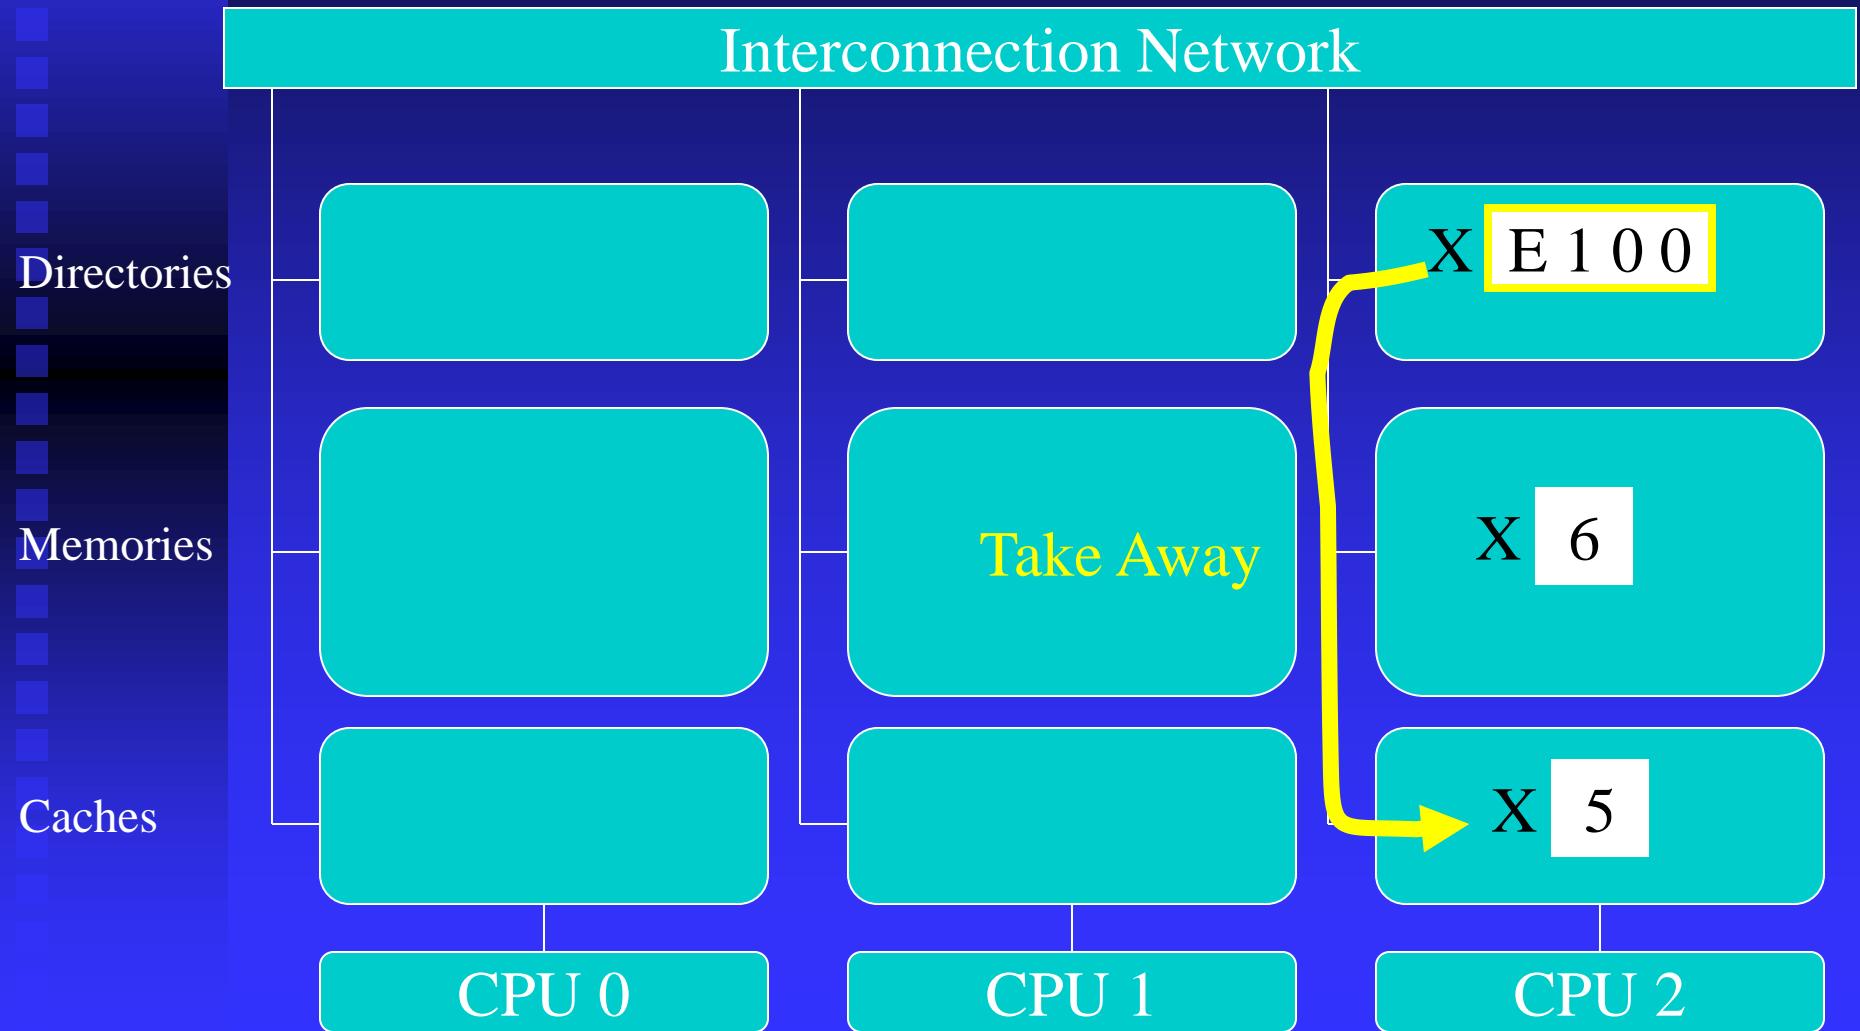
CPU 2 Writes 5 to X



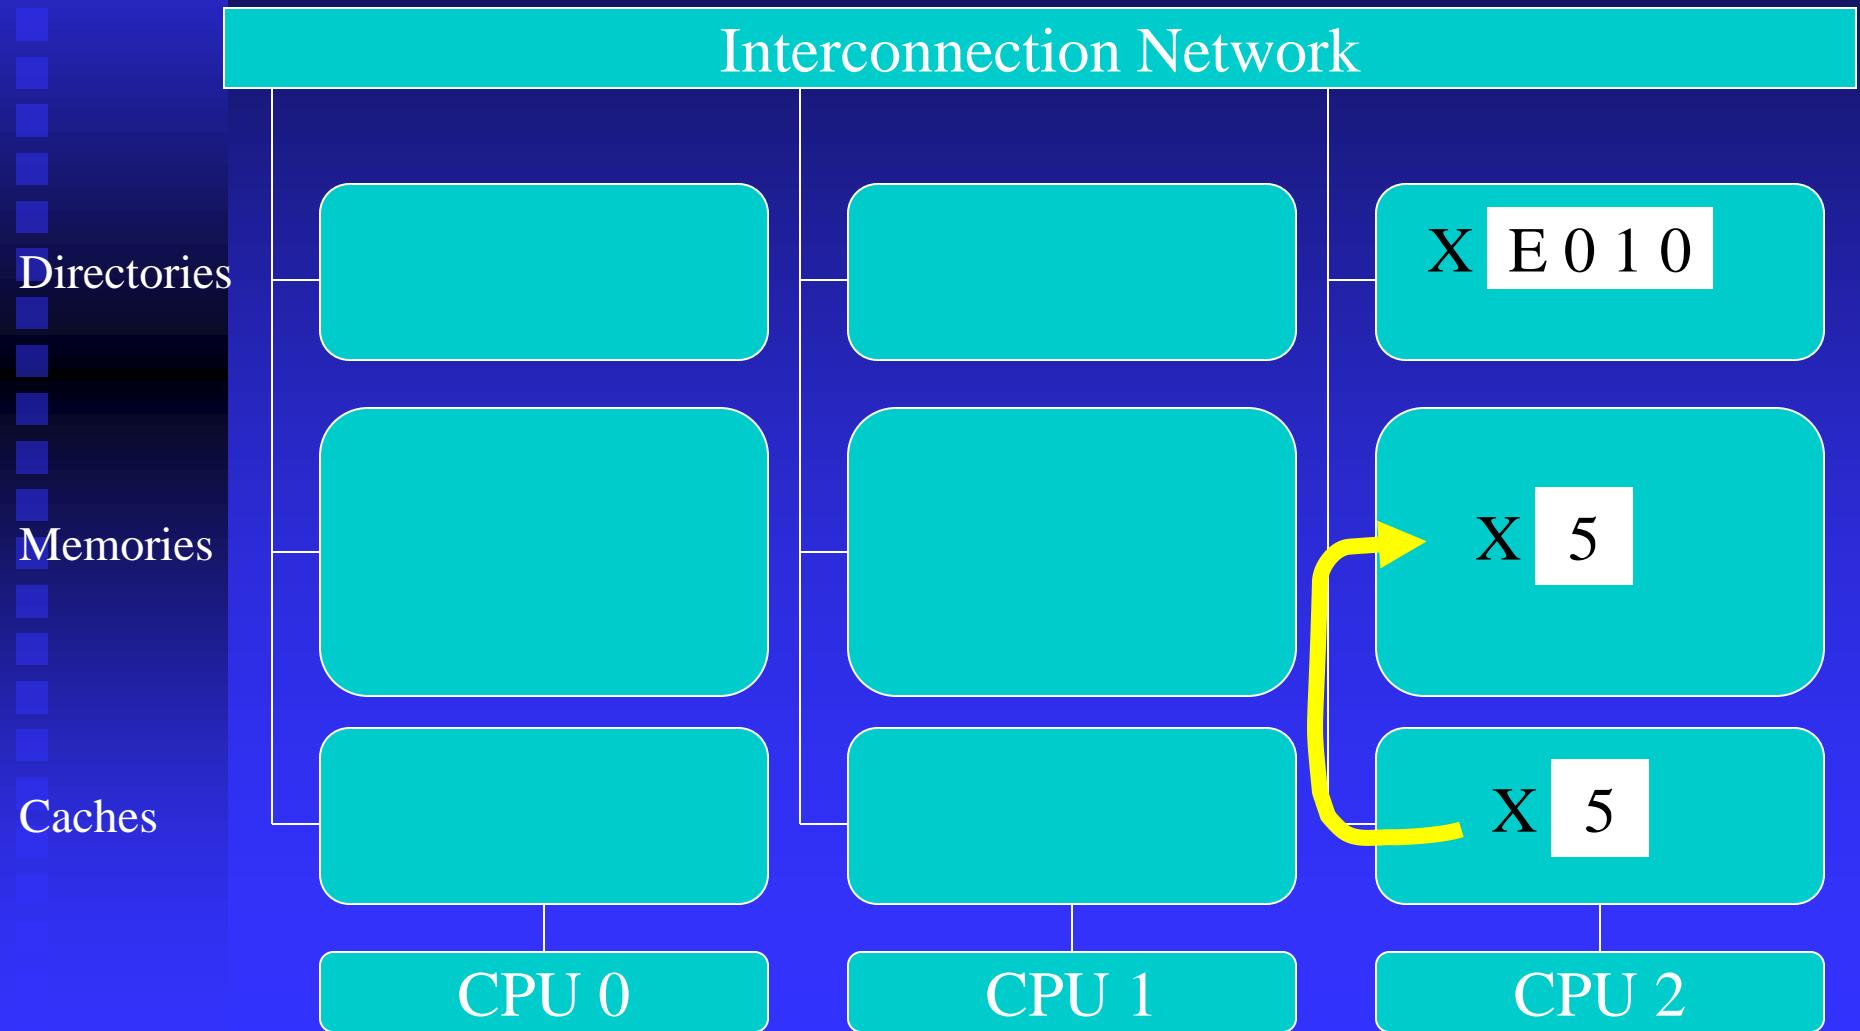
CPU 0 Writes 4 to X



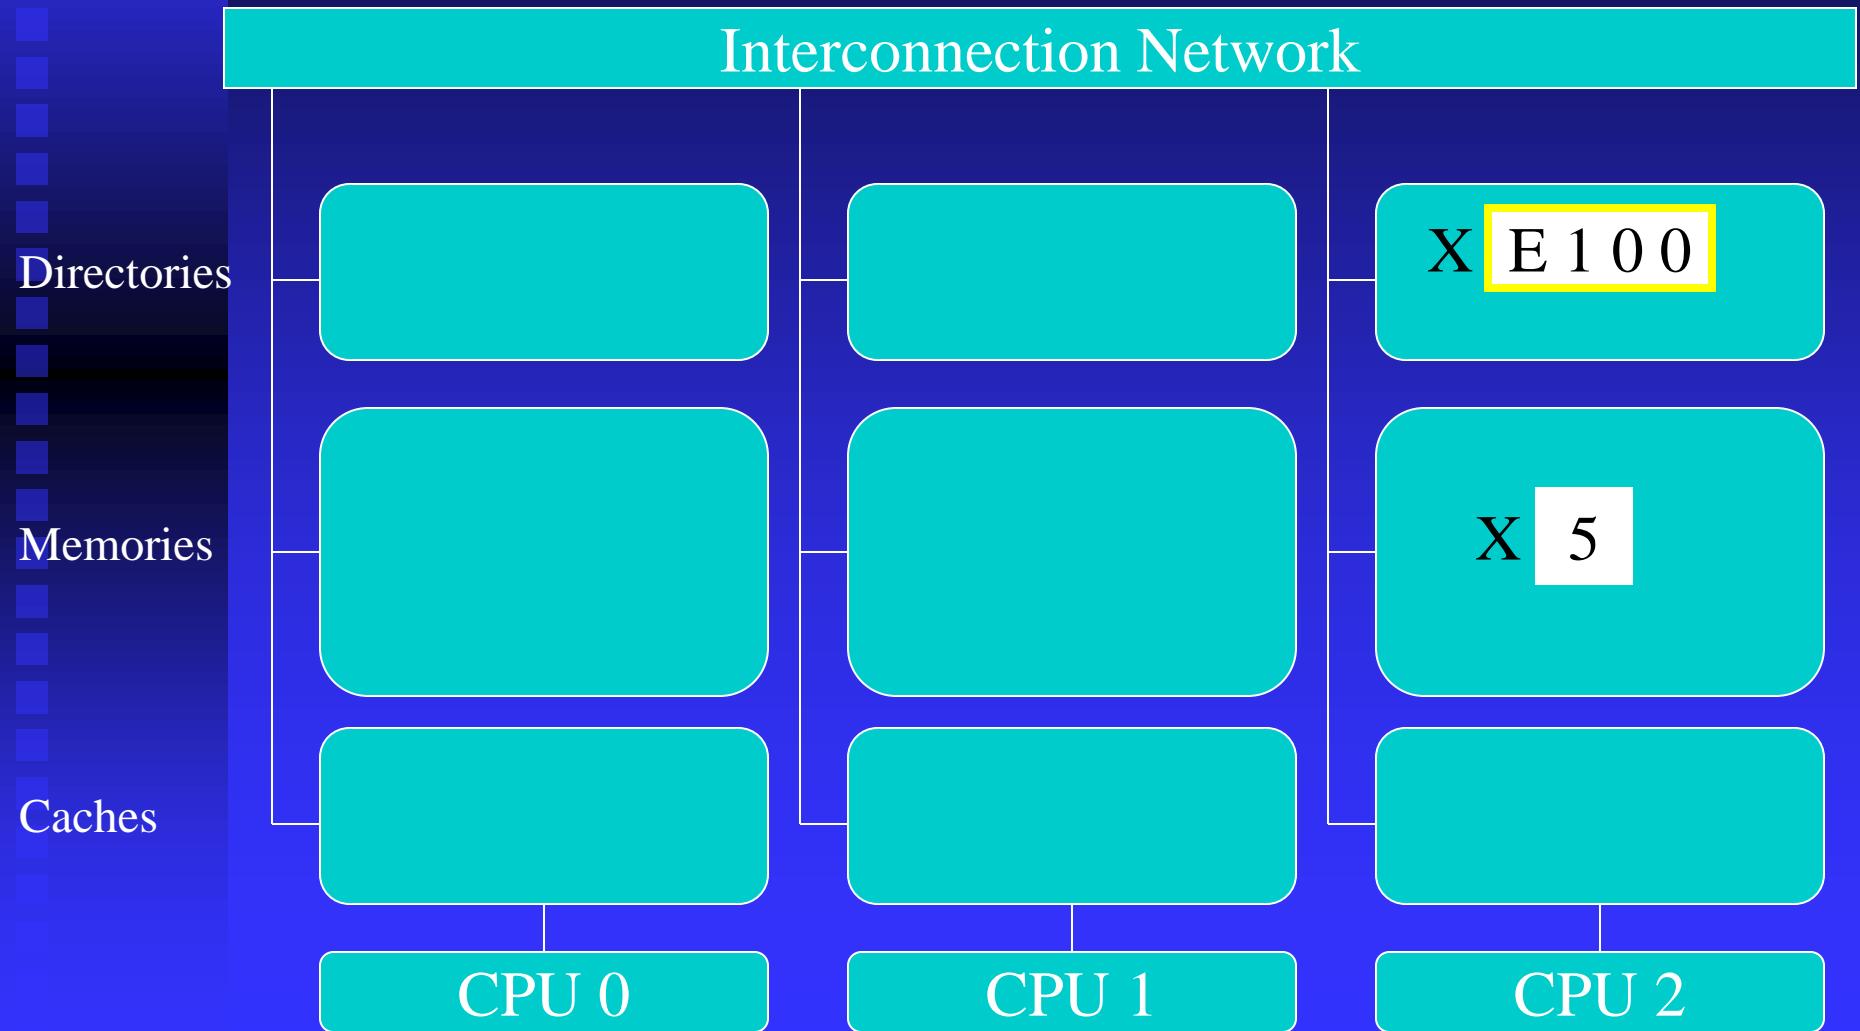
CPU 0 Writes 4 to X



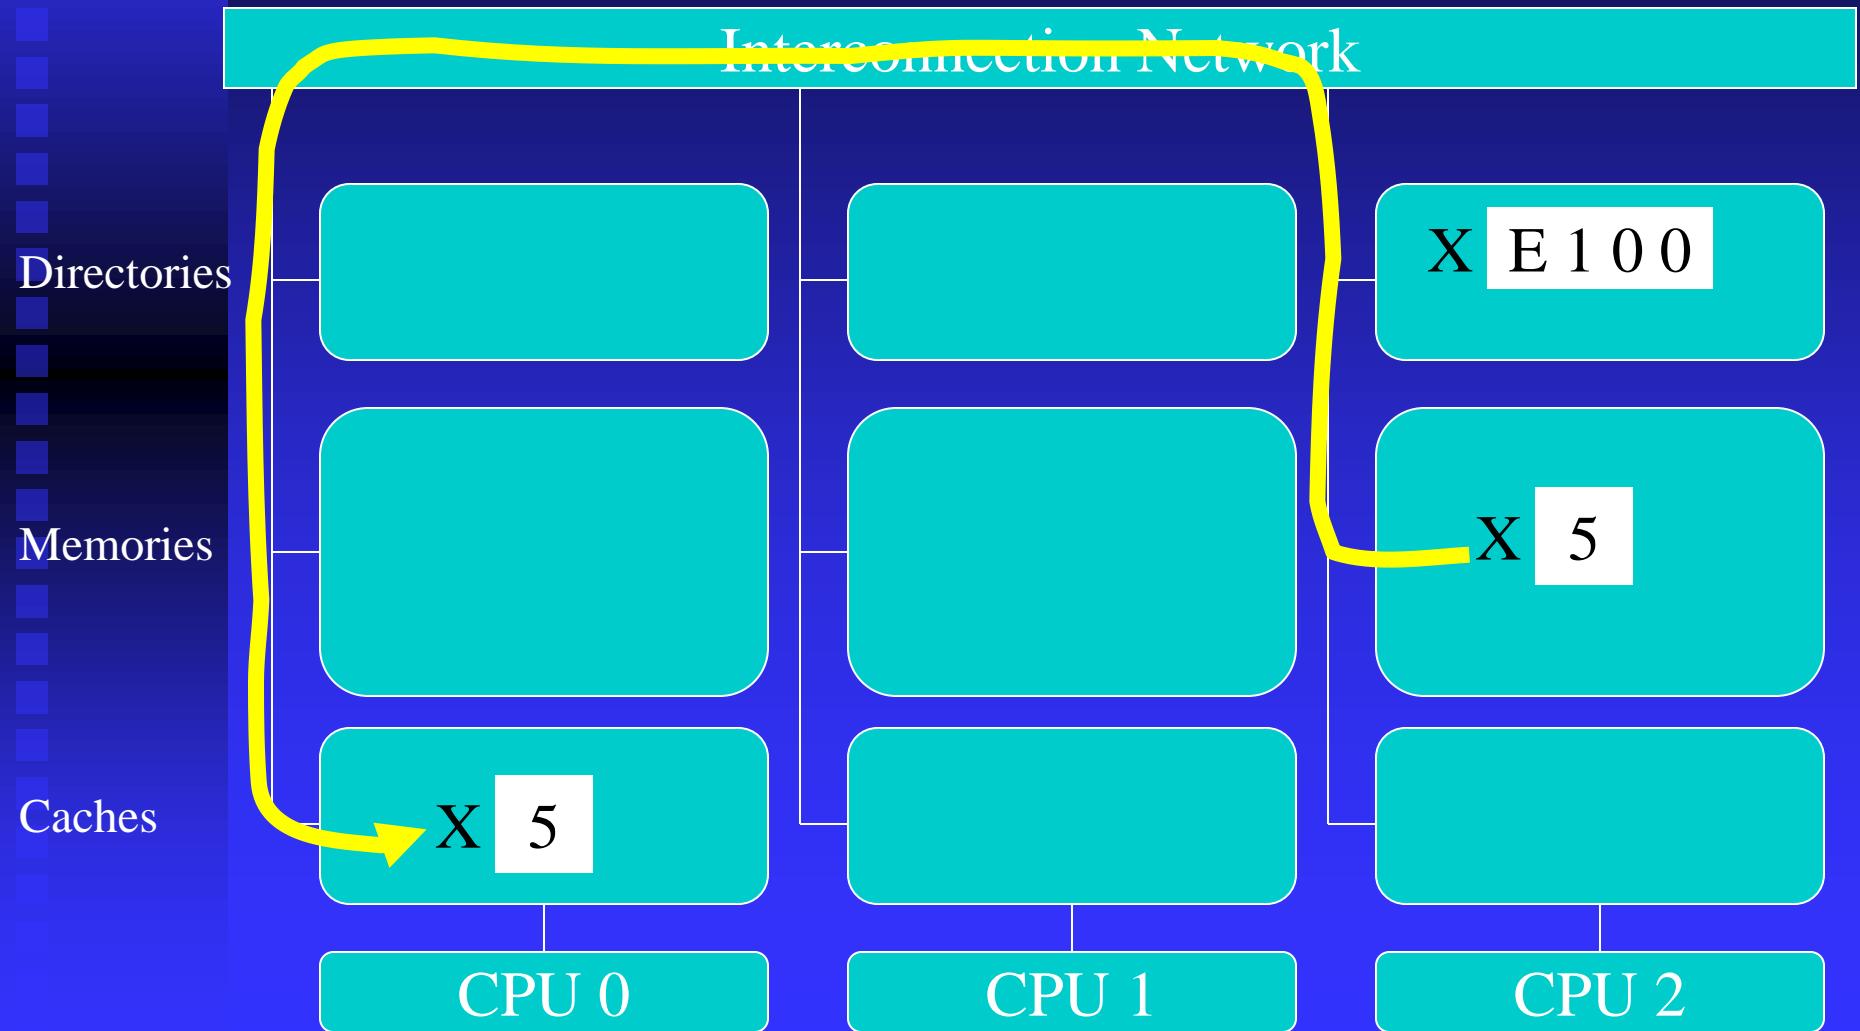
CPU 0 Writes 4 to X



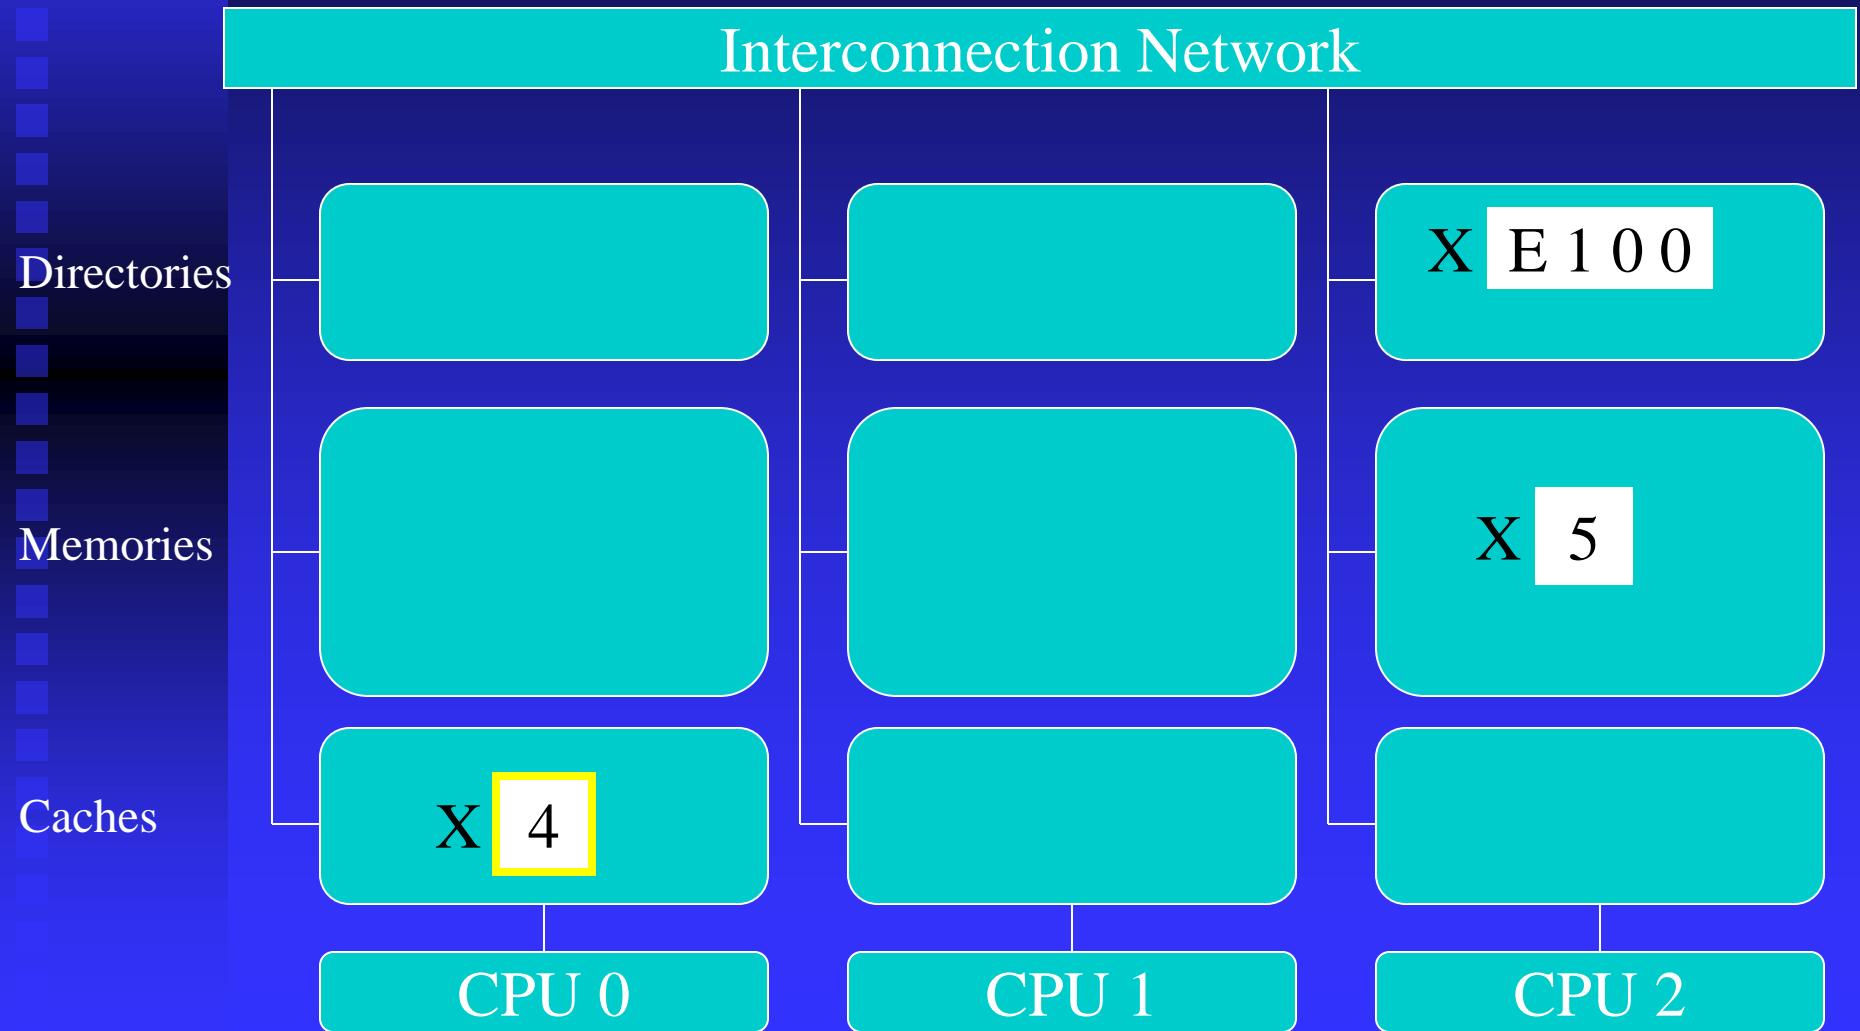
CPU 0 Writes 4 to X



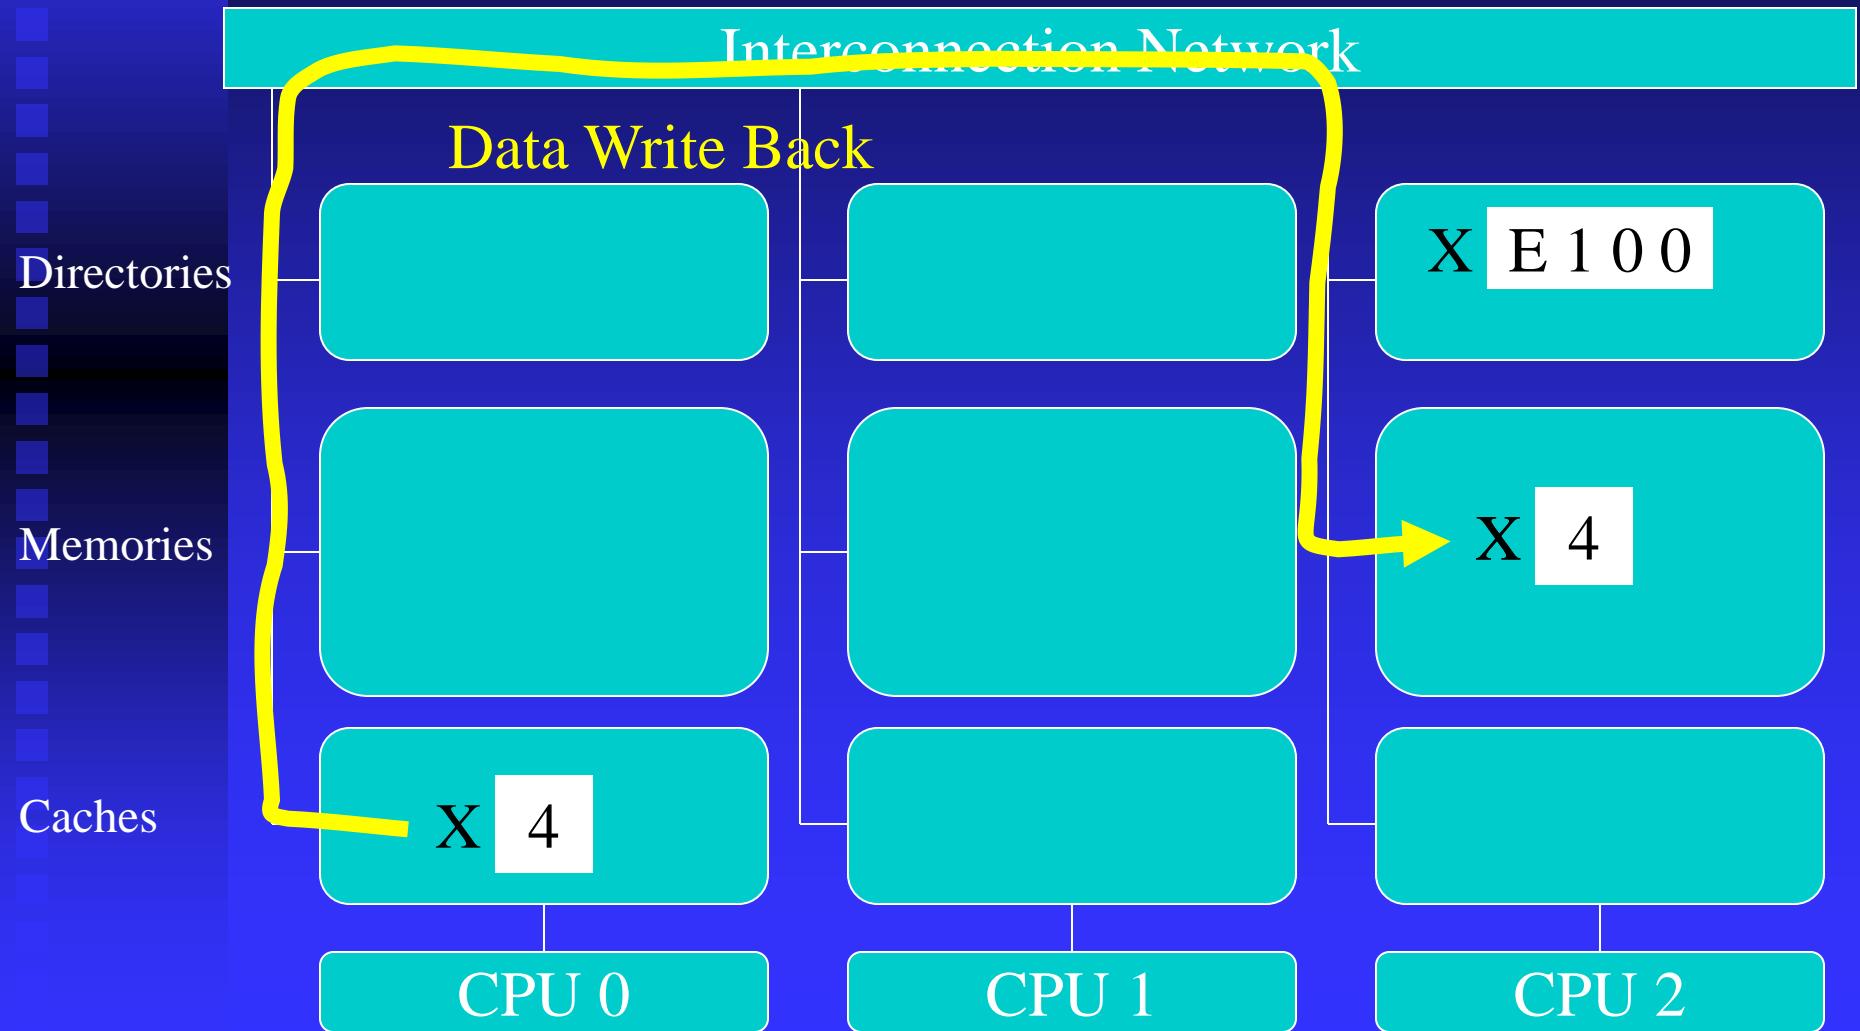
CPU 0 Writes 4 to X



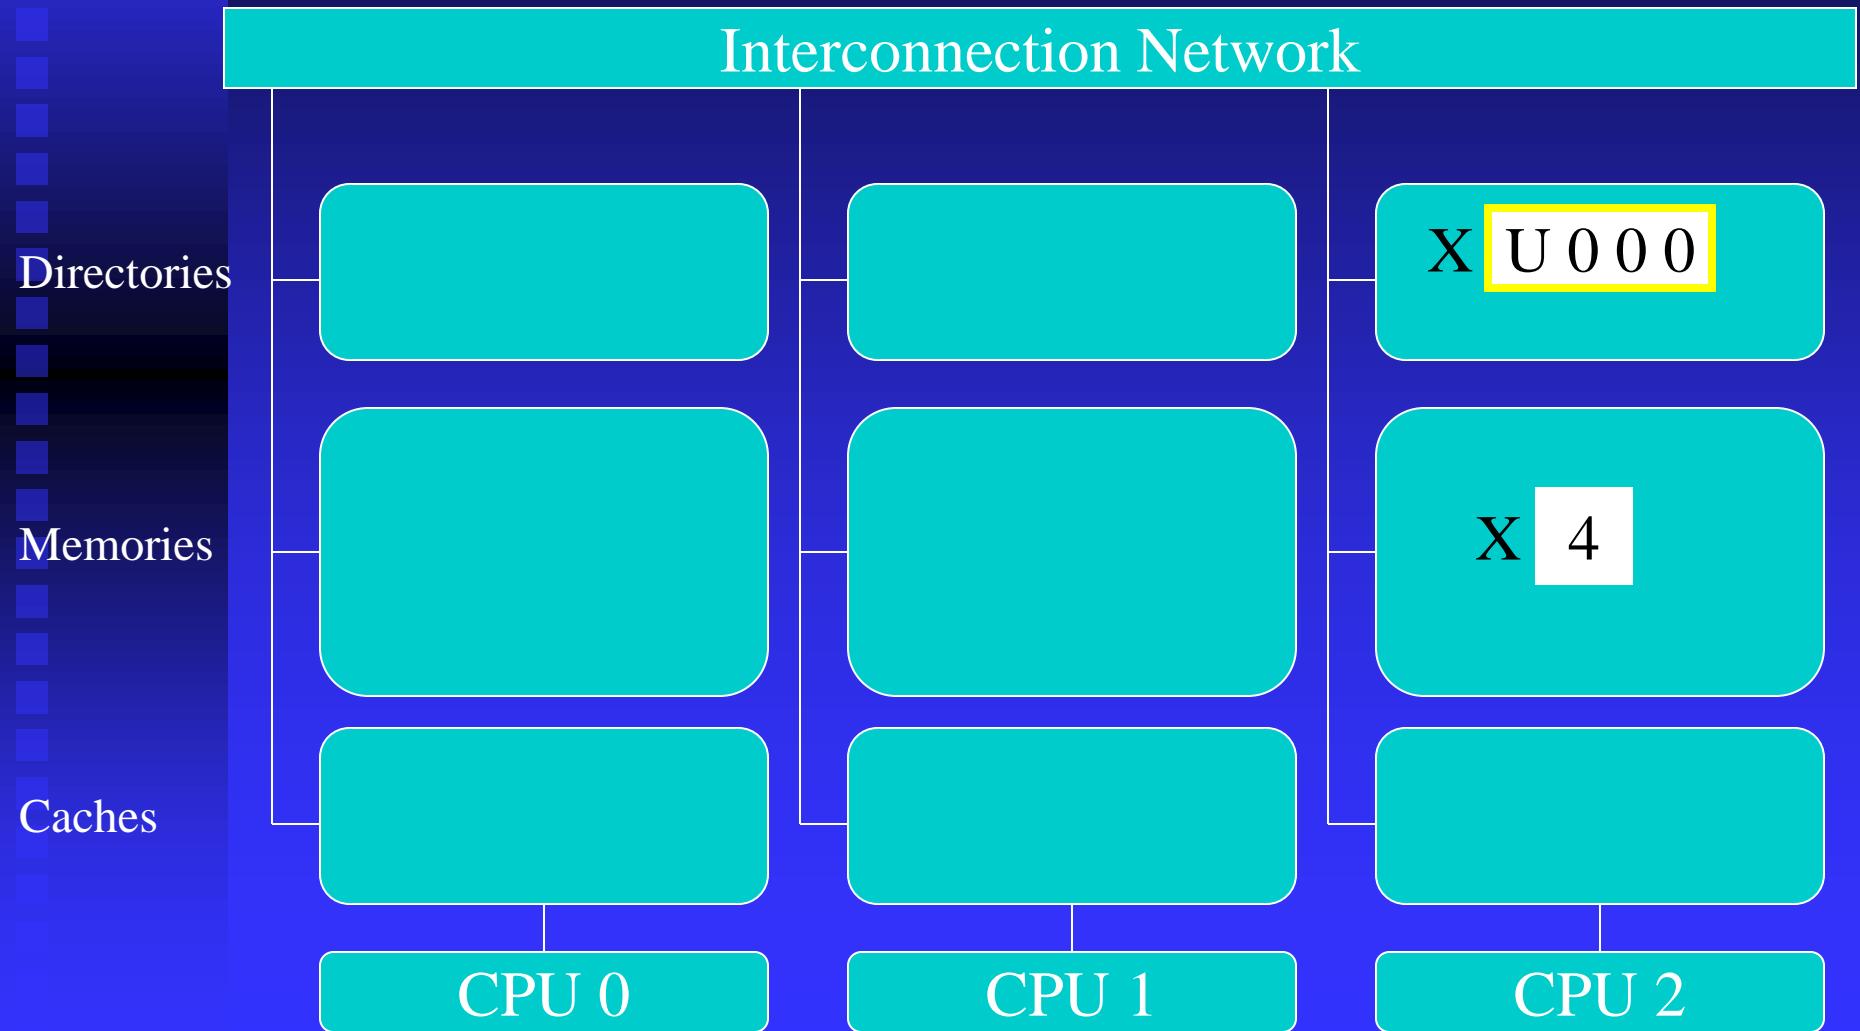
CPU 0 Writes 4 to X



CPU 0 Writes Back X Block



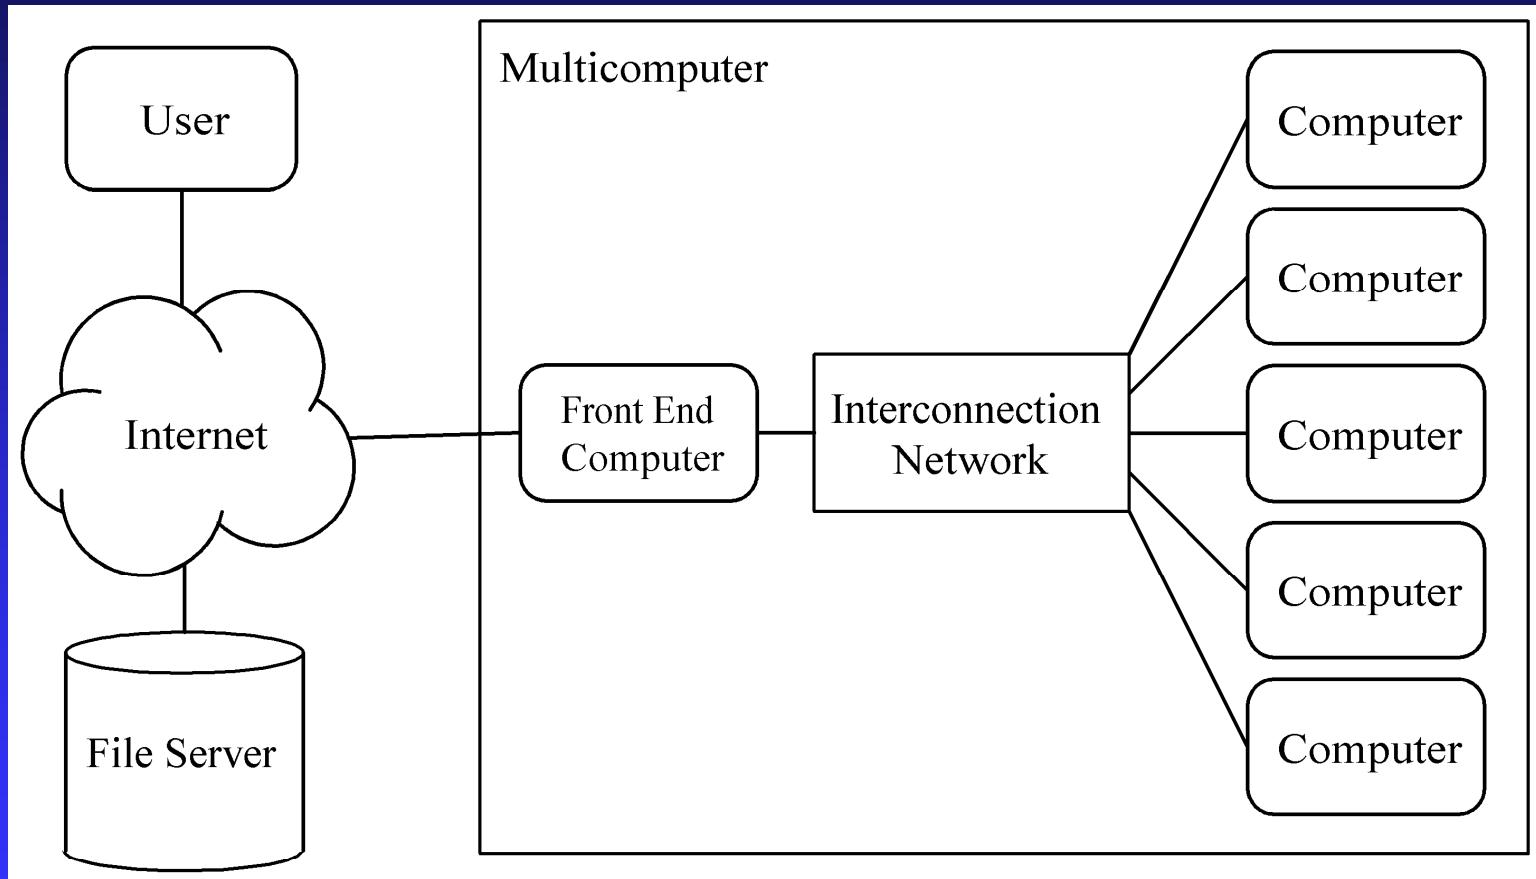
CPU 0 Writes Back X Block



Multicomputer

- Distributed memory multiple-CPU computer
- Same address on different processors refers to different physical memory locations
- Processors interact through message passing
- Commercial multicomputers
- Commodity clusters

Asymmetrical Multicomputer



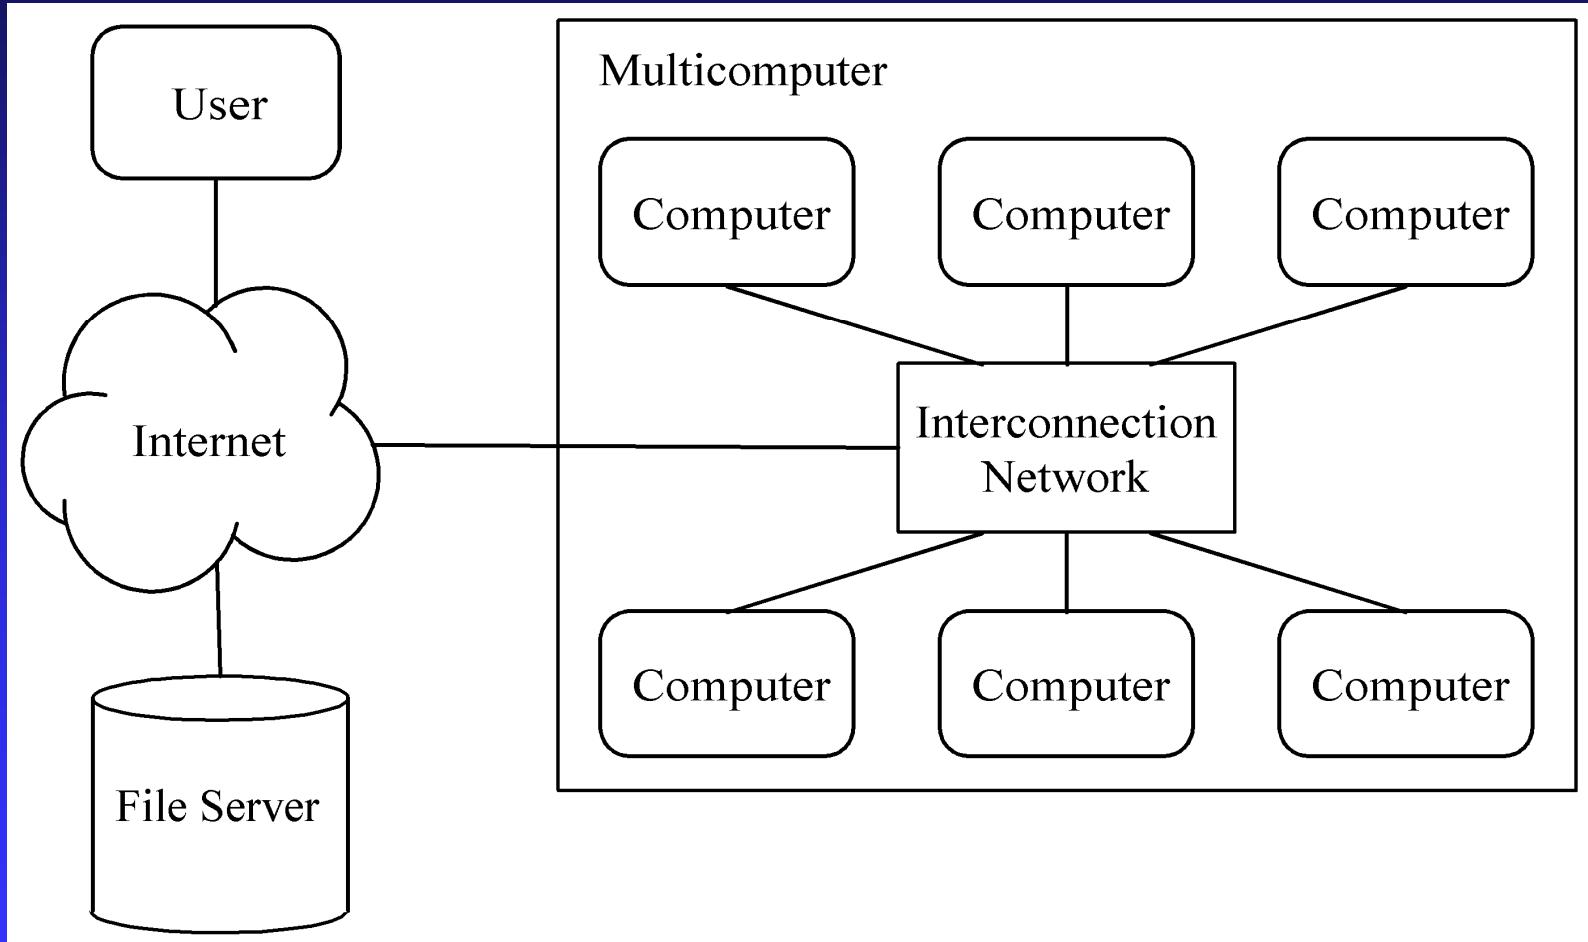
Asymmetrical MC Advantages

- Back-end processors dedicated to parallel computations ⇒ Easier to understand, model, tune performance
- Only a simple back-end operating system needed ⇒ Easy for a vendor to create

Asymmetrical MC Disadvantages

- Front-end computer is a single point of failure
- Single front-end computer limits scalability of system
- Primitive operating system in back-end processors makes debugging difficult
- Every application requires development of both front-end and back-end program

Symmetrical Multicomputer



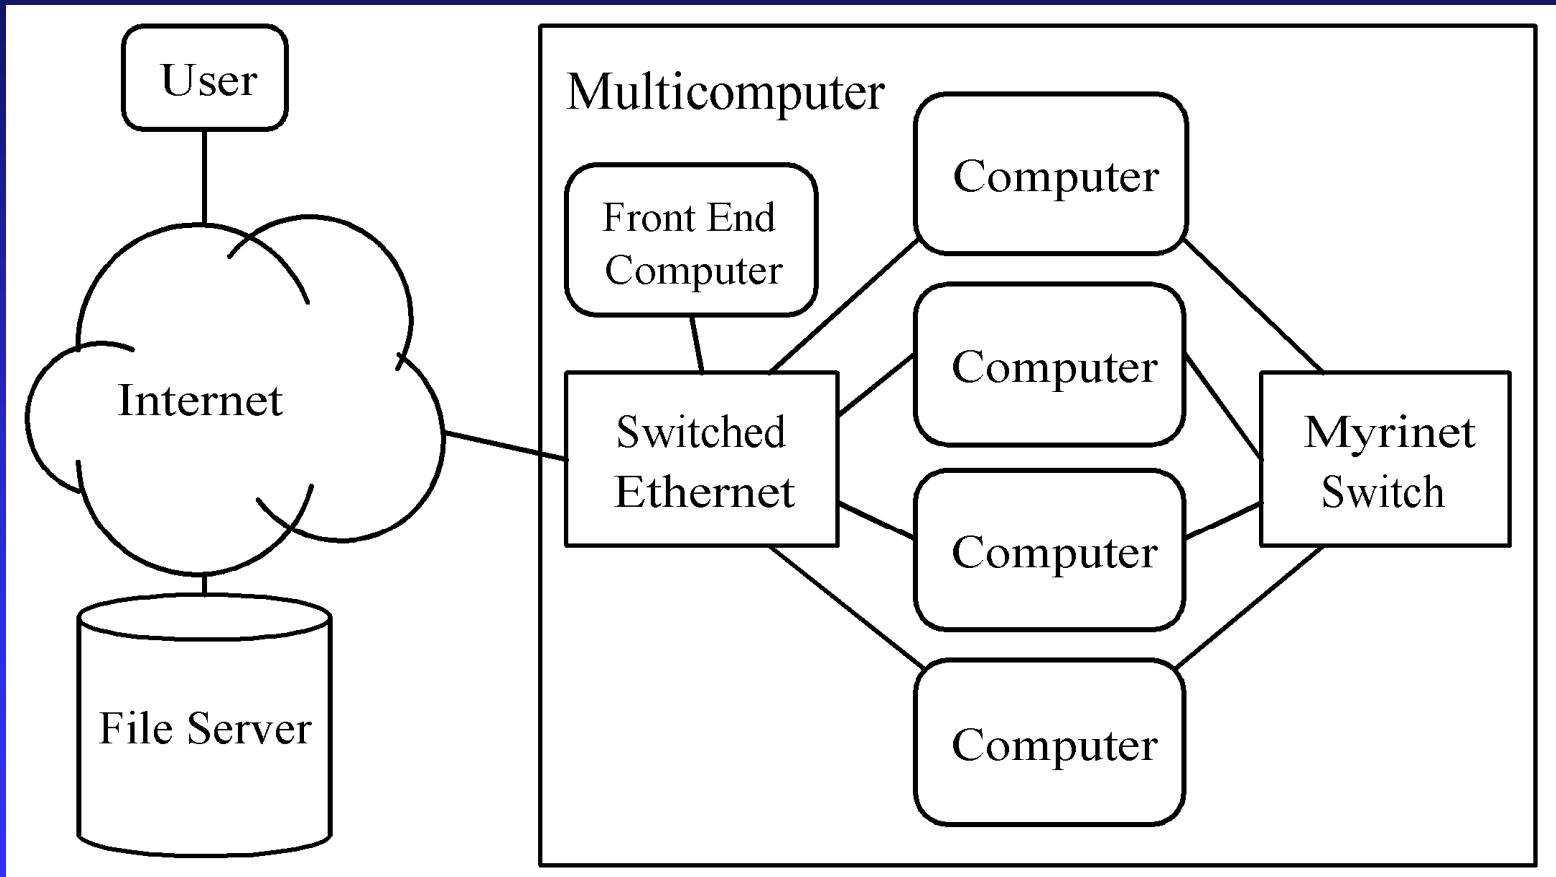
Symmetrical MC Advantages

- Alleviate performance bottleneck caused by single front-end computer
- Better support for debugging
- Every processor executes same program

Symmetrical MC Disadvantages

- More difficult to maintain illusion of single “parallel computer”
- No simple way to balance program development workload among processors
- More difficult to achieve high performance when multiple processes on each processor

ParPar Cluster, A Mixed Model



Commodity Cluster

- Co-located computers
- Dedicated to running parallel jobs
- No keyboards or displays
- Identical operating system
- Identical local disk images
- Administered as an entity

Network of Workstations

- Dispersed computers
- First priority: person at keyboard
- Parallel jobs run in background
- Different operating systems
- Different local images
- Checkpointing and restarting important

Flynn's Taxonomy

- Instruction stream
- Data stream
- Single vs. multiple
- Four combinations
 - ◆ SISD
 - ◆ SIMD
 - ◆ MISD
 - ◆ MIMD

SISD

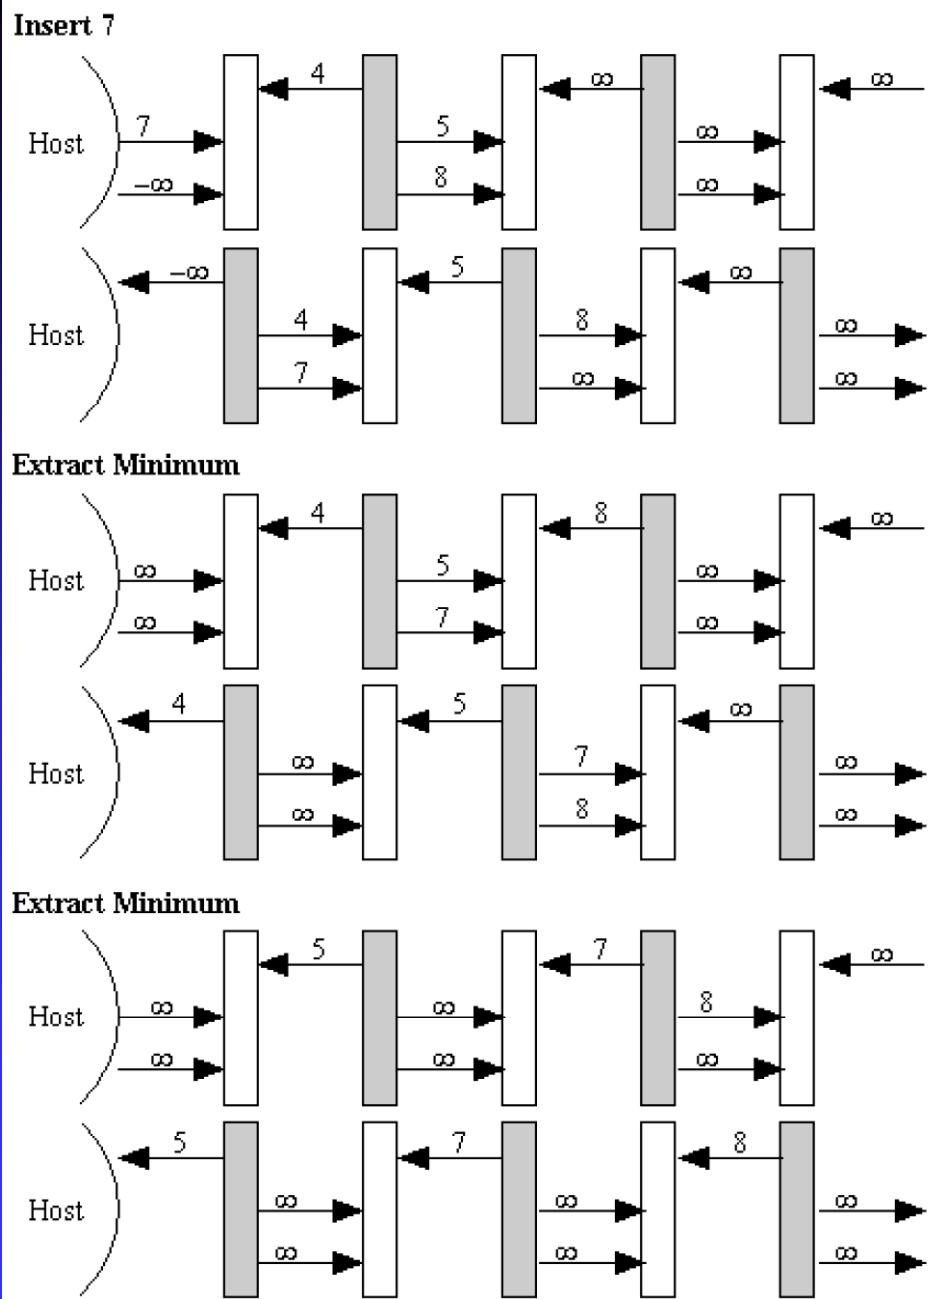
- Single Instruction, Single Data
- Single-CPU systems
- Note: co-processors don't count
 - ◆ Functional
 - ◆ I/O
- Example: PCs

SIMD

- Single Instruction, Multiple Data
- Two architectures fit this category
 - ◆ Pipelined vector processor
(e.g., Cray-1)
 - ◆ Processor array
(e.g., Connection Machine)

MISD

- Multiple Instruction, Single Data
- Example: systolic array



MIMD

- Multiple Instruction, Multiple Data
- Multiple-CPU computers
 - ◆ Multiprocessors
 - ◆ Multicomputers

Summary

- Commercial parallel computers appeared in 1980s
- Multiple-CPU computers now dominate
- Small-scale: Centralized multiprocessors
- Large-scale: Distributed memory architectures (multiprocessors or multicamputers)